

## Список литературы

- [AAH<sup>+</sup>85] M. W. Alford, J. P. Ansart, G. Hommel, L. Lamport, B. Liskov, G. P. Mullery, and F. B. Schneider. *Distributed Systems - Methods and Tools for Specification - An Advanced Course, Lecture Notes in Computer Science*. Springer-Verlag, 1985.
- [AB87] N. Ahituv and O. Berman. Devising a cooperation policy for emergency networks. *Journal of the Operational Research Society*, pages 1015–1029, 1987.
- [AB93] Paul Adler and Bryan Borys. Materialism and idealism in organizational research. *Organization Studies*, 14(5):657–669, 1993.
- [ABCL90] J. R. Anderson, C. F. Boyle, A. T. Corbett, and M. W. Lewis. Cognitive modeling and intelligent tutoring. *Artificial Intelligence*, 42(1):7–50, February 1990.
- [ABD85] James K. Aragonés, Piero P. Bonissone, and Keith S. Decker. Lotta: An object based simulator for reasoning in antagonistic situations. Working Paper, GE Corporate Research and Development, 1985.
- [Abr63] Norman Abramson. *Information Theory and Coding*. McGraw-Hill, New York, 1963.
- [AC87] P. Agre and D. Chapman. Pengi: An implementation of a theory of activity. In *Proceedings of the National Conference on Artificial Intelligence*, pages 268–272, 1987.
- [AC88] Philip E. Agre and David Chapman. What are plans for? Technical Report AI Memo 1050, Massachusetts Institute of Technology, September 1988.
- [ACC<sup>+</sup>88] Daniel E. Atkins, Michael D. Cohen, Lynn A. Conway, Gary M. Olson, Judith R. Olson, John Laird, and Steven Lytinen. A center for collaboration science and technology. NSF Proposal, 1988.
- [ACF<sup>+</sup>84] J. M. Abram, C. Y. Chong, M. R. Fehling, J. S. Rosenschein, and E. Tse. Distributed decision making environment. Technical Report TR-3013-3, AI & DS, April 1984.
- [ACP95] Matthew Anderton, Jim Cunningham, and Jeremy Pitt. A framework for multi-agent inter-organizational applications: A position paper. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 435, San Francisco, CA, 1995. MIT Press. (poster).
- [ACR87] John R. Anderson, Albert T. Corbett, and Brian J. Reiser. *Essential Lisp*. Addison-Wesley, 1987.
- [ADWW89] Mark R. Adler, Alvah B. Davis, Robert Weihmayer, and Ralph Worrest. Conflict-resolution strategies for nonhierarchical distributed agents. In Michael N. Huhns and Les Gasser, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*, pages 139–162. Pitman, 1989. Also COINS Technical Report 88-89, University of Massachusetts, 1988.
- [AES90] C. Applegate, C. Elsaesser, and J. Sanborn. An architecture for adversarial planning. *IEEE Transactions on Systems, Man, and Cybernetics*, 20(1):186–194, January 1990.
- [Agh86] G. Agha. *Actors—A Model of Concurrent Computation for Distributed Systems*. MIT Press, 1986.
- [Agr91] Philip E. Agre. Title undecided. To be published by Cambridge University Press. Position paper presented at AAAI Fall Symposium on Knowledge and Action at Social and Organizational Levels, 1991.
- [AH95] Syed S. Ali and Susan Haller. Interpreting spread sheet data for human-agent interactions. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.

- [Aik80] Jan S. Aikins. *Prototypes and Production Rules: A Knowledge Representation for Computer Consultations*. PhD thesis, Stanford University, 1980. (Also published as Technical Report STAN-CSD-80-814, Computer Science Department, Stanford University.).
- [AJ95] Yongyuth Aramkulchai and John Yuejun Jiang. Generalised proof-theoretic multi-agent autoepistemic reasoning. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 304–311, San Francisco, CA, 1995. MIT Press.
- [AK83] James F. Allen and Johannes A. Koomen. Planning using a temporal world model. In *Proceedings of the Eighth International Joint Conference on Artificial Intelligence*, pages 741–747, August 1983.
- [AKA91] David W. Aha, Dennis Kibler, and Marc K. Albert. Instance-based learning algorithms. *Machine Learning*, 6(1):37–66, January 1991.
- [All71] Graham T. Allison. *Essence of Decision: Explaining the Cuban Missile Crisis*. Little, Brown and Company, Boston, 1971.
- [All78] John R. Allen. *Anatomy of Lisp*. McGraw-Hill, 1978.
- [All79] James F. Allen. *A Plan-Based Approach to Speech Act Recognition*. PhD thesis, University of Toronto, February 1979. (Also published as Technical Report 131/79, Department of Computer Science, University of Toronto, Toronto, Canada, February 1979.).
- [All83] James F. Allen. Maintaining knowledge about temporal intervals. *Communications of the ACM*, 26(11):832–843, November 1983.
- [All84] James F. Allen. Towards a general theory of action and time. *Artificial Intelligence*, 23:123–154, 1984.
- [ALS95] N. M. Avouris, M. H. V. Liederkerke, and L. Sommaruga. CooperA: A cooperating expert system shell and its testing in the domain of environmental emergencies management. In G. O’Hare and N. Jennings, editors, *Foundations of Distributed Artificial Intelligence*, chapter 14. Wiley Inter-Science, 1995. Forthcoming.
- [Amb95] Stanislaw Ambroszkiewicz. Knowledge and behavior in multiagent systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 434, San Francisco, CA, 1995. MIT Press. (poster).
- [And95] David Andre. The evolution of agents that build mental models and create simple plans using genetic programming. In *Proceedings of the Sixth International Conference on Genetic Algorithms*, pages 248–255. Morgan Kaufmann Publishers, Inc., 1995.
- [Ant65] Robert N. Anthony. Planning and control systems: A framework for analysis. Technical report, Division of Research, Harvard Business School, Harvard University, 1965.
- [App82] D. E. Appelt. Planning natural language utterances to satisfy multiple goals. Technical Note 259, SRI International, Menlo Park, CA, 1982.
- [Ara95] Sachiyo Arai. Rational balancing of planning and communication in the dynamic multi-agents world. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 436, San Francisco, CA, 1995. MIT Press. (poster).
- [Arb72] Michael A. Arbib. *The Metaphorical Brain*. Wiley, 1972.
- [ARH87] Ronald C. Arkin, Edward M. Riseman, and Allen R. Hanson. ArRA: An architecture for vision-based robot navigation. In *Proceedings of the DARPA Image Understanding Workshop*, pages 17–431, Los Angeles, CA, February 1987.

- [Ark87] Ronald C. Arkin. *Towards Cosmopolitan Robots: Intelligent Navigation in Extended Man-Made Environments*. PhD thesis, University of Massachusetts, September 1987.
- [AS85] Harold Abelson and Gerald Jay Sussman. *Structure and Interpretation of Computer Programs*. MIT Press and McGraw-Hill (text orders), 1985.
- [AS90] Mark R. Adler and Evangelos Simoudis. Integrating distributed expertise. In *Proceedings of the 1990 Distributed AI Workshop*, Bandara, TX, October 1990.
- [AS95] Michel Aubé and Alain Senteni. A foundation for commitments as resource management in multi-agent systems. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [AS96a] Neeraj Arora and Sandip Sen. Resolving social dilemmas using genetic algorithms. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [AS96b] Michel Aubi and Alain Senteni. Emotions as commitments operators: A foundation for control structure in multi-agents systems. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [Ash88] Kevin D. Ashley. *Modelling Legal Argument: Reasoning with Cases and Hypotheticals*. PhD thesis, University of Massachusetts, February 1988.
- [AUH95] Minoru Asada, Eiji Uchibe, and Koh Hosoda. Agents that learn from other competitive agents. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [aut77] Various authors. Special issue on microelectronics. *Scientific American*, 237(3):62–245, September 1977.
- [AUY79] Alfred V. Aho, J. D. Ullman, and M. Yannakakis. Modeling communications protocols by automata. In *Proceedings of the 20th Symposium on the Foundations of Computer Science*, pages 267–273, October 1979.
- [AVA95] Pierre Azema, Francois Vernadat, and Jean Luc Albacete. A communication protocol for conflict resolution. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 437, San Francisco, CA, 1995. MIT Press. (poster).
- [Axe84] Robert Axelrod. *The Evolution of Cooperation*. Basic Books, 1984.
- [AZL<sup>+</sup>86] R. Adelman, L. Zirk, D. Lehner, P. Moffet, and R. Hall. Distributed tactical decision making: Conceptual framework and empirical results. *IEEE Transactions on Systems, Man and Cybernetics*, pages 794–804, 1986.
- [BA95] Tucker Balch and Ronald C. Arkin. Motor schema-based formation control for multiagent robot teams. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 10–16, San Francisco, CA, 1995. MIT Press.
- [BA96] Omar Belakhdar and Jacqueline Ayel. Modelling approach and tool for designing protocols for automated cooperation in a multi-agent systems. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [BAC<sup>+</sup>87] R. Bisiani, F. Alleva, F. Correrini, A. Forin, F. Lecouat, and R. Lerner. Heterogeneous parallel processing: The Agora shared memory. Technical Report CMU-CS-87-112, Computer Science Department, Carnegie-Mellon University, Pittsburgh, PA, March 1987.

- [BAF<sup>+</sup>87] R. Bisiani, F. Alleva, A. Forin, R. Lerner, and M. Bauer. The architecture of the agora environment. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, Research Notes in Artificial Intelligence, chapter 4, pages 99–118. Pitman, 1987.
- [Bak74] K. R. Baker. *Introduction to Sequencing and Scheduling*. John Wiley & Sons, 1974.
- [Bak88] A. D. Baker. Complete manufacturing control using a contract net: A simulation study. In *Proceedings of the 1988 International Conference on Computer Integrated Manufacturing*, pages 100–109, May 1988.
- [Bar81] Jeffrey A. Barnett. Computational methods for a mathematical theory of evidence. In *Proceedings of the Seventh International Joint Conference on Artificial Intelligence*, pages 868–875, Tokyo, Japan, 1981.
- [Bar82] Jeffrey A. Barnett. Some issues of control in expert systems. Technical report, USC/Information Sciences Institute, Marina del Rey, CA, 1982.
- [Bar84] Jeffrey A. Barnett. How much is control knowledge worth: A primitive example. *Artificial Intelligence*, 22:77–89, 1984.
- [Bar85] Andrew G. Barto. Learning by statistical cooperation of self-interested neuron-like computing elements. Technical Report 85-11, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, April 1985.
- [Bat86] Peter Charles Bates. *Debugging Programs in a Distributed System Environment*. PhD thesis, University of Massachusetts, February 1986. (Also published as Technical Report 86-05, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, February 1986.).
- [Bat96] Joe Bates. Believable autonomous agents. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [Bau78] Gerard M. Baudet. Asynchronous iterative methods for multiprocessors. *Journal of the ACM*, 25(2):226–244, April 1978.
- [Baz95] Ana L. C. Bazan. A game-theoretic approach to distributed control of traffic signals. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 439, San Francisco, CA, 1995. MIT Press. (poster).
- [BBMN91] J. C. Bean, J. R. Birge, J. Mittenthal, and C. E. Noon. Matchup scheduling with multiple resources, release dates and disruptions. *Operations Research*, 39(3):470–483, May-June 1991.
- [BC86] David C. Brown and B. Chandrasekaran. Knowledge and control for a mechanical design expert system. *IEEE Computer*, pages 92–100, July 1986.
- [BC95a] Chumki Basu and Paul B. Cantor. When just one is enough. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [BC95b] Will Briggs and Diane Cook. Flexible social laws. In Chris S. Mellish, editor, *Proceedings of the Fourteenth International Joint Conference on Artificial Intelligence*, pages 688–693, Montréal, Canada, August 1995.
- [BD86] Piero P. Bonissone and Keith S. Decker. Selecting uncertainty calculi and granularity: An experiment in trading-off precision and complexity. In L. N. Kanal and J. F. Lemmer, editors, *Uncertainty in Artificial Intelligence*. North Holland, 1986.
- [BD89] Mark Boddy and Thomas Dean. Solving time-dependent planning problems. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, pages 979–984, Detroit, MI, August 1989.

- [BDA96] Christof Baeijs, Yves Demazeau, and Luis Alvares. SIGMA: Application of multi-agent systems to cartographic generalization. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [BDC<sup>+</sup>95] Jeffrey M. Bradshaw, Stewart Dutfield, Bob Carpenter, Renia Jeffers, and Tom Robinson. KAoS: A generic agent architecture for aerospace applications. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [BDJ86] L. Baum, R. Dodhiawala, and V. Jagannathan. Boeing Blackboard System, version 1.0. Technical Report BCS-G2010-31, Boeing Computer Services, P.O. Box 24346, Seattle, Washington 98124, July 1986.
- [BDKJT95] Frances Brazier, Barbara Dunin-Keplicz, Nick R. Jennings, and Jan Treur. Formal specification of multi-agent systems: a real-world case. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 25–32, San Francisco, CA, 1995. MIT Press.
- [BdRCHV95] Rafael Heitor Bordini, Antônio Carlos da Rocha Costa, Jomi Fred Hübner, and Rosa Maria Viccari. Linguistic support for agent migration. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 441, San Francisco, CA, 1995. MIT Press. (poster).
- [BE82] Jeffrey A. Barnett and Lee D. Erman. Making control decisions in an expert system is a problem-solving task. Unfinished technical report, USC/Information Sciences Institute, Marina del Rey, CA, April 1982.
- [BE89] H. Berliner and C. Ekling. Pattern knowledge and search : The SUPREM architecture. *Artificial Intelligence*, 38(2):?, March 1989.
- [BE95a] Ciara Byrne and Pete Edwards. Refinement in agent groups. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 1–7, 1995.
- [BE95b] Ciara Byrne and Peter Edwards. Refining the knowledge of multiple agents. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [Bee78] Stafford Beer. *Platform for Change*. John Wiley & Sons, corrected reprint edition, 1978.
- [Bee79] Stafford Beer. *The Heart of Enterprise*. John Wiley & Sons, 1979.
- [Bee81] Stafford Beer. *Brain of the Firm*. John Wiley & Sons, second edition, 1981.
- [Bel89] Collin E. Bell. Maintaining project networks in automated artificial intelligence planning. *Management Science*, 35(10):1192–1214, October 1989.
- [BELW80] Robert Balzer, Lee Erman, Philip London, and Chuck Williams. Hearsay-III: A domain-independent framework for expert systems. In *Proceedings of the First Annual National Conference on Artificial Intelligence*, pages 108–110, Stanford, CA, August 1980.
- [BF79] Jon L. Bentley and Jerome H. Friedman. Data structures for range searching. *Computing Surveys*, 11(4):397–413, December 1979.
- [BF95] Mihai Barbuceanu and Mark S. Fox. COOL: A language for describing coordination in multiagent systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 17–24, San Francisco, CA, 1995. MIT Press.
- [BF96] Lawrence Bull and Terence C. Fogarty. Evolution in cooperative multiagent environments. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.

- [BFL83] Ronald J. Brachman, Richard E. Fikes, and Hector J. Levesque. Krypton: A functional approach to knowledge representation. *IEEE Computer*, pages 67–73, October 1983.
- [BFS91] T. Bouron, J. Ferber, and F. Samuel. MAGES: A multi-agent testbed for heterogeneous agents. In Y. Demazeau and J-P Muller, editors, *Decentralised Artificial Intelligence 2*. Elsevier, 1991.
- [BG88a] Alan H. Bond and Les Gasser. An analysis of problems and research in DAI. In Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 3–35. Morgan Kaufmann Publishers, San Mateo, CA, 1988.
- [BG88b] Alan H. Bond and Les Gasser, editors. *Readings in Distributed Artificial Intelligence*. Morgan Kaufmann, 1988.
- [BGD87] Piero P. Bonissone, Steven S. Gans, and Keith S. Decker. RUM: A layered architecture for reasoning with uncertainty. In *Proceedings of the Tenth International Joint Conference on Artificial Intelligence*, August 1987.
- [BGG84] D. Bertsekas, E. Gafni, and R. Gallager. Second derivative algorithms for minimum delay routing in networks. *IEEE Transactions on Communications*, COM-32(8), 1984.
- [BGS<sup>+</sup>91] P. Brazdil, M. Gams, S. Sian, L. Torgo, and W. van de Velde. Learning in distributed systems and multi-agent environments. In *European Working Session on Learning*, Lecture Notes in AI, 482, Berlin, March 1991. Springer Verlag.
- [BH82] P. G. Bennett and C. S. Huxham. Hypergames and what they do: A ‘soft OR’ approach. *Journal of Operational Research Society*, 33:41–50, 1982.
- [BH90] David Murray Bridgeland and Michael N. Huhns. Distributed truth maintenance. In *Proceedings of the Eighth National Conference on Artificial Intelligence*, pages 72–77, July 1990.
- [Bha95] Raghu R. Bhat. An agent system for case adaptation. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 440, San Francisco, CA, 1995. MIT Press. (poster).
- [BHL87] C. A. Broverman, K. E. Huff, and V. R. Lesser. The role of plan recognition in design of an intelligent user interface. In *Proceedings of the IEEE Systems, Man, and Cybernetics Conference*, pages 863–868, Atlanta, Georgia, 1987.
- [BHS<sup>+</sup>91] E. Busch, M. Hamalainen, Y. Suh, A Whinston, and C. Holsapple. Issues and obstacles in the development of team support systems. *Journal of Organizational Computing*, 2(1):161–186, 1991.
- [Bis86] R. Bisiani. A software and hardware environment for developing AI applications on parallel processors. In *Proceedings of the National Conference on Artificial Intelligence*, pages 742–747, Philadelphia, PA, August 1986. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 451–456, Morgan Kaufmann, 1988.).
- [BJD86] M. Benda, V. Jagannathan, and R. Dodhiawala. On optimal cooperation of knowledge sources - an empirical investigation. Technical Report BCS-G2010-28, Boeing Advanced Technology Center, Boeing Computing Services, Seattle, Washington, July 1986.
- [BJSS86] Kenneth P. Birman, Thomas A. Joseph, Frank Schmuck, and Pat Stephenson. Programming with shared bulletin boards in asynchronous distributed systems. Department of Computer Science, Cornell Univ., Ithaca, NY, December 1986.
- [BK96] H. H. Bui and D. Kieronska. Negotiating agents that learn about other’s preferences. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.

- [BL79] Richard S. Brooks and Victor R. Lesser. Distributed problem solving using iterative refinement. Technical Report 79-14, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, May 1979.
- [BM76] J. A. Bondy and U. S. R. Murty. *Graph Theory with Applications*. North-Holland, New York, NY, 1976.
- [BMM95] Charles Brooks, W. Scott Meeks, and Murray S. Mazer. An architecture for supporting quasi-agent entities in the WWW. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [BMS86] Daniel G. Bobrow, Sanjay Mittal, and Mark J. Stefik. Expert system : Perils and promise. *Communications of the ACM*, 29(9):880–894, September 1986.
- [BN78] Bertram Bruce and Dennis Newman. Interacting plans. *Cognitive Science*, pages 195–233, 1978. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 248–267, Morgan Kaufmann, 1988.).
- [BN95a] Sivaram Balasubramanian and Douglas H. Norrie. A multi-agent intelligent design system integrating manufacturing and shop-floor control. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 3–9, San Francisco, CA, 1995. MIT Press.
- [BN95b] Orlando Belo and Jose Neves. Cooperation among opportunistic agents on a distributed environment. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [BO84] Richard M. Burton and Børge Obel. *Designing Efficient Organizations: Modelling and Experimentation*. North Holland, Amsterdam, 1984.
- [Bob91] D. Bobrow. Dimensions of interaction: A shift of perspective in artificial intelligence. *AI Magazine*, 12(3), 1991.
- [Boe79] W. Boehm. Software engineering: R and D trends and defense needs. In P. (Ed) Wegener, editor, *Research Directions in Software Technology*. M. I. T. Press, Cambridge, Mass., 1979.
- [Boe88] W. Boehm. A spiral model of software development and enhancement. *IEEE Computer*, pages 61–72, May 1988.
- [Bog83] K. P. Bogart. *Introductory Combinatorics*. Pitman, Boston, MA, 1983.
- [Bon86a] Piero P. Bonissone. Plausible reasoning: Coping with uncertainty in expert systems. In S. C. Shapiro, editor, *Encyclopedia of Artificial Intelligence*. John Wiley and Sons, 1986.
- [Bon86b] Piero P. Bonissone. Reasoning with uncertain and incomplete information: DARPA demonstration road map. Working Paper, GE Corporate Research and Development, 1986.
- [Bon89] Alan H. Bond. The cooperation of experts in engineering design. In Michael N. Huhns and Les Gasser, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [Bon90] Alan H. Bond. Commitment: A computational model for organizations of cooperating intelligent agents. In *Proceedings of the 1990 Conference on Office Automation Systems*, pages 21–30, Cambridge, MA, 1990.
- [BPA<sup>+</sup>95] Uwe Borghoff, Remo Pareschi, Jean-Marc Andreoli, Francesca Arcelli, Ferrante Formato, and Hans Schlichter. Agent-based document construction. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.

- [Bra83] Ronald J. Brachman. What IS-A is and isn't: An analysis of taxonomic links in semantic networks. *IEEE Computer*, pages 30–36, October 1983.
- [BRA<sup>+</sup>94] G. Burrell, M. Reed, M. Alverson, M. Calás, and L. Smircich. Why Organization? why now? *Organization*, 1(1):5–17, 1994.
- [Bro80] Richard S. Brooks. A balance principle for optimal access control. Technical Report 80-20, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, November 1980.
- [Bro83] Richard S. Brooks. *Experiments in Distributed Problem Solving with Iterative Refinement*. PhD thesis, University of Massachusetts, February 1983. (Also published as Technical Report 82-25, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, October 1982.).
- [Bro85] Rodney A. Brooks. *Programming in Common Lisp*. John Wiley & Sons, 1985.
- [Bro86a] Hank Bromley. *Lisp Lore: A guide to programming on the Lisp machine*. Kluwer Academic Publishers, 1986.
- [Bro86b] Rodney A. Brooks. A robust layered control system for a mobile robot. *IEEE Journal of Robotics and Automation*, RA-2(1):14–23, March 1986.
- [Bro95] Shirley Browne. Need for a security profile for agent environments. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [BS84] P. M. Blau and J. E. Schwartz. *Crosscutting Social Circles: Testing a Macrostructural Theory of Group Relationships*. Academic Press, New York, 1984.
- [BT95] Joachim Baczewski and Mario Tokoro. Linear prediction in a multi-agent environment. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 438, San Francisco, CA, 1995. MIT Press. (poster).
- [BT96] Wolfgang Balzer and Raimo Tuomela. The structure and build-up of plan-based joint intentions. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [BTvE96] Frances Brazier, Jan Treur, and Pascal van Eck. Modelling cooperative behaviour for resource access. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [Bud91] Timothy A. Budd. *An Introduction to Object-Oriented Programming*. Addison-Wesley, Reading, Mass., 1991.
- [Bur82] R. S. Burt. *Toward a structural theory of action : network models of social structure, perception, and action*. Quantitative studies in social relations. Academic Press, New York, 1982.
- [Bur91] Mark Burstein. Integrated feasibility demonstrations within the DARPA/RL planning and scheduling initiative. Section 5 of the Working Group Reports for the DARPA/RL Knowledge Based Scheduling and Planning Initiative, 1991.
- [BV94] J. A. Breuker and W. Van de Velde, editors. *CommonKADS Library for Expertise Modelling: Reusable problem solving components*, volume 21 of *Frontiers in Artificial Intelligence and Applications*. IOS-Press, Amsterdam, August 1994.
- [BW82] Peter C. Bates and Jack C. Wileden. Event Definition Language: An aid to monitoring and debugging of complex software systems. In *Proceedings of the Fifteenth Hawaii International Conference on System Science*, pages 86–93, January 1982.
- [BW83] Peter C. Bates and Jack C. Wileden. High-level debugging of distributed systems: The behavioral abstraction approach. *Journal of Systems and Software*, 3(4):255–264, 1983.



- [BWL81] Peter Bates, Jack C. Wileden, and Victor R. Lesser. A language to support debugging in distributed systems. Technical Report 81-07, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, February 1981.
- [BWL82] Peter C. Bates, Jack C. Wileden, and Victor R. Lesser. A debugging tool for distributed systems. Technical Report 82-34, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, December 1982.
- [BWL83] P. C. Bates, J. C. Wileden, and V. R. Lesser. A debugging tool for distributed systems. In *Proceedings of the Second Annual Phoenix Conference on Computers and Communications*, pages 311–315, March 1983.
- [BZ83] Daniel Brand and Pitro Zaffropulo. On communicating finite-state machines. *Journal of the ACM*, 30(2):323–342, April 1983.
- [Car88] Norman Carver. Evidence based plan recognition. COINS Technical Report 88-13, University of Massachusetts, 1988.
- [Car93] Claire Cardie. Using decision trees to improve case-based learning. In *Proceedings of the Tenth International Conference on Machine Learning*, pages 25–32. Morgan Kaufmann Publishers, Inc., 1993.
- [Car95] Bengt Carlsson. An evolutionary model of multi-agent systems. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [Cas93] C. Castelfranchi. Commitments: from individual intentions to groups and organizations. In Michael Prietula, editor, *AI and theories of groups & organizations: Conceptual and Empirical Research*. AAAI Workshop, 1993. Working Notes.
- [Cas95] Cristiano Castelfranchi. Commitments: from individual intentions to groups and organizations. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 41–48, San Francisco, CA, 1995. MIT Press.
- [CBK89] Gregg Collins, Lawrence Birnbaum, and Bruce Krulwich. An adaptive model of decision-making in planning. In *IJCAI*, pages 511–516, 1989.
- [CC96] R. Conte and C. Castelfranchi. From conventions to prescriptions. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [CCC95] Jennifer Chu-Carroll and Sandra Carberry. Communication for conflict resolution in multi-agent collaborative planning. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 49–56, San Francisco, CA, 1995. MIT Press.
- [CCL91] Norman Carver, Zarko Cvetanovic, and Victor Lesser. Sophisticated cooperation in FA/C distributed problem solving systems. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 191–198, Anaheim, July 1991.
- [CDJ84] F. B. Chambers, D. A. Duce, and G. P. Jones, editors. *Distributed Computing*, volume 20 of *APIC Studies in Data Processing*. Academic Press, London, 1984.
- [CdMM92] B. Chaib-draa, R. Mandiau, and P. Millot. Distributed artificial intelligence: An annotated bibliography. *SIGART Bulletin*, 3(3):20–37, August 1992.
- [CG88] Daniel D. Corkill and Kevin Q. Gallagher. Tuning a blackboard-based application: A case study using GBB. In *Proceedings of the National Conference on Artificial Intelligence*, pages 671–676, St. Paul, MN, August 1988.
- [CG90] R. Clark and C. Grossner. Generation of organizations for cooperating expert systems. DAI Technical Report DAI-0290-0002, Concordia University, Montreal Quebec, February 1990.

- [CG91] Eugene Charniak and Robert Goldman. A probabilistic model of plan recognition. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 160–165, Anaheim, CA, July 1991.
- [CGHH89] Paul R. Cohen, Michael L. Greenberg, David M. Hart, and Adele E. Howe. Trial by fire: Understanding the design requirements for agents in complex environments. *AI Magazine*, 10(3):32–48, Fall 1989. Also COINS-TR-89-61.
- [CGJ86] Daniel D. Corkill, Kevin Q. Gallagher, and Philip M. Johnson. From prototype to product: Evolutionary development within the blackboard paradigm. Technical Report 86-46, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, October 1986. (Presented at Workshop on High Level Tools for Knowledge Based Systems, Columbus, Ohio, October 7–8, 1986.).
- [CGJ87] Daniel D. Corkill, Kevin Q. Gallagher, and Philip M. Johnson. Achieving flexibility, efficiency, and generality in blackboard architectures. In *Proceedings of the National Conference on Artificial Intelligence*, pages 18–23, Seattle, WA, July 1987. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 451–456, Morgan Kaufmann, 1988.).
- [CGM86a] D. D. Corkill, K. Q. Gallagher, and K. E. Murray, editors. *GBB: A Generic Blackboard Development System*, Philadelphia, August 1986.
- [CGM86b] Daniel D. Corkill, Kevin Q. Gallagher, and Kelly E. Murray. GBB: A generic blackboard development system. In *Proceedings of the Fifth National Conference on Artificial Intelligence*, pages 1008–1014, Philadelphia, PA, August 1986. (Also published in *Blackboard Systems*, Robert S. Englemore and Anthony Morgan, editors, pages 503–518, Addison-Wesley, 1988.).
- [CGR92] R. Clark, C. Grossner, and T. Radhakrishnan. CONSENSUS: A planning protocol for cooperating expert systems. In *11th International Workshop on Distributed Artificial Intelligence, Glen Arbor, Michigan*, February 1992.
- [CGR93] R. Clark, C. Grossner, and T. Radhakrishnan. CONSENSUS and COMPROMISE: Planning in cooperating expert systems. In *Submitted for review to Int. Journal of Intelligent and Cooperative Information Systems*, 1993.
- [cH80] Yu chi Ho. Team decision theory and information structures. *Proceedings of the IEEE*, 68(6):644–654, June 1980.
- [CH89] Caroline Clarke-Hayes. A model of planning for plan efficiency: Taking advantage of operator overlap. In *IJCAI*, pages 949–953, 1989.
- [Cha75] Alphonse Chapanis. Interactive human communication. *Scientific American*, 232(3):36–42, March 1975.
- [Cha81] B. Chandrasekaran. Natural and social system metaphors for distributed problem solving: Introduction to the issue. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):1–5, January 1981.
- [Cha83] B. Chandrasekaran. Towards a taxonomy of problem solving types. *AI Magazine*, 4(1):9–17, Winter-Spring 1983.
- [Cha85] Ernest Chang. Participant systems. In *Proceedings of the 1985 Distributed Artificial Intelligence Workshop*, December 1985.
- [Cha87] Ernest Chang. Participant systems. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, Research Notes in Artificial Intelligence, chapter 11, pages 311–339. Pitman, 1987.
- [Cha93] B. Chandrasekaran. An interview with Allen Newell. *IEEE Expert*, 8(3):5–12, June 1993.

- [CHE] CHECKMATE. Air campaign planning—an approach for the future. White Paper.
- [Che84] Peter Cheeseman. A representation of time for automatic planning. In *Proceedings of the IEEE Conference on Robotics*, pages 513–518, 1984.
- [Che85] Peter Cheeseman. In defense of probability. In *Proceedings of the Ninth International Joint Conference on Artificial Intelligence*, pages 1002–1009, Los Angeles, CA, August 1985.
- [Che95] Brant A. Cheikes. GIA: An agent-based architecture for intelligent tutoring systems. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Chi89] Donald Chisholm. *Coordination Without Hierarchy: Informal Structures in Multiorganizational Systems*. University of California Press, Berkeley, 1989.
- [CHL89] M. Connell, K. E. Huff, and V. R. Lesser. Implementing an incremental hierarchical plan recognition system. In *Proceedings of the 2nd AAAI Workshop on Plan Recognition*, IJCAI-89, Detroit, 1989.
- [CHL91] M. E. Connell, K. E. Huff, and V. R. Lesser. Planning in the context of an intelligent assistant. Technical Report 91-44, Department of Computer Science, University of Massachusetts, 1991.
- [CHR90] Anne Collinot and Barbara Hayes-Roth. Real-time control of reasoning: Experiments with two control models. In *Proceedings of the 1990 DARPA Workshop on Innovative Approaches to Planning, Scheduling, and Control*, pages 263–270, November 1990.
- [Chu77] W. W. Chu. *Advances in Computer Communications*. Artech House, Dedham, Massachusetts, 1977.
- [CI89] Jie Cheng and Keki B. Irani. Ordering problem subgoals. In *IJCAI*, pages 931–936, 1989.
- [CK88] Thomas L. Casavant and Jon G. Kuhl. A taxonomy of scheduling in general-purpose distributed computing systems. *IEEE Transactions on Software Engineering*, 14(2):141–154, February 1988.
- [CK90] Thomas L. Casavant and Jon G. Kuhl. A communicating finite automata approach to modeling distributed computation and its application to distributed decision-making. *IEEE Transactions on Computers*, 39(5):628–639, May 1990.
- [CKLM91] S. E. Conry, K. Kuwabara, V. R. Lesser, and R. A. Meyer. Multistage negotiation for distributed constraint satisfaction. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6):1462–1477, December 1991. (Special Issue on Distributed AI).
- [CL81a] William J. Clancey and Reed Letsinger. NEOMYCIN: Reconfiguring a rule-based expert system for application to teaching. In *Proceedings of the Seventh International Joint Conference on Artificial Intelligence*, pages 829–836, August 1981.
- [CL81b] Daniel D. Corkill and Victor R. Lesser. A goal-directed Hearsay-II architecture: Unifying data-directed and goal-directed control. Technical Report 81-15, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, June 1981.
- [CL83a] Daniel D. Corkill and Victor R. Lesser. Coordination in a distributed problem solving network. In *Proceedings of the Conference on Artificial Intelligence*, pages 637–648, Oakland University, Rochester, Michigan, April 1983.

- [CL83b] Daniel D. Corkill and Victor R. Lesser. The use of meta-level control for coordination in a distributed problem solving network. In *Proceedings of the Eighth International Joint Conference on Artificial Intelligence*, pages 748–756, Karlsruhe, Federal Republic of Germany, August 1983. (Also appeared in *Computer Architectures for Artificial Intelligence Applications*, Benjamin W. Wah and G.-J. Li, editors, IEEE Computer Society Press, pages 507–515, 1986).
- [CL86] Philip R. Cohen and Hector J. Levesque. Persistence, intention, and commitment. In *Proceedings of the 1986 Workshop on Reasoning About Actions and Plans*, July 1986.
- [CL87] Philip R. Cohen and Hector J. Levesque. Intention = choice + commitment. In *Proceedings of the National Conference on Artificial Intelligence*, pages 410–415, July 1987.
- [CL88] W. B. Croft and L. S. Lefkowitz. Knowledge-based support of cooperative activities. In *Proceedings of the Twenty-first Annual Hawaii International Conference on System Sciences*, volume 3, pages 312–318, 1988. (Published by IEEE Computer Society Press, Catalog Number 88TH0213-9.).
- [CL89] Norman Carver and Victor R. Lesser. Planning for the control of an interpretation system. COINS Technical Report 89-39, University of Massachusetts, April 1989.
- [CL90a] P. R. Cohen and H. J. Levesque. Rational interaction as the basis for communication. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*. MIT Press, 1990.
- [CL90b] Philip R. Cohen and Hector J. Levesque. Intention is choice with commitment. *Artificial Intelligence*, 42(3):213–261, 1990.
- [CL91] Norman Carver and Victor Lesser. A new framework for sensor interpretation: Planning to resolve sources of uncertainty. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 724–731, August 1991.
- [CL95a] N. Carver and V. Lesser. The DRESUN testbed for research in FA/C distributed situation assessment: Extensions to the model of external evidence. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 33–40, San Francisco, CA, 1995. MIT Press.
- [CL95b] Philip R. Cohen and Hector J. Levesque. Communicative actions for artificial agents. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 65–72, San Francisco, CA, 1995. MIT Press.
- [Cla83] William J. Clancey. The epistemology of a rule-based expert system—A framework for explanation. *Artificial Intelligence*, 20:215–251, 1983.
- [Cla87] Phil Clark. Office automation: Automation gains municipal ground. *American City & County*, page 10, January 1987.
- [Cla92a] R. Clark. Automatic generation of proposals. DAI Technical Report DAI-0892-0015, Concordia University, Montreal Quebec, August 1992.
- [Cla92b] R. Clark. CONSENSUS: A planning protocol for cooperating expert systems. Master’s thesis, Concordia University, 1992.
- [Cla92c] R. Clark. Distributed artificial intelligence reference list. DAI Technical Report DAI-0692-0014, Concordia University, Montreal Quebec, June 1992.
- [Cla95] Keith Clark. A language for building intelligent distributed information retrieval systems. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM ’95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.

- [CLH82] Daniel D. Corkill, Victor R. Lesser, and Eva Hudlická. Unifying data-directed and goal-directed control: An example and experiments. In *Proceedings of the National Conference on Artificial Intelligence*, pages 143–147, Pittsburgh, PA, August 1982.
- [Cli81] W. Clinger. Foundations of actor semantics. Technical Report 633, MIT Artificial Intelligence Laboratory, Cambridge, MA, May 1981.
- [CLM84] Norman F. Carver, Victor R. Lesser, and Daniel L. McCue. Focusing in plan recognition. In *Proceedings of the National Conference on Artificial Intelligence*, pages 42–48, Austin, TX, August 1984.
- [Clo95] Jeffrey Clouse. Learning from an automated training agent. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [CLR90] Thomas H. Cormen, Charles E. Leiserson, and Ronald L. Rivest. *Introduction to Algorithms*. MIT Press, Cambridge, MA, 1990.
- [CM81] K. M. Chandy and I. Misra. Asynchronous distributed simulation via a sequence of parallel computations. *Comm. of the ACM*, 24:198–206, April 1981.
- [CM88] Daniel D. Corkill and Kelly E. Murray. Issues in common lisp multiprocessing. March 1988. (In preparation.).
- [CM95] David Carmel and Shaul Markovitch. Opponent modeling in a multi-agent system. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 8–13, 1995.
- [CMC91] R. Conte, M. Miceli, and C. Castelfranchi. Limits and levels of cooperation: Disentangling various types of prosocial interaction. In Y. Demazeau and J.-P. Muller, editors, *Decentralized AI*, volume 2. North Holland, 1991.
- [CML86] Susan E. Conry, Robert A. Meyer, and Victor R. Lesser. Multistage negotiation in distributed planning. In *COINS TR86-87*, pages 1–17, 1986. also in Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 367–384. Morgan Kaufman Publishers, Inc., San Mateo, California, 1988.
- [CML88] Susan E. Conry, Robert A. Meyer, and Victor R. Lesser. Multistage negotiation in distributed planning. In Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 367–384. Morgan Kaufmann, 1988.
- [CMM90] Susan E. Conry, Douglas J. MacIntosh, and Robert A. Meyer. DARES: A Distributed Automated REasoning System. In *Proceedings of the National Conference on Artificial Intelligence*, pages 78–85, July 1990.
- [CMP89] S. E. Conry, R. A. Meyer, and R. P. Pope. Mechanisms for assessing nonlocal impact of local decisions in distributed planning. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [CMR96a] A. Cesta, M. Miceli, and P. Rizzo. Comparing interaction attitudes in a simulated multi-agent environment. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [CMR96b] Amedeo Cesta, Maria Miceli, and Paola Rizzo. Effects of different interaction attitudes on a multi-agent system performance. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.

- [CMS83] Stephanie Cammarata, David McArthur, and Randall Steeb. Strategies of cooperation in distributed problem solving. *Proceedings of the 8th International Joint Conference on Artificial Intelligence*, pages 767–770, 1983. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 102–105, Morgan Kaufmann, 1988.).
- [CMS85] S. Conry, R. Meyer, and J. Searlemen. A shared knowledge base for independent problem solving agents. In *Proceedings of the IEEE Expert Systems in Government Symposium*, Mc Lean, Virginia, October 1985.
- [CN95] Grace Crowder and Charles Nicholas. An approach to large scale distributed information systems using statistical properties of text to guide agent search. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Coh78] Philip R. Cohen. *On Knowing What to Say: Planning Speech Acts*. PhD thesis, University of Toronto, January 1978. (Also published as Technical Report 118, Department of Computer Science, University of Toronto, Toronto, Canada, January 1978.).
- [Coh85] Paul R. Cohen. *Heuristic Reasoning about Uncertainty : An Artificial Intelligence Approach*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman Publishing Inc, 1985.
- [Coh88] Paul R. Cohen. How evaluation guides AI research. *AI Magazine*, 9(4):35–43, Winter 1988.
- [Coh91] Paul R. Cohen. A survey of the Eighth National Conference on Artificial Intelligence: Pulling together or pulling apart? *AI Magazine*, 12(1):16–41, Spring 1991.
- [Coh92] Paul R. Cohen. Tactics for generalization. Course notes for Experimental Design in AI, 1992.
- [Col89] Anne Collinot. Revising the BB1 basic control loop to control the behavior of knowledge sources. In V. Jagannathan, Rajendra Dodhiawala, and Lawrence S. Baum, editors, *Blackboard Architectures and Applications*, pages 27–43. Academic Press, Inc., 1989.
- [Cor] Corporation:19xx:AAC Corporation. ACPT—the air campaign planning tool. White Paper.
- [Cor76a] Daniel D. Corkill. PARTNER: Pattern actuated responses to natural English ramblings. *Nebraska Blue Print*, 65(3):4–7, February 1976.
- [Cor76b] Daniel David Corkill. A model for the design of conversational programs. Master’s thesis, University of Nebraska, December 1976.
- [Cor77a] Daniel D. Corkill. FUTIL: FORTRAN utilities for APL files. Technical report, APL Group, University Computing Center, University of Massachusetts, Amherst, MA 01003, July 1977.
- [Cor77b] Daniel D. Corkill. Indexed sequential files for APL. Technical report, APL Group, University Computing Center, University of Massachusetts, Amherst, MA 01003, September 1977.
- [Cor79] Daniel D. Corkill. Hierarchical planning in a distributed environment. In *Proceedings of the Sixth International Joint Conference on Artificial Intelligence*, pages 168–175, Tbilisi, Georgia, USSR, August 1979. (An extended version was published as Technical Report 79-13, Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts 01003, February 1979.).
- [Cor80a] Daniel D. Corkill. *CLisp Reference Manual*. Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts, 1980. (Periodically revised through 1984).

- [Cor80b] Daniel D. Corkill. An organizational approach to planning in distributed problem solving systems. Technical Report 80-13, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, May 1980.
- [Cor83] Daniel David Corkill. *A Framework for Organizational Self-Design in Distributed Problem Solving Networks*. PhD thesis, University of Massachusetts, February 1983. (Also published as Technical Report 82-33, Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts 01003, December 1982.).
- [Cor86] Daniel D. Corkill. An interactive graphic modeling program for arbitrary closed queuing systems with exponential servers. In Lester Lipsky, editor, *Queuing Theory and System Modeling*, pages 1–19. Technical Report, Department of Computer Science, University of Nebraska, May 1986.
- [Cor88] Daniel D. Corkill. Design alternatives for parallel and distributed blackboard systems. Technical Report 88-38, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, August 1988. (Presented at the AAAI-88 Blackboard Workshop, St. Paul, Minnesota, August 24, 1988. Also to appear in *Current Trends in Blackboard Systems*, V. Jagannathan and Raj T. Dodhiawala, editors, Academic Press, 1989.).
- [Cor89] Daniel D. Corkill. Embedable problem solving architectures: A study of integrating OPS5 with GBB. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, Detroit, MI, August 1989. (Submitted.).
- [Cou90] Gregory Courand. Cooperation via consensus formation. In *Proceedings of the Tenth International Workshop on Distributed AI*, Texas, October 1990.
- [CP79] Philip R. Cohen and C. Raymond Perrault. Elements of a plan-based theory of speech acts. *Cognitive Science*, 3(3):177–212, 1979.
- [CP94a] K. M. Carley and M. J. Prietula. ACTS theory: Extending the model of bounded rationality. In K. M. Carley and M. J. Prietula, editors, *Computational Organization Theory*, chapter 4, pages 55–89. Lawrence Erlbaum Associates, Hillsdale, NJ, 1994.
- [CP94b] K. M. Carley and M. J. Prietula, editors. *Computational Organization Theory*. Lawrence Erlbaum Associates, Hillsdale, NJ, 1994.
- [CP95] Stephen Cranefield and Martin Purvis. Agent-based integration of general-purpose tools. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Cra78] Hewitt D. Crane. Beyond the seventh synapse: The neural marketplace of the mind. Research memorandum, Menlo Park, CA, December 1978.
- [Cra80] Hewitt D. Crane. *The New Social Marketplace: Notes on effecting social change in America's third century*. Ablex Publishing, 1980.
- [Cro94] Kevin Crowston. Coordinating restaurant service: An analysis of TGI Friday's process. Presentation given at the TIMS/ORSA Workshop on Mathematical and Computational Organization Theory, April 1994. Crowston is a professor at the University of Michigan School of Business.
- [CS77] L. Casey and N. Shelness. A domain structure for distributed computer systems. In *Proceedings of the Sixth ACM Symposium on Operating System Principles*, pages 101–108, November 1977.
- [CS95] Alessandro Cimatti and Luciano Serafini. Multiagent reasoning with belief contexts II: Elaboration tolerance. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 57–64, San Francisco, CA, 1995. MIT Press.

- [CS96] Stefan Conrad and Gunter Saake. Specifying evolving temporal behaviour. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [CSR86] S. Cheng, J. A. Stankovic, and K. Ramamritham. Dynamic scheduling of groups of tasks with precedence constraints in distributed hard real-time systems. In *Real-time Systems Symposium*, December 1986.
- [CSR88] S. Cheng, J. Stankovic, and K. Ramamritham. Scheduling algorithms for hard real-time systems. In *Hard Real-Time Systems*. IEEE Press, 1988.
- [CT95a] Lawrence Cavedon and Gil Tidhar. A logical framework for multi-agent systems and joint attitudes. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [CT95b] Fumihiko Chimura and Mario Tokoro. Map making as a support for cooperation. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 442, San Francisco, CA, 1995. MIT Press. (poster).
- [Cul81] Richard E. Cullingford. Integrating knowledge sources for computer “understanding” tasks. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):52–60, January 1981.
- [Cul89] C. Culbert. *Artificial Intelligence Section, Johnson Space Center*. Houston, 1989.
- [CW75] R. T. Chien and S. Weissman. Planning and execution in incompletely specified environments. In *Proceedings of the Fourth International Joint Conference on Artificial Intelligence*, pages 169–174, August 1975.
- [CW88] Kai-Hsuing Chang and William G. Wee. A planning model with problem analysis and operator hierarchy. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 10(5):672–675, September 1988.
- [DAD86] R. J. Doyle, D. J. Atkinson, and R. J. Doshi. Generating perception requests and expectations to verify the execution of plans. In *Proceedings of the National Conference on Artificial Intelligence*, pages 81–88, Philadelphia, PA, August 1986.
- [Dan78] W. W. Daniel. *Applied Nonparametric Statistics*. Houghton-Mifflin, Boston, 1978.
- [Dan88] B. M. Dang. Cleopatra: A distributed communication platform. Computer science technical report, Concordia University, Montreal, Quebec, 1988.
- [Dav77] Keith Davis. *Organizational Behavior: A Book of Readings*. McGraw Hill, 1977.
- [Dav80a] Randall Davis. Meta-rules: Reasoning about control. *Artificial Intelligence*, 15:179–222, 1980.
- [Dav80b] Randall Davis. Models of problem solving: Why cooperate? *SIGART Newsletter*, (73):50–51, October 1980.
- [Dav80c] Randall Davis. Report on the workshop on Distributed AI. *SIGART Newsletter*, (73):42–43, October 1980.
- [Dav81] Randall Davis. A model for planning in a multiagent environment: Steps toward principles for teamwork. AI Working Paper 217, Artificial Intelligence Laboratory, Massachusetts Institute of Technology, Cambridge, MA 02139, June 1981.
- [Dav82] Randall Davis. Report on the second workshop on Distributed AI. *SIGART Newsletter*, (80):13–23, April 1982.
- [Dav95] Paul Davidsson. A linearly quasi-anticipatory autonomous agent architecture: Some preliminary experiments. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.



- [Daw76] Richard Dawkins. *The Selfish Gene*. Oxford University Press, 1976.
- [Day95] William B. Day. The consensus scheduling model: Negotiation among peers. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 444, San Francisco, CA, 1995. MIT Press. (poster).
- [DB88a] T. Dean and M. Boddy. Reasoning about partially ordered events. *Artificial Intelligence*, 36(3):375–400, October 1988.
- [DB88b] Thomas Dean and Mark Boddy. An analysis of time-dependent planning. In *Proceedings of the Seventh National Conference on Artificial Intelligence*, pages 49–54, St. Paul, MN, August 1988.
- [DBM<sup>+</sup>92] Lisa Dent, Jesus Boticario, John McDermott, Tom Mitchell, and David Zabowski. A personal learning apprentice. In *Proceedings of the Tenth National Conference on Artificial Intelligence*, pages 96–103, July 1992.
- [DBS77] Randall Davis, Bruce Buchanan, and Edward Shortliffe. Production rules as a representation for a knowledge-based consultation program. *Artificial Intelligence*, 8(1):15–45, 1977.
- [DC89] Mark Drummond and Ken Currie. Goal ordering in partially ordered plans. In *IJCAI*, pages 960–965, 1989.
- [DC96] Paul Dongha and Cristiano Castelfranchi. Rational commitment in resource bounded agents. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [DCCL91] J. Doran, H. Carvajal, Y. J. Choo, and Y. Li. The MCS multi-agent testbed: Developments and experiments. In S. M. Deen, editor, *Cooperating Knowledge Based Systems 1990*, pages 240–251. Springer-Verlag, 1991.
- [DCL84] E. H. Durfee, D. D. Corkill, and V. R. Lesser. Distributing a distributed problem solving network simulator. In *Proceedings of the IEEE Computer Society’s Fifth Real-Time Systems Symposium*, pages 237–246, Austin, Texas, December 1984.
- [DCT87] Mark Drummond, Ken Currie, and Austin Tate. Contingent plan structures for spacecraft. In *Proceedings of the Space Telerobotics Workshop*, January 1987.
- [DD86] Hubert Dreyfus and Stuart Dreyfus. Why expert systems do not exhibit expertise. *IEEE Expert*, pages 86–90, 1986.
- [DD95] Leslie L. Daigle and Peter Deutsch. Agents for internet information clients. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM ’95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [DDL89] Keith S. Decker, Edmund H. Durfee, and Victor R. Lesser. Evaluating research in cooperative distributed problem solving. In L. Gasser and M. N. Huhns, editors, *Distributed Artificial Intelligence, Vol. II*, pages 485–519. Pitman Publishing Ltd., 1989. Also COINS Technical Report 88-89, University of Massachusetts, 1988.
- [DE95] Winton Davies and Peter Edwards. Distributed learning: An agent-based approach to data-mining. In Diana Gordon, editor, *Working Notes of the ICML ’95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [Dea86] Thomas L. Dean. Intractability and time dependent planning. In *Proceedings of the 1986 Workshop on Reasoning About Actions and Plans*, July 1986.
- [Dea87a] Thomas Dean. Large-scale temporal data bases for planning in complex domains. In *Proceedings of the Tenth International Joint Conference on Artificial Intelligence*, pages 860–866, Milan, Italy, August 1987.

- [Dea87b] Thomas Dean. Planning, execution, and control. In *Proceedings of the DARPA Knowledge-based Planning Workshop*, December 1987.
- [Dec87a] Keith S. Decker. DESCANT: A distributed expert system control architecture for non-decomposable tasks. Technical Report 87-CRD-030, GE Corporate Research and Development, April 1987.
- [Dec87b] Keith S. Decker. DESCANT: A distributed expert system control architecture for non-independant tasks. Master's thesis, Rensselaer Polytechnic Institute, May 1987. Also GE Corporate R&D Tech Report 87-CRD-030.
- [Dec87c] Keith S. Decker. *DESCANT User's Guide*. Artificial Intelligence Branch, GE Corporate Research and Development, version 1.0 edition, February 1987. GE Class 2 Proprietary.
- [Dec87d] Keith S. Decker. Distributed problem solving: A survey. *IEEE Transactions on Systems, Man, and Cybernetics*, 17(5):729–740, September 1987.
- [Dec91] Keith S. Decker. Blackboard systems. *IEEE Expert*, 6(5):71–72, October 1991.
- [Dec95a] Keith S. Decker. Distributed artificial intelligence testbeds. In G. O'Hare and N. Jennings, editors, *Foundations of Distributed Artificial Intelligence*, chapter 3. Wiley Inter-Science, 1995. Forthcoming.
- [Dec95b] Keith S. Decker. *Environment Centered Analysis and Design of Coordination Mechanisms*. PhD thesis, University of Massachusetts, 1995.
- [Dec95c] Keith S. Decker. Tæ MS: A framework for analysis and design of coordination mechanisms. In G. O'Hare and N. Jennings, editors, *Foundations of Distributed Artificial Intelligence*, chapter 16. Wiley Inter-Science, 1995. Forthcoming.
- [DeG75] Morris H. DeGroot. *Probability and Statistics*. Addison-Wesley, Reading, MA, 1975.
- [Den82] Eric Denardo. *Dynamic Programming : Models and Applications*. Prentice-Hall, Englewood Cliffs, N.J., 1982.
- [Den86a] D. Dennett. Intentional systems. In D. Dennett, editor, *Brainstorms*. MIT Press, 1986.
- [Den86b] Peter J. Denning. Towards a science of expert systems. *IEEE Expert*, 1986.
- [Den95] Jörg Denzinger. Knowledge-based distributed search using teamwork. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 81–88, San Francisco, CA, 1995. MIT Press.
- [DFF<sup>+</sup>87] Bruce D'Ambrosio, Michael R. Fehling, Stephanie Forrest, Peter Raulefs, and B. Michael Wilber. Real-time process management for materials composition in chemical manufacturing. *IEEE Expert*, pages 80–93, 1987.
- [DG82] V. J. Derlega and J. Grzelak, editors. *Cooperation and Helping Behavior: Theories and Research*. Academic Press, New York, 1982.
- [DG87] G. DeSanctis and R. B. Gallupe. A foundation for the study of group decision support systems. *Management Science*, 33:589–609, 1987.
- [DGHL90] Keith S. Decker, Alan J. Garvey, Marty A. Humphrey, and Victor R. Lesser. Effects of parallelism on blackboard system scheduling. In *Proceedings of the Fourth Annual AAAI Workshop on Blackboard Systems*, Boston, August 1990. Also COINS TR-90-54.
- [DGHL91] Keith S. Decker, Alan J. Garvey, Marty A. Humphrey, and Victor R. Lesser. Effects of parallelism on blackboard system scheduling. In *Proceedings of the Twelfth International Joint Conference on Artificial Intelligence*, pages 15–21, Sydney, Australia, August 1991. Extended version to appear in the *International Journal of Pattern Recognition and Artificial Intelligence* 7(2) 1993.

- [DGH92] Keith Decker, Alan Garvey, Marty Humphrey, and Victor Lesser. A blackboard system for real-time control of approximate processing. In *Proceedings of the 25th Hawaii International Conference on System Sciences*, January 1992. Extended version to appear in the *International Journal of Pattern Recognition and Artificial Intelligence* 7(2) 1993.
- [DGH93a] Keith S. Decker, Alan J. Garvey, Marty A. Humphrey, and Victor R. Lesser. Control heuristics for scheduling in a parallel blackboard system. *International Journal of Pattern Recognition and Artificial Intelligence*, 7(2):243–264, 1993.
- [DGH93b] Keith S. Decker, Alan J. Garvey, Marty A. Humphrey, and Victor R. Lesser. A real-time control architecture for an approximate processing blackboard system. *International Journal of Pattern Recognition and Artificial Intelligence*, 7(2):265–284, 1993.
- [DGL92a] Keith S. Decker, Alan J. Garvey, Victor R. Lesser, and Marty A. Humphrey. An approach to modeling environment and task characteristics for coordination. In Charles J. Petrie, Jr., editor, *Enterprise Integration Modeling: Proceedings of the First International Conference*, pages 379–388. MIT Press, 1992.
- [DGL92b] Keith S. Decker, Alan J. Garvey, Victor R. Lesser, and Marty A. Humphrey. Modeling environment and task characteristics for coordination and scheduling. Working Paper, Cooperative Distributed Problem Solving Laboratory, 1992.
- [DGM95] Thouraya Daouas, Khaled Ghedira, and Jean Pierre Muller. Distributed flow shop scheduling problem: Global versus local optimization. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 443, San Francisco, CA, 1995. MIT Press. (poster).
- [DH89] Keith S. Decker and Marty A. Humphrey. The new DMVT simulator. Working Paper, Cooperative Distributed Problem Solving Laboratory, 1989.
- [DHL89] Keith S. Decker, Marty A. Humphrey, and Victor R. Lesser. Experimenting with control in the DVMT. In *Proceedings of the Third Annual AAAI Workshop on Blackboard Systems*, Detroit, August 1989. Also COINS TR-89-85.
- [DHW90] Pi-Sheng Deng, Clyde W. Holsapple, and Andrew B. Whinston. A skill refinement learning model for rule-based expert systems. *IEEE Expert*, 5(2):15–28, April 1990.
- [DKR95] Barbara Dunin-Keplicz and Anna Radzikowska. Actions with typical effects: Epistemic characterization of scenarios. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 445, San Francisco, CA, 1995. MIT Press. (poster).
- [dKW87] Johan de Kleer and Brian C. Williams. Diagnosing multiple faults. *Artificial Intelligence*, 32:97–130, 1987.
- [DL86] Edmund H. Durfee and Victor R. Lesser. Incremental planning to control a blackboard-based problem solver. In *Proceedings of the National Conference on Artificial Intelligence*, pages 58–64, Philadelphia, PA, August 1986.
- [DL87a] Edmund H. Durfee and Victor R. Lesser. Planning coordinated actions in dynamic domains. In *Proceedings of the DARPA Knowledge-Based Planning Workshop*, pages 18.1–18.10, December 1987. Also COINS-TR-87-130.
- [DL87b] Edmund H. Durfee and Victor R. Lesser. Planning to meet deadlines in a blackboard-based problem solver. In J. Stankovic and K. Ramamritham, editors, *Tutorial on Hard Real-Time Systems*, pages 595–608. IEEE Computer Society Press, 1987.
- [DL87c] Edmund H. Durfee and Victor R. Lesser. Using partial global plans to coordinate distributed problem solvers. In *Proceedings of the Tenth International Joint Conference on Artificial Intelligence*, pages 875–883, Milan, Italy, August 1987. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 285–293, Morgan Kaufmann, 1988.).

- [DL88a] E. Durfee and V. Lesser. Predictability vs. responsiveness: Coordinating problem solvers in dynamic domains. In *Proceedings of the Seventh National Conference on Artificial Intelligence*, pages 66–71, St. Paul, MN, August 1988.
- [DL88b] Edmund H. Durfee and Victor R. Lesser. Incremental planning to control a time-constrained, blackboard-based problem solver. *IEEE Transactions on Aerospace and Electronic Systems*, 24(5):647–662, September 1988.
- [DL88c] Edmund H. Durfee and Victor R. Lesser. Negotiation through partial global planning. In *Proceedings of the 1988 Distributed AI Workshop*, May 1988.
- [DL88d] Edmund H. Durfee and Victor R. Lesser. Predictability versus responsiveness: Coordinating problem solvers in dynamic domains. In *Proceedings of the National Conference on Artificial Intelligence*, pages 66–71, August 1988.
- [DL89a] Keith S. Decker and Victor R. Lesser. Some initial thoughts on a generic architecture for CDPS network control. In *Proceedings of the Ninth Workshop on Distributed AI*, September 1989. Also in COINS TR-90-81.
- [DL89b] Edmund H. Durfee and Victor R. Lesser. Negotiating task decomposition and allocation using partial global planning. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*, pages 229–243. Pitman, 1989.
- [DL90a] Keith S. Decker and Victor R. Lesser. Extending the partial global planning framework for cooperative distributed problem solving network control. In *Proceedings of the Workshop on Innovative Approaches to Planning, Scheduling and Control*, pages 396–408, San Diego, November 1990. Morgan Kaufmann. Also COINS TR-90-81.
- [DL90b] Keith S. Decker and Victor R. Lesser. A scenario for cooperative distributed problem solving. In *Proceedings of the Tenth International Workshop on Distributed AI*, Texas, October 1990.
- [DL91] Edmund H. Durfee and Victor R. Lesser. Partial global planning: A coordination framework for distributed hypothesis formation. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(5), September 1991. (Special Issue on Distributed Sensor Networks).
- [DL92a] Keith Decker and Victor Lesser. The analysis of quantitative coordination relationships. In *Working Papers of the 11th International Workshop on Distributed Artificial Intelligence*, pages 59–75, February 1992.
- [DL92b] Keith S. Decker and Victor R. Lesser. Generalizing the partial global planning algorithm. *International Journal of Intelligent and Cooperative Information Systems*, 1(2):319–346, June 1992.
- [DL93a] Keith S. Decker and Victor R. Lesser. Analyzing a quantitative coordination relationship. *Group Decision and Negotiation*, 2(3):195–217, 1993.
- [DL93b] Keith S. Decker and Victor R. Lesser. An approach to analyzing the need for meta-level communication. In *Proceedings of the Thirteenth International Joint Conference on Artificial Intelligence*, pages 360–366, Chambéry, France, August 1993.
- [DL93c] Keith S. Decker and Victor R. Lesser. A one-shot dynamic coordination algorithm for distributed sensor networks. In *Proceedings of the Eleventh National Conference on Artificial Intelligence*, pages 210–216, Washington, July 1993.
- [DL93d] Keith S. Decker and Victor R. Lesser. Quantitative modeling of complex computational task environments. In *Proceedings of the Eleventh National Conference on Artificial Intelligence*, pages 217–224, Washington, July 1993.

- [DL93e] Keith S. Decker and Victor R. Lesser. Quantitative modeling of complex environments. *International Journal of Intelligent Systems in Accounting, Finance, and Management*, 2(4):215–234, December 1993. Special issue on “Mathematical and Computational Models of Organizations: Models and Characteristics of Agent Behavior”.
- [DL94a] Keith S. Decker and Victor R. Lesser. Designing a family of coordination algorithms. In *Proceedings of the Thirteenth International Workshop on Distributed AI*, pages 65–84, Seattle, WA, July 1994. AAAI Press Technical Report WS-94-02. Also UMass CS-TR-94-14. To appear, Proceedings of the First International Conference on Multi-Agent Systems, San Francisco, AAAI Press, 1995.
- [DL94b] Keith S. Decker and Victor R. Lesser. Experimenting with a family of coordination algorithms. Submitted to IJCAI-95, 1994.
- [DL94c] Keith S. Decker and Victor R. Lesser. Task environment centered design of organizations. In Ingemar Hulthage, editor, *Computational Organization Design*. AAAI Spring Symposium, 1994. Working Notes.
- [DL95a] Keith Decker and Victor Lesser. Designing a family of coordination algorithms. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 73–80, San Francisco, CA, 1995. MIT Press. Longer version available as UMass CS-TR 94–14.
- [DL95b] Keith S. Decker and Victor R. Lesser. Coordination assistance for mixed human and computational agent systems. CS Technical Report 95–31, UMass, 1995.
- [dL96] Mark d’Inverno and Michael Luck. Formalising the contract net as a goal-directed system. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [DLC85a] Edmund H. Durfee, Victor R. Lesser, and Daniel D. Corkill. Increasing coherence in a distributed problem solving network. In *Proceedings of the Ninth International Joint Conference on Artificial Intelligence*, pages 1025–1030, Los Angeles, CA, August 1985.
- [DLC85b] Edmund H. Durfee, Victor R. Lesser, and Daniel D. Corkill. Making coherent communication decisions in a distributed problem solving network. In *Proceedings of the Workshop on AI and Distributed Problem Solving*, pages 59–75. National Academy of Sciences, May 1985.
- [DLC87a] Edmund H. Durfee, Victor R. Lesser, and Daniel D. Corkill. Coherent cooperation among communicating problem solvers. *IEEE Transactions on Computers*, C-36(11):1275–1291, November 1987. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 268–284, Morgan Kaufmann, 1988.).
- [DLC87b] Edmund H. Durfee, Victor R. Lesser, and Daniel D. Corkill. Cooperation through communication in a distributed problem solving network. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, Research Notes in Artificial Intelligence, chapter 2, pages 29–58. Pitman, 1987. (Also in S. Robertson, W. Zachary, and J. Black (eds.), *Cognition, Computing, and Cooperation*, Ablex 1990.).
- [DLC89a] Edmund H. Durfee, Victor R. Lesser, and Daniel D. Corkill. Cooperative distributed problem solving. In Avron Barr, Paul R. Cohen, and Edward A. Feigenbaum, editors, *The Handbook of Artificial Intelligence*, volume IV, chapter XVII, pages 83–137. Addison-Wesley, 1989.
- [DLC89b] Edmund H. Durfee, Victor R. Lesser, and Daniel D. Corkill. Trends in cooperative distributed problem solving. *IEEE Transactions on Knowledge and Data Engineering*, 1(1):63–83, March 1989.
- [DLC91] Edmund H. Durfee, Victor R. Lesser, and Daniel D. Corkill. Distributed problem solving. In S. Shapiro, editor, *The Encyclopedia of Artificial Intelligence, Second Edition*. John Wiley & Sons, 1991.

- [DLG93] E. H. Durfee, J. Lee, and P. J. Gmytrasiewicz. Overeager reciprocal rationality and mixed strategy equilibria. In *Proceedings of the Eleventh National Conference on Artificial Intelligence*, pages 225–230, Washington D.C., 1993.
- [DLPW95] Keith Decker, Victor Lesser, M. V. Nagendra Prasad, and Thomas Wagner. MACRON: An architecture for multi-agent cooperative information gathering. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [DLW90] Keith S. Decker, Victor R. Lesser, and Robert C. Whitehair. Extending a blackboard architecture for approximate processing. *The Journal of Real-Time Systems*, 2(1/2):47–79, 1990.
- [DM89] E. Durfee and T. Montgomery. MICE: A flexible testbed for intelligent coordination experiments. In *Proceedings of the Ninth Workshop on Distributed AI*, pages 25–40, Rosario, Washington, September 1989.
- [DM90] Edmund H. Durfee and Thomas A. Montgomery. A hierarchical protocol for coordinating multiagent behaviors. In *Proceedings of the Eighth National Conference on Artificial Intelligence*, pages 86–93, July 1990.
- [DM91] Edmund H. Durfee and Thomas A. Montgomery. Coordination as distributed search in a hierarchical behavior space. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6):1363–1378, November/December 1991.
- [DMGLS82] Renato De-Mori, Attilio Giordana, Pietro Laface, and Lorenza Saitta. An expert systems for interpreting speech patterns. In *AAAI*, pages 107–110, 1982.
- [Dor93] J. Doran. Using DAI software testbeds. In *Proceedings of the CKBS SIG Workshop 1993*, University of Keele, Keele, UK, 1993.
- [Doy79] Jon Doyle. A truth maintenance system. *Artificial Intelligence*, 12:231–272, 1979.
- [Doy83] Jon Doyle. A society of mind. Technical Report CMU-CS-83-127, CMU Computer Science Department, Pittsburgh, PA, 1983.
- [DP84] D. Dubois and H. Prade. Criteria aggregation and ranking of alternatives in the framework of fuzzy set theory. In H. J. Zimmermen, L. A. Zadeh, and B. R. Gains, editors, *TIMS/Studies in Management Science*, volume 20, pages 209–240. Elsevier Science Publishers, 1984.
- [DPGM94] J. Doran, M. Palmer, N. Gilbert, and P. Mellars. The EOS project: Modelling upper palaeolithic social change. In N. Gilbert and J. Doran, editors, *Simulating Societies*. UCL Press, 1994. Forthcoming.
- [DR89] C. DeKoven and T. Radhakrishnan. Cooperative game playing in a computer-supported environment. Computer science technical report, Concordia University, Montreal, Quebec, 1989.
- [Dre81] Melvin Dresher. *The Mathematics of Games of Strategy: Theory and Applications*. Dover, 1981.
- [DS81] Randall Davis and Reid G. Smith. Negotiation as a metaphor for distributed problem solving. AI Memo 624, Artificial Intelligence Laboratory, Massachusetts Institute of Technology, Cambridge, MA 02139, May 1981. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 333–356, Morgan Kaufmann, 1988.).
- [DS83a] Randall Davis and Howard Shrobe. Representing structure and behavior of digital hardware. *IEEE Computer*, pages 75–81, October 1983.
- [DS83b] Randall Davis and Reid Smith. Negotiation as a metaphor for distributed problem solving. *Artificial Intelligence*, 20(1):63–109, January 1983.

- [Dur84] Edmund Howell Durfee. A parallel simulation of a distributed problem solving network. Master's thesis, University of Massachusetts, September 1984. (Also published as Technical Report 84-19, Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts 01003, September 1984.).
- [Dur86] Edmund H. Durfee. An approach to cooperation: Planning and communication in a distributed problem solving network. Technical Report 86-09, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, March 1986.
- [Dur87] Edmund Howell Durfee. *A Unified Approach to Dynamic Coordination: Planning actions and interactions in a distributed problem solving network*. PhD thesis, University of Massachusetts, September 1987. (Also published as Technical Report 87-84, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, September 1987.).
- [Dur88] Edmund H. Durfee. *Coordination of Distributed Problem Solvers*. Kluwer Academic Publishers, 1988.
- [Dur90a] Edmund H. Durfee. A cooperative approach to planning for real-time control. In *Proceedings of the 1990 DARPA Workshop on Innovative Approaches to Planning, Scheduling, and Control*, pages 277–283, November 1990.
- [Dur90b] Edmund H. Durfee. Towards intelligent real-time cooperative systems. In *Notes of the 1990 AAAI Spring Symposium on Planning in Uncertain, Unpredictable, or Changing Environments*, pages 29–33, March 1990. (Available in University of Maryland Systems Research Center technical report *Planning in Uncertain, Unpredictable, or Changing Environments* (SRC-TR-90-45), James Hendler (ed.)).
- [Dur91a] Edmund H. Durfee. The distributed artificial intelligence melting pot. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6), December 1991. (Special Issue on Distributed AI).
- [Dur91b] Edmund H. Durfee. Organizations, plans, and schedules: An interdisciplinary perspective on coordinating AI agents. *Journal of Intelligent Systems*, 1991. (Special Issue on the Social Context of Intelligent Systems).
- [Dur95] Edmund H. Durfee. Blissful ignorance: Knowing just enough to coordinate well. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 406–413, San Francisco, CA, 1995. MIT Press. (invited speaker talk).
- [DW90] Jon Doyle and Michael P. Wellman. Rational distributed reason maintenance for planning and replanning of large-scale activities (preliminary report). In *Proceedings of the Workshop on Innovative Approaches to Planning, Scheduling and Control*, pages 28–36, San Diego, November 1990. Morgan Kaufmann.
- [DW91] T. L. Dean and M. P. Wellman. *Planning and Control*. Morgan Kaufman Publishers, San Mateo, California, 1991.
- [DW96] Eric Dubois and Suchun Wu. An agent-oriented framework for security requirements. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [Dyb87] R. Kent Dybvig. *The Scheme Programming Language*. Prentice Hall, 1987.
- [EA89] M. Evans and J. Anderson. A constraint-directed architecture for multi-agent planning. In *Proceedings of the 9th Workshop on Distributed Artificial Intelligence*, Seattle, Washington, 1989.
- [ea90] Natraj Arni et al. Overview of RAD: A hybrid and distributed reasoning tool. Technical Report ACT-RA-098-90, Microelectronics and Computer Technology Corporation, Austin, TX, March 1990.

- [EBD95] Love Ekenberg, Magnus Boman, and Mats Danielson. A tool for coordinating autonomous agents with conflicting goals. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 89–93, San Francisco, CA, 1995. MIT Press.
- [(ed84] David Maines (editor). Urban life, 1984. Special issue on negotiated order theory.
- [EG85] J. Robert Ensor and John D. Gabbe. Transactional blackboards. In *Proceedings of the Ninth International Joint Conference on Artificial Intelligence*, pages 340–344, Los Angeles, CA, August 1985. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 557–562, Morgan Kaufmann, 1988.).
- [EHRLR80] Lee D. Erman, Frederick Hayes-Roth, Victor R. Lesser, and D. Raj Reddy. The Hearsay-II speech-understanding system: Integrating knowledge to resolve uncertainty. *Computing Surveys*, 12(2):213–253, June 1980.
- [EK90] Kaan Egilmez and Steven H. Kim. Coordination of multiple agents through modal state logic and influence specifications. Technical report, MIT, Laboratory for Manufacturing and Productivity, 1990.
- [EL75] Lee D. Erman and Victor R. Lesser. A multi-level organization for problem solving using many diverse cooperating sources of knowledge. In *Proceedings of the Fourth International Joint Conference on Artificial Intelligence*, pages 483–490, Stanford, CA, August 1975.
- [ELF81] Lee D. Erman, Philip E. London, and Stephen F. Fickas. The design and an example use of Hearsay-III. In *Proceedings of the Seventh International Joint Conference on Artificial Intelligence*, pages 409–415, Tokyo, Japan, August 1981.
- [EM88] Robert Englemore and Tony Morgan. *Blackboard Systems*. Addison Wesley, Reading Mass., 1988.
- [Ema91] *DAI Group Discussion by EMAIL*, 1991.
- [EN69] George W. Ernst and Allen Newell. *GPS: A Case Study in Generality and Problem Solving*. Academic Press, New York, 1969.
- [EN77] Robert S. Englemore and H. Nii. A knowledge-based system for the interpretation of protein X-ray crystallographic data. Technical Report CS-77-589, Computer Science Department, Stanford University, Stanford, CA 94305, February 1977.
- [EPR95] Eithan Ephrati, Martha Pollack, and Jeffrey S. Rosenschein. A tractable heuristic that maximizes global utility through local plan combination. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 94–101, San Francisco, CA, 1995. MIT Press.
- [EPU95] Eithan Ephrati, Martha E. Pollack, and Sigalit Ur. Deriving multi-agent coordination through filtering strategies. In Chris S. Mellish, editor, *Proceedings of the Fourteenth International Joint Conference on Artificial Intelligence*, pages 679–685, Montréal, Canada, August 1995.
- [ER91] Eithan Ephrati and Jeffrey S. Rosenschein. The clarke tax as a consensus mechanism among automated agents. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 173–178, Anaheim, CA, July 1991.
- [ER92] Eithan Ephrati and Jeffrey S. Rosenschein. Constrained intelligent action: Planning under the influence of a master agent. In *Proceedings of the National Conference on Artificial Intelligence*, pages 263–268, San Jose, CA, July 1992.
- [ER94] Eithan Ephrati and Jeffrey S. Rosenschein. Divide and conquer in multi-agent planning. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 375–380, Seattle, WA, 1994.



- [Eri82] Lars Warren Ericson. DPL-82: A language for distributed processing. In *Proceedings of the 3rd International Conference on Distributed Computing Systems*, pages 526–531, Miami/Ft. Lauderdale, Florida, October 1982.
- [Eve79] Shimon Even. *Graph Algorithms*. Computer Science Press, Potomac, Maryland, 1979.
- [EW81] Manfred Eigen and Ruthild Winkler. *Laws of the Game*. Harper and Row, 1981.
- [FAS82] M. S. Fox, B. Allen, and G. Strohm. Job-shop scheduling: An investigation in constraint-directed reasoning. In *AAAI-82*, Pittsburgh, PA, 1982. AAAI.
- [FB82] J. A. Feldman and D. H. Ballard. Connectionist models and their properties. *Cognitive Science*, 6(3):205–254, July–September 1982.
- [FC90] Jaques Ferber and Patrice Carle. Actors and agents as reflective concurrent objects: A MERING IV perspective. In *Proceedings of the Tenth International Workshop on Distributed AI*, Bandera, Texas, October 1990.
- [FC91] Jacques Ferber and Patrice Carle. Actors and agents as reflective concurrent objects: A MERING IV perspective. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6), December 1991. (Special Issue on Distributed AI).
- [FE83] Michael Fehling and Lee Erman. Report on the third annual workshop on Distributed Artificial Intelligence. *SIGART Newsletter*, (84):3–12, April 1983.
- [Fei77] Edward A. Feigenbaum. The art of Artificial Intelligence I: Themes and cases studies of knowledge engineering. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, pages 1014–1029, Tbilisi, Georgia, USSR, August 1977.
- [Fel79] Jerome A. Feldman. High level programming for distributed computing. *Communications of the ACM*, 22(6):353–368, June 1979.
- [Fel87] M. Feldman. Electronic mail and weak ties in organizations. *Office Technology and People*, 3:83–101, 1987.
- [Fen75] Richard Dean Fennell. *Multiprocess Software Architecture for AI Problem Solving*. PhD thesis, Carnegie-Mellon University, May 1975.
- [Fer96] Jacques Ferber. Cooperation strategies in collective intelligence. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [FF84] Robert E. Filman and Daniel P. Friedman. *Coordinated Computing: Tools and Techniques for Distributed Software*. McGraw-Hill, New York, 1984.
- [FFMM94] Tim Finin, Rich Fritzson, Don McKay, and Robin McEntire. KQML - A language and protocol for knowledge and information exchange. Technical Report CS-94-02, Computer Science Department, University of Maryland and Valley Forge Engineering Center, Unisys Corporation, Computer Science Department, University of Maryland, UMBC Baltimore MD 21228, 1994.
- [FH87] Robert F. Franklin and Laurel A. Harmon. Elements of cooperative behavior. Technical report, Environmental Research Institute of Michigan, Ann Arbor, MI 48107, August 1987.
- [FHN72] Richard E. Fikes, Peter E. Hart, and Nils J. Nilsson. Learning and executing generalized robot plans. *Artificial Intelligence*, 3, 1972.
- [FHS83] Scott E. Fahlman, Geoffrey E. Hinton, and Terrence J. Sejnowski. Massively parallel architectures for AI: NETL, THISTLE and Boltzmann Machines. In *Proceedings of the Eighth International Joint Conference on Artificial Intelligence*, pages 109–113, Karlsruhe, Federal Republic of Germany, August 1983.

- [FI83] Peter E. Friedland and Yumi Iwasaki. The concept and implementation of skeletal plans. Draft technical report, Computer Science Department, Stanford University, Stanford, CA 94305, April 1983.
- [Fik82] R. E. Fikes. A commitment-based framework for describing informal cooperative work. *Cognitive Science*, 6:331–347, 1982.
- [Fir87] R. James Firby. An investigation into reactive planning in complex domains. In *Proceedings of the National Conference on Artificial Intelligence*, pages 202–206, Seattle, WA, August 1987.
- [Fis95] K. Fisher. MAGSY. In G. O’Hare and N. Jennings, editors, *Foundations of Distributed Artificial Intelligence*, chapter 15. Wiley Inter-Science, 1995. Forthcoming.
- [FK96] Innes A. Ferguson and Grigoris J. Karakoulas. Multiagent learning and adaptation in an information filtering market. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [FKR95] Maier Fenster, Sarit Kraus, and Jeffrey S. Rosenschein. Coordination without communication: Experimental validation of focal point techniques. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 102–108, San Francisco, CA, 1995. MIT Press.
- [FL77a] R. D Fennell and V. R. Lesser. Parallelism in AI problem solving: A case study of Hearsay-II. *IEEE Transactions on Computers*, C-26(2):98–111, February 1977.
- [FL77b] Richard D. Fennell and Victor R. Lesser. Parallelism in Artificial Intelligence problem solving: A case study of Hearsay II. *IEEE Transactions on Computers*, C-26(2):98–111, February 1977. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 106–119, Morgan Kaufmann, 1988.).
- [FL86] Nicholas V. Findler and Ron Lo. An examination of distributed planning in the world of air traffic control. *Journal of Parallel and Distributed Computing*, 3:411–431, 1986.
- [FM96] Klaus Fischer and Jvrg Muller. A decision-theoretic model for cooperative transportation scheduling. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [FMB82] M. Flowers, R. McGuire, and L. Birnbaum. Adversary arguments and the logic of personal attacks. In W. G. Lehnert and M. H. Ringle, editors, *Strategies for Natural Language Processing*. Lawrence Erlbaum Associates, Hillsdale, N.J., 1982.
- [FMP95] Klaus Fischer, Jörg P. Müller, and Markus Pischel. A model for cooperative transportation scheduling. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 109–116, San Francisco, CA, 1995. MIT Press.
- [FMPS95] Klaus Fischer, Jörg P. Müller, Markus Pischel, and Darius Schier. Unifying control in A layered agent architecture. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 446, San Francisco, CA, 1995. MIT Press. (poster).
- [FN71] R. E. Fikes and N. J. Nilsson. STRIPS: A new approach to the application of theorem proving to problem solving. *Artificial Intelligence*, 2:189–208, 1971.
- [For81] C. L. Forgy. OPS5 reference manual. Technical Report CMU-CS-81-135, Computer Science Department, Carnegie-Mellon University, Pittsburgh, PA, 1981.
- [For82] C. L. Forgy. RETE: A fast algorithm for the many pattern/many object pattern matching problem. *Artificial Intelligence*, 19:17–37, 1982.

- [Fox79] Mark S. Fox. Organization structuring: Designing large complex software. Technical Report 79-155, Computer Science Department, Carnegie-Mellon University, Pittsburgh, PA, December 1979.
- [Fox81] Mark S. Fox. An organizational view of distributed systems. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):70–80, January 1981. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 140–150, Morgan Kaufmann, 1988.).
- [Fox82] Mark S. Fox. The intelligent management system: An overview. Technical Report CMU-RI-81-4, Carnegie-Mellon Intelligent Systems Laboratory, Pittsburgh, PA, December 1982.
- [Fox83] Mark S. Fox. *Constraint-directed Search: A Case Study of Job-Shop Scheduling*. PhD thesis, Carnegie-Mellon University, 1983. (Also published as Technical Report CMU-CS-83-161, Computer Science Department, Carnegie-Mellon University, Pittsburgh, Pennsylvania 15213.).
- [Fox87] Mark S. Fox. *Constraint-directed Search: A Case Study of Job-Shop Scheduling*. Morgan Kaufmann Publishers, 1987.
- [Fox90] Mark S. Fox. AI and expert system myths, legends, and facts. *IEEE Expert*, 5(1):8–20, February 1990.
- [FPT<sup>+</sup>95] Tim Finin, Anupama Potluri, Chelliah Thirunavukkarasu, Don McKay, and Robin McEntire. On agent domains, agent names and proxy agents. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [FPWT91] F. Farhoodi, J. Proffitt, P. Woodman, and A. Tunnicliffe. Design of organizations in distributed decision systems. In *Proceedings of the AAAI Workshop on Cooperation Among Heterogeneous Intelligent Systems*, Anaheim, California, July 1991.
- [Fre82] S. French. *Sequencing and Scheduling: An Introduction to the Mathematics of the Job-Shop*. Ellis Horwood, Chichester, 1982.
- [Fre86] Simon French. *Decision Theory: An Introduction to the Mathematics of Rationality*. Ellis Horwood, Chichester, England, 1986.
- [Fre90] Yehudah Freundlich. Transfer pricing : Integrating expert systems in MIS environments. *IEEE Expert*, 5(1):54–62, February 1990.
- [Fri79] Peter E. Friedland. *Knowledge-based Experiment Design in Molecular Genetics*. PhD thesis, Stanford University, 1979. (Also published as Technical Report 79-771, Computer Science Department, Stanford University, Stanford California 94305, 1979.).
- [FS77] Jerome A. Feldman and Robert F. Sproull. Decision theory and Artificial Intelligence II: The hungry monkey. *Cognitive Science*, 1(2):158–192, April 1977.
- [FS84] Mark S. Fox and Stephen F. Smith. ISIS—A knowledge-based system for factory scheduling. *Expert Systems*, 1(1):25–49, July 1984.
- [FS88] M. S. Fox and N. Sadeh. Problem space textures and topologies: Viewing heuristic search as constraint satisfaction. Robotics Institute, Carnegie Mellon University, December 1988.
- [FS91] Nicholas Findler and U. Sengupta. An overview of some recent and current research in the AI lab at arizona state university. *AI Magazine*, 12(3), 1991.
- [FSB89] Mark S. Fox, Norman Sadeh, and Can Baykan. Constrained heuristic search. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, pages 309–315, Detroit, MI, August 1989.

- [FSH96] Amal El Fallah-Seghrouchni and Serge Haddad. A coordination algorithm for multi-agent planning. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [FSO86] M. S. Fox, S. F. Smith, and P. S. Ow. Generating and maintaining detailed production plans: Investigations into the development of knowledge-based factory scheduling systems. *AI Magazine*, 7(4), Fall 1986.
- [FT84] S. Fiske and S. E. Taylor. *Social Cognition*. Addison Wesley, New York, N.Y., 1984.
- [FW77] Daniel P. Friedman and David S. Wise. Functional combination. *Computer Languages*, 3:31–35, 1977.
- [GA95] Andrew Garland and Richard Alterman. Preparation of multi-agent knowledge for reuse. In David W. Aha and Ashwin Ram, editors, *Working Notes for the AAAI Symposium on Adaptation of Knowledge for Reuse*, Cambridge, MA, November 1995. AAAI.
- [GA96] Andrew Garland and Richard Alterman. Multiagent learning through collective memory. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [Gal58] J. Galbraith. *Organizations*. John Wiley and Sons, New York, 1958.
- [Gal73] Jay Galbraith. *Designing Complex Organizations*. Addison-Wesley, 1973.
- [Gal77] Jay R. Galbraith. *Organization Design*. Addison-Wesley, 1977.
- [Gar87] Alan Garvey. Implementing diverse forms of control knowledge in multiple control architectures. Technical Report KSL-87-40, Knowledge Systems Laboratory, Stanford University, 1987.
- [Gar96] Alan J. Garvey. *Design-to-time Real-time Scheduling*. PhD thesis, Department of Computer Science, University of Massachusetts, Amherst, February 1996.
- [Gas85a] Saul Gass. *Linear Programming : Methods and Applications*. McGraw-Hill, 1985.
- [Gas85b] L. Gasser. The 1985 workshop on distributed artificial intelligence. *AI Magazine*, 8(2), 1985.
- [Gas86] L. Gasser. The integration of computing and routine work. *ACM Transactions on Office Information Systems*, 4(3):205–225, July 1986.
- [Gas87] Les Gasser. Report on the 1985 workshop on distributed artificial intelligence. *AI Magazine*, 8(2):91–97, Summer 1987.
- [Gas91a] Les Gasser. Social conceptions of knowledge and action. *Artificial Intelligence*, 47(1):107–138, 1991.
- [Gas91b] Les Gasser. Social conceptions of knowledge and action: DAI foundations and open systems semantics. *Artificial Intelligence*, 47(1-3):107–138, 1991.
- [Gas95] Les Gasser. Computational organization research. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 414–415, San Francisco, CA, 1995. MIT Press. (invited speaker talk).
- [GB96] Dan L. Greco and David C. Brown. Learning to design together. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [GBH87a] L. Gasser, C. Braganza, and N. Herman. Implementing distributed artificial intelligence systems using MACE. In *Proceedings of the 3rd IEEE Conference on Artificial Intelligence Applications*, pages 315–320, 1987. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 445–450, Morgan Kaufmann, 1988.).

- [GBH87b] Les Gasser, Carl Braganza, and Nava Herman. MACE: A flexible testbed for distributed AI research. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, Research Notes in Artificial Intelligence, chapter 5, pages 119–152. Pitman, 1987.
- [GBH88] L. Gasser, C. Braganza, and N. Herman. Implementing distributed AI systems using MACE. In Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 445–450. Morgan Kaufmann, 1988.
- [GC81] Fernando Gomez and B. Chandrasekaran. Knowledge organization and distribution for medical diagnosis. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):34–42, January 1981.
- [GC88] Kevin Q. Gallagher and Daniel D. Corkill. Blackboard retrieval strategies in GBB. Technical Report 88-39, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, December 1988.
- [GC89] Kevin Q. Gallagher and Daniel D. Corkill. Performance aspects of GBB. In V. Jagannathan and Raj T. Dodhiawala, editors, *Current Trends in Blackboard Systems*. Academic Press, 1989. (In press.).
- [GC95] Graça Gaspar and Helder Coelho. Deriving goals and intentions. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 447, San Francisco, CA, 1995. MIT Press. (poster).
- [GCHR87] Alan Garvey, Craig Cornelius, and Barbara Hayes-Roth. Computational costs versus benefits of control reasoning. In *Proceedings of the Sixth National Conference on Artificial Intelligence*, pages 110–115, July 1987.
- [GCJ88] Kevin Q. Gallagher, Daniel D. Corkill, and Philip M. Johnson. *GBB Reference Manual*. Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, GBB Version 1.2 edition, September 1988. (Published as Technical Report 88-66, Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts 01003, September 1988.).
- [GD92] P. J. Gmytrasiewicz and E. H. Durfee. Logic of knowledge and belief for recursive modeling: Preliminary report. In *Proceedings of the National Conference on Artificial Intelligence*, pages 628–634, July 1992.
- [GD95] Piotr J. Gmytrasiewicz and Edmund H. Durfee. A rigorous, operational formalization of recursive modeling. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 125–132, San Francisco, CA, 1995. MIT Press.
- [GD96a] John Grefenstette and Robert Daley. Methods for competitive and cooperative co-evolution. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [GD96b] Zahia Guessoum and M. Dojat. A real-time agent model in an asynchronous-object environment. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [GDW91a] Piotr J. Gmytrasiewicz, Edmund H. Durfee, and David K. Wehe. Combining decision theory and hierarchical planning for a time-dependent robotic application. In *Proceedings of the Seventh IEEE Conference on AI Applications*, pages 282–288, February 1991.
- [GDW91b] Piotr J. Gmytrasiewicz, Edmund H. Durfee, and David K. Wehe. A decision-theoretic approach to coordinating multiagent interactions. In *Proceedings of the Twelfth International Joint Conference on Artificial Intelligence*, pages 62–68, Sydney, Australia, August 1991.

- [GDW91c] Piotr J. Gmytrasiewicz, Edmund H. Durfee, and David K. Wehe. The utility of communication in coordinating intelligent agents. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 166–172, July 1991.
- [Gen83] Michael R. Genesereth. An overview of meta-level architecture. In *Proceedings of the National Conference on Artificial Intelligence*, pages 119–124, Washington, D.C., August 1983.
- [Geo83a] Michael Georgeff. Communication and interaction in multi-agent planning. In *Proceedings of the National Conference on Artificial Intelligence*, pages 125–129, Washington, D.C., August 1983. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 200–204, Morgan Kaufmann, 1988.).
- [Geo83b] Michael Georgeff. Communication and interaction in multiagent planning. In *Proceedings of the Eighth International Joint Conference on Artificial Intelligence*, pages 125–129, August 1983.
- [Geo84] Michael Georgeff. A theory of action for multiagent planning. In *Proceedings of the Fourth National Conference on Artificial Intelligence*, pages 121–125, Austin, TX, August 1984. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 205–209, Morgan Kaufmann, 1988.).
- [Geo86] Michael Georgeff. A representation of events in multi-agent domains. In *Proceedings of the National Conference on Artificial Intelligence*, pages 70–75, Philadelphia, PA, August 1986. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 210–215, Morgan Kaufmann, 1988.).
- [Geo95] Michael Georgeff. Agents and their plans. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [Ger76] Elihu M. Gerson. On ‘quality of life’. *American Sociological Review*, 41:793–806, 1976.
- [GF<sup>+</sup>92] M. R. Genesereth, R. E. Fikes, et al. Knowledge interchange format, version 3.0 reference manual. Technical Report Logic-92-1, Computer Science Department, Stanford University, 1992.
- [GG96] Fausto Giunchiglia and Enrico Giunchiglia. Ideal and real belief about belief. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [GGKS95] Donald F. Geddis, Michael R. Genesereth, Arthur M. Keller, and Narinder P. Singh. Infomaster: A virtual information system. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM ’95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [GGP92] C. Grossner, P. Gokulchander, and A. Preece. On the structure of rule based expert systems. DAI Technical Report DAI-0592-0013, Concordia University, Montreal Quebec, May 1992.
- [GGPR93] C. Grossner, P. Gokulchander, A. Preece, and T. Radhakrishnan. Data distribution in organizations of cooperating expert systems. In *12th International Workshop on Distributed Artificial Intelligence*, May 1993.
- [GGR84] M. R. Genesereth, M. L. Ginsberg, and J. S. Rosenschein. Cooperation without communications. Technical Report 84-36, Stanford Heuristic Programming Project, Computer Science Department, Stanford University, Stanford, CA 94305, 1984. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 220–226, Morgan Kaufmann, 1988.).

- [GGS95] Enrico Giunchiglia, Fausto Giunchiglia, and Luciano Serafini. Agents as reasoners, observers, or arbitrary believers. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 448, San Francisco, CA, 1995. MIT Press. (poster).
- [GH86] A. L. Glass and K. J. Holyoak. *Cognition*. Random House, New York, 1986.
- [GH89] Les Gasser and Michael N. Huhns, editors. *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [GH95] Natalie S. Glance and Tad Hogg. Dilemmas in computational societies. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 117–124, San Francisco, CA, 1995. MIT Press.
- [GHJ<sup>+</sup>86] Alan Garvey, Michael Hewett, M. Vaughan Johnson, Robert Schulman, and Barbara Hayes-Roth. *BB1 User Manual*. Knowledge Systems Laboratory, Departments of Medical and Computer Science, Stanford, CA 94305, Common Lisp edition, October 1986. (Published as Working Paper KSL 86-61, Knowledge Systems Laboratory, Departments of Medical and Computer Science, Stanford University, Stanford, California 94305.).
- [GHL93] Alan Garvey, Marty Humphrey, and Victor Lesser. Task interdependencies in design-to-time real-time scheduling. In *Proceedings of the Eleventh National Conference on Artificial Intelligence*, pages 580–585, Washington, July 1993.
- [GHR89] Alan Garvey and Barbara Hayes-Roth. An empirical analysis of explicit vs. implicit control architectures. In V. Jagannathan, Rajendra Dodhiawala, and Lawrence S. Baum, editors, *Blackboard Architectures and Applications*, pages 43–56. Academic Press, Inc., 1989.
- [GI91] Les Gasser and Toru Ishida. A dynamic organizational architecture for adaptive problem solving. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 185–190, Anaheim, July 1991.
- [GI95] Cheng Gu and Toru Ishida. A quantitative analysis of the contract net protocol. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 449, San Francisco, CA, 1995. MIT Press. (poster).
- [GI96] Cheng Gu and Toru Ishida. Analyzing the social behavior of contract net protocol. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [Gil95] Yolanda Gil. Acquiring knowledge from users in a reflective architecture. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [Gin87] Matthew L. Ginsberg. Decision procedures. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, volume 1 of *Research Notes in Artificial Intelligence*. Pitman, 1987.
- [GJ79] Michael R. Garey and David S. Johnson. *Computers and Intractability: A Guide to the Theory of NP-Completeness*. W.H. Freeman and Co., San Francisco, CA, 1979.
- [GK93] B. Grosz and S. Kraus. Collaborative plans for group activities. In *Proceedings of the Thirteenth International Joint Conference on Artificial Intelligence*, Chambéry, France, August 1993.
- [GKE90] Jolene Galegher, Robert E. Kraut, and Carmen Egidio. *Intellectual Teamwork: Social and Technological Foundations of Cooperative Work*. Lawrence Erlbaum Associates, Hillsdale, NJ, 1990.

- [GL90] R. V. Guha and Douglas B. Lenat. Cyc: A midterm report. *AI Magazine*, 11(3):33–59, Fall 1990.
- [GL91] Alan Garvey and Victor Lesser. Design-to-time real-time scheduling. COINS Technical Report 91–72, University of Massachusetts, 1991. To appear, *IEEE Transactions on Systems, Man and Cybernetics*, Special Issue on Scheduling, Planning, and Control, 23(6), 1993.
- [GL93] Alan Garvey and Victor Lesser. Design-to-time real-time scheduling. *IEEE Transactions on Systems, Man, and Cybernetics*, 23(6):1491–1502, 1993.
- [GLC<sup>+</sup>95] Benjamin N. Grosz, David W. Levine, Hoi Y. Chan, Colin J. Parris, and Joshua S. Auerbach. Reusable architecture for embedding rule-based intelligence in information agents. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [GLF81] Thomas D. Garvey, John D. Lowrance, and Martin A. Fischler. An inference technique for integrating knowledge from disparate sources. In *Proceedings of the Seventh International Joint Conference on Artificial Intelligence*, pages 319–325, Tokyo, Japan, August 1981.
- [GLR91] C. Grossner, J. Lyons, and T. Radhakrishnan. Validation of an expert system intended for research in distributed artificial intelligence. In *2nd CLIPS Conference, Johnson Space Center*, September 1991.
- [GLR92] C. Grossner, J. Lyons, and T. Radhakrishnan. Towards a tool for design of cooperating expert systems. In *4th International Conference on Tools for Artificial Intelligence*, November 1992.
- [GM84] Richard P. Gabriel and John McCarthy. Queue-based multi-processing Lisp. In *Proceedings of the 1984 ACM Symposium on Lisp and Functional Programming*, pages 25–43, Austin, Texas, August 1984.
- [GM95] Pan Gu and Anthony B. Maddox. A framework for distributed reinforcement learning. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 26–31, 1995.
- [Gmy91] Piotr J. Gmytrasiewicz. Rational reasoning system: Application of decision theory in autonomous robotics. Technical report, University of Michigan, In preparation 1991.
- [Got95] William H. Gotwald, Jr. *Army Ants: The Biology of Social Predation*. Cornell University, 1995.
- [GPG92] P. Gokulchander, A. Preece, and C. Grossner. Path hunter: A tool for finding the paths in a rule based expert system. DAI Technical Report DAI-0592-0012, Concordia University, Montreal Quebec, May 1992.
- [GPG<sup>+</sup>93] C. Grossner, A. Preece, P. Gokulchander, T. Radhakrishnan, and C. Y. Suen. Exploring the structure of rule based systems. In *Proc. Eleventh National Conference on Artificial Intelligence (AAAI 93)*, 1993.
- [GR88] Les Gasser and Nicolas Rouquette. Representing and using organizational knowledge in distributed AI systems. In *Proceedings of the 1988 Distributed AI Workshop*, May 1988.
- [GR89] J. Giarratano and G. Riley. *Expert Systems: Principles & Programming*. PWS-KENT, 1989.
- [GR90] C. Grossner and T. Radhakrishnan. Organizations for cooperating expert systems. In *22nd Southeastern Symposium on System Theory*, March 1990.
- [GR91] Andrew R. Golding and Paul S. Rosenbloom. Improving rule-based systems through case-based reasoning. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 22–27, 1991.



- [GR93a] Piotr J. Gmytrasiewicz and Jeffrey S. Rosenschein. The utility of embedded knowledge-oriented actions. In *Proceedings of the 12th International Workshop on Distributed Artificial Intelligence*, pages 155–169, May 1993.
- [GR93b] Claudia Goldman and Jeffrey S. Rosenschein. Emergent coordination through the use of cooperative state-changing rules. In *Working Papers of the 12th International Workshop on Distributed Artificial Intelligence*, pages 171–185, May 1993.
- [GR94] Claudia V. Goldman and Jeffrey S. Rosenschein. Emergent coordination through the use of cooperative state-changing rules. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 408–413, Seattle, WA, 1994.
- [GR95a] Michael P. Georgeff and Anand S. Rao. The semantics of intention maintenance for rational agents. In Chris S. Mellish, editor, *Proceedings of the Fourteenth International Joint Conference on Artificial Intelligence*, pages 704–710, Montréal, Canada, August 1995.
- [GR95b] Claudia V. Goldman and Jeffrey S. Rosenschein. Mutually supervised learning in multiagent systems. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 20–25, 1995.
- [Gra81] S. C. Graves. A review of production scheduling. *Operations Research*, 29:646–675, 1981.
- [Gra86] T. J. Grant. Lessons for OR from AI: A scheduling case study. *Operational Research Society*, 37:41–57, 1986.
- [Gra90] N. A. B. Gray. Capturing knowledge through top-down induction of decision trees. *IEEE Expert*, 5(3):41–50, June 1990.
- [Gra95] Robert S. Gray. Agent Tcl: A transportable agent system. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Gre82a] Peter E. Green. Distributed acoustic surveillance and tracking. In *Proceedings of the Distributed Sensor Networks Workshop*, 1982. Published by the Department of Computer Science, Carnegie-Mellon University, Pittsburgh, Pennsylvania.
- [Gre82b] Irene Greif. PCAL: A personal calendar. Technical Report TM-213, MIT Laboratory for Computer Science, Cambridge, Mass, 1982.
- [Gre87] Peter E. Green. AF: A flexible testbed for distributed AI research. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, Research Notes in Artificial Intelligence, chapter 6, pages 153–175. Pitman, 1987.
- [GRHL89] Les Gasser, Nicolas Rouquette, Randall W. Hill, and John Lieb. Representing and using organizational knowledge in DAI systems. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*, pages 55–78. Pitman, 1989.
- [Gro90a] C. Grossner. Ill structured problems. DAI Technical Report DAI-0690-0004, Concordia University, Montreal Quebec, June 1990.
- [Gro90b] C. Grossner. Information deficit: A metric for the data distribution in an organization. DAI Technical Report DAI-1190-0008, Concordia University, Montreal Quebec, November 1990.
- [Gro90c] C. Grossner. Organizations for cooperating expert systems. DAI Technical Report DAI-0790-0006, Concordia University, Montreal Quebec, July 1990.
- [Gro91a] C. Grossner. Blackbox training and testing software for human subjects. DAI Technical Report DAI-0791-0011, Concordia University, Montreal Quebec, July 1991.

- [Gro91b] C. Grossner. Optimality of an organization. DAI Technical Report DAI-0391-0009, Concordia University, Montreal Quebec, March 1991.
- [Gro91c] C. Grossner. Towards a tool for the design of organizations. DAI Technical Report DAI-0591-0010, Concordia University, Montreal Quebec, May 1991.
- [Gro95] Benjamin Grosf. Conflict resolution in advice taking and instruction for learning agents. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [Gru87] Jonathan Grudin. Social evaluation of the user interface: Who does the work and who gets the benefit? In H. Bullinger and B. Shackel, editors, *Human Computer Interaction – INTERACT87*, pages 805–811. North Holland, 1987.
- [GS82] Michael R. Genesereth and David E. Smith. Meta-level architecture. Technical report, Computer Science Department, Stanford University, Stanford, CA 94305, December 1982.
- [GS85] Barbara J. Grosz and Candace L. Sidner. Discourse structure and the proper treatment of interruptions. In *Proceedings of the Ninth International Joint Conference on Artificial Intelligence*, pages 832–839, August 1985.
- [GS88] B. J. Grosz and C. Sidner. Plans for discourse. In Cohen, Morgan, and Pollack, editors, *Intentions in Communication*. Massachusetts Institute of Technology, 1988.
- [GS90] B. J. Grosz and C. Sidner. Plans for discourse. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*. MIT Press, 1990.
- [GS94] Rose Gamble and Sandip Sen. Using formal specification to resolve conflicts between contracting agents. In *Proceedings of the AAAI-94 Workshop on Conflict Management in Cooperative Problem Solving*, pages 33–38, Seattle, WA, August 1994. AAAI Press.
- [GSH95] Laurence Glicoes, Richard Staats, and Michael N. Huhns. A multi-agent environment for department of defense distribution. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 14–19, 1995.
- [GW86] S. K. Goyal and W. Worrest. Expert systems in network management and maintenance. In *Proc. ICC*, pages 1225–1229, 1986.
- [GW88] S. Goyal and R. Worrest. Expert system applications to network management. In J. Leibowitz, editor, *Expert System Applications to Telecommunications*, volume 1, pages 3–44. John Wiley & Sons, 1988.
- [HA90] James Hendler and Ashok Agrawala. Mission critical planning: AI on the MARUTI real-time operating system. In *Proceedings of the 1990 DARPA Workshop on Innovative Approaches to Planning, Scheduling, and Control*, pages 77–84, November 1990.
- [HAC90] D. M. Hart, S. D. Anderson, and P. R. Cohen. Envelopes as a vehicle for improving the efficiency of plan execution. COINS Technical Report 90-21, University of Massachusetts, 1990.
- [Had95] Afsaneh Haddadi. Towards a pragmatic theory of interactions. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 133–139, San Francisco, CA, 1995. MIT Press.
- [Hag82] Roger Hagafors. The character of organizational problems: A classification system for organizational decision-making. Technical Report 317, University of Uppsala, Department of Psychology, Uppsala, Sweden, 1982.
- [HAG86] S. Y. Harmon, W. A. Aviles, and D. W. Gage. A technique for coordinating autonomous robots. In *Proceedings of the 1986 IEEE International Conference on Robotics and Automation*, pages 2029–2034, 1986.

- [Hal84] Robert H. Halstead, Jr. Implementation of Multilisp: Lisp on a multiprocessor. In *Proceedings of the 1984 ACM Symposium on Lisp and Functional Programming*, pages 9–17, Austin, Texas, August 1984.
- [Hal85] Robert H. Halstead, Jr. Multilisp: A language for concurrent symbolic computation. *ACM Transactions on Programming Languages and Systems*, 7(4):501–538, October 1985.
- [Hal86a] Joseph Y. Halpern, editor. *Theoretical Aspects of Reasoning About Knowledge: Proceedings of the 1986 Conference*. Morgan Kaufmann Publishers, Inc., March 1986.
- [Hal86b] Robert H. Halstead, Jr. An assessment of multilisp: Lessons from experience. *International Journal of Parallel Programming*, 15(6):459–501, 1986. (Due to publishing delays, this issue did not appear until late 1987.).
- [Ham86] Kristian J. Hammond. CHEF: A model of case-based planning. In *Proceedings of the National Conference on Artificial Intelligence*, pages 267–271, Philadelphia, PA, August 1986.
- [Han88] Steve Hanks. Representing and computing temporally scoped beliefs. In *Proceedings of the National Conference on Artificial Intelligence*, pages 501–505, August 1988.
- [Han94] Mark S. Hanh. Simulating evolution in a kolmogorov predator-prey model with genetic extensions. In John R. Koza, editor, *Artificial Life at Stanford 1994*, pages 44–53, Stanford, California, June 1994. Stanford Bookstore.
- [Hay45] F. Hayek. The uses of knowledge in society. *American Economic Review*, 35:519–530, September 1945.
- [Hay74] Patrick J. Hayes. Some problems and non-problems in representation theory. In *Proc. AISB Summer Conference*, pages 63–79, Univ. of Sussex, 1974.
- [Hay94] Thomas D. Haynes. A simulation of adaptive agents in a hostile environment. Master’s thesis, University of Tulsa, Tulsa, OK., April 1994.
- [HB77] Carl Hewitt and Henry Baker. Laws for communicating parallel processes. *Information Processing*, 1977.
- [HB91] Michael N. Huhns and David M. Bridgeland. Multiagent truth maintenance. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6), December 1991. (Special Issue on Distributed AI).
- [HB96] Salima Hassas and Jacques Bonneville. Towards a self-organizational approach for a parallel computation in a distributed production rule based system. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [HC90] David W. Hildum and Daniel D. Corkill. Solving dynamic sequencing problems. COINS Technical Report 90–63, University of Massachusetts, 1990.
- [HC95] Bernardo A. Huberman and Scott H. Clearwater. A multiagent system for controlling building environments. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 171–176, San Francisco, CA, 1995. MIT Press.
- [HD95] Marcus J. Huber and Edmund H. Durfee. Deciding when to commit to action during observation-based coordination. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 163–170, San Francisco, CA, 1995. MIT Press.
- [HdJ83] Carl Hewitt and Peter de Jong. Analyzing the roles of descriptions and actions in open systems. In *Proceedings of the National Conference on Artificial Intelligence*, pages 162–166, Washington, D.C., August 1983.

- [HE80] R. M. Haralick and G. L. Elliott. Increasing tree search efficiency for constraint satisfaction problems. *Artificial Intelligence*, 14(3):263–313, 1980.
- [HE84] Donna L. Hudson and Thelma Estrin. EMERGE - A data-driven medical decision making aid. *IEEE Transaction on Pattern Analysis and Machine Intelligence*, PAMI-6(1), January 1984.
- [Hea73] G. M. Heal. *The Theory of Economic Planning*. American Elsevier, 1973.
- [Hel84] John Joseph Helly, Jr. A distributed expert system for space shuttle flight control. Technical Report CSD-840038, University of California Computer Science Department, Los Angeles, California, 1984.
- [Hen80] Peter Henderson. *Functional Programming: Application and implementation*. Prentice-Hall International, 1980.
- [Hen90] James Hendler. Abstraction and reaction. In *Working Notes of the 1990 AAAI Spring Symposium on Planning*, pages 54–56, March 1990.
- [Hew76] Carl Hewitt. Viewing control structures as patterns of passing messages. Memo 410, MIT Artificial Intelligence Laboratory, December 1976.
- [Hew77] Carl Hewitt. Viewing control structures as patterns of passing messages. *Artificial Intelligence*, 8(3):323–364, Fall 1977.
- [Hew85] Carl Hewitt. The challenge of open systems. *Byte*, 10(4):223–242, April 1985.
- [Hew86] Carl Hewitt. Offices are open systems. *ACM Transactions on Office Information Systems*, 4(3):271–287, July 1986. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 321–330, Morgan Kaufmann, 1988.).
- [Hew91] Carl Hewitt. Open information systems semantics for distributed artificial intelligence. *Artificial Intelligence*, 47(1):79–106, 1991.
- [HG80] S. Harmon and D. Gage. Protocols for robot communications: Transport and content layers. In *Proceedings of the 1980 International Conference on Cybernetics and Society*, pages 1090–1097, 1980.
- [HGAB84] S. Y. Harmon, D. W. Gage, W. A. Aviles, and G. L. Bianchini. Coordination of intelligent subsystems in complex robots. In *Proceedings of the First IEEE Conference on Artificial Intelligence Applications*, pages 64–69, December 1984.
- [HH91] Tad Hogg and Bernardo A. Huberman. Controlling chaos in distributed systems. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6):1325–1332, December 1991. (Special Issue on Distributed AI).
- [HI91] Carl Hewitt and Jeff Inman. DAI betwixt and between: From “intelligent agents” to open systems science. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6):1409–1419, December 1991. (Special Issue on Distributed AI).
- [Hil94] David W. Hildum. *Flexibility in a Knowledge-Based System for Solving Dynamic Resource-Constrained Scheduling Problems*. PhD thesis, Department of Computer Science, University of Massachusetts, Amherst, September 1994.
- [HJR95] Matthew M Huntbach, Nick R Jennings, and Graem A Ringwood. How agents do it in stream logic programming. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 177–184, San Francisco, CA, 1995. MIT Press.
- [HK74] J. R. Hackman and R. E. Kaplan. Interventions into group processes: An approach to improving the effectiveness of groups. *Decision Science*, 5:459–480, 1974.

- [HK80] Carl Hewitt and Bill Kornfeld. Message passing semantics. *SIGART Newsletter*, (73):48, October 1980.
- [HL83] Carl Hewitt and Henry Lieberman. Design issues in parallel architectures for artificial intelligence. Memo 750, MIT Artificial Intelligence Laboratory, November 1983.
- [HL84a] Carl Hewitt and Henry Lieberman. Design issues in parallel architectures for artificial intelligence. In *Proceedings of the 28th IEEE Computer Society International Conference*, pages 418–423, San Francisco, CA, February 1984.
- [HL84b] Eva Hudlická and Victor Lesser. Design of a knowledge-based fault detection and diagnosis system. In *Proceedings of the Seventeenth Hawaiian International Conference System Sciences*, pages 224–230, January 1984.
- [HL84c] Eva Hudlická and Victor R. Lesser. Diagnosing the behavior of a distributed problem solving system. Technical Report 84-03, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, May 1984.
- [HL84d] Eva Hudlická and Victor R. Lesser. Meta-level control through fault detection and diagnosis. In *Proceedings of the Fourth National Conference on Artificial Intelligence*, pages 153–161, Austin, TX, August 1984.
- [HL85] Eva Hudlická and Victor R. Lesser. Construction and use of a causal model of a problem-solving system. *SIGART Newsletter*, (93):48–55, July 1985.
- [HL86] Eva Hudlická and Victor R. Lesser. Modeling and reasoning about problem-solving system behavior. Technical Report 86-04, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, February 1986.
- [HL87] Eva Hudlická and Victor R. Lesser. Modeling and diagnosing problem solving system behavior. *IEEE Transactions on Systems, Man, and Cybernetics*, 17(3):407–419, May 1987. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 490–502, Morgan Kaufmann, 1988.).
- [HL88] K. E. Huff and V. R. Lesser. A plan-based intelligent assistant that supports the software development process. In *Proceedings of the Third ACM Symposium on Software Development Environments*, Boston, MA, November 1988.
- [HL92] Marty A. Humphrey and Victor R. Lesser. Parallel scheduling based on abstract task models. Working Paper, Cooperative Distributed Problem Solving Laboratory, 1992.
- [HLPR86] Eva Hudlicka, Victor R. Lesser, Jasmina Pavlin, and Anil Rewari. Design of a distributed diagnosis system. Technical Report 86-63, Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts 01003, December 1986.
- [HLS96] Thomas Haynes, Kit Lau, and Sandip Sen. Learning cases to compliment rules for conflict resolution in multiagent systems. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [HM84a] Joseph Y. Halpern and Yoram Moses. Knowledge and common knowledge in a distributed environment. In *Third ACM Conference on Principles of Distributed Computing*, 1984.
- [HM84b] Joseph Y. Halpern and Yoram Moses. Towards a theory of knowledge and ignorance: Preliminary report. In *Proceedings of the 1984 Non-Monotonic Reasoning Workshop*, pages 125–143, October 1984.
- [HM85] Joseph Y. Halpern and Yoram Moses. A guide to the modal logics of knowledge and belief. In *Proceedings of the Ninth International Joint Conference on Artificial Intelligence*, pages 480–490, Los Angeles, CA, August 1985.

- [HM86] G. Huber and R. McDaniel. The decision making paradigm for organizational design. *Management Science*, pages 572–589, 1986.
- [HMSB87] Michael N. Huhns, Uttam Mukhopadhyay, Larry M. Stephens, and Ronald D. Bonnell. DAI for document retrieval: The MINDS project. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, Research Notes in Artificial Intelligence, chapter 9, pages 249–284. Pitman, 1987.
- [HNM95] Kôiti Hasida, Katashi Nagao, and Takashi Miyata. A game-theoretic account of cooperation in communication. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 140–147, San Francisco, CA, 1995. MIT Press.
- [Ho80] Y.-C. Ho. Team decision theory and information structures. *Proceedings of the IEEE*, 68, June 1980.
- [Hof83] Douglas R. Hofstadter. Metamagical themas: Computer tournaments of the prisoner’s dilemma suggest how cooperation evolves. *Scientific American*, 1983.
- [Hog95] Tad Hogg. Social dilemmas in computational ecosystems. In Chris S. Mellish, editor, *Proceedings of the Fourteenth International Joint Conference on Artificial Intelligence*, pages 711–716, Montréal, Canada, August 1995.
- [Hol96] Owen Holland. Multiagent systems: Lessons from social insects and collective robotics. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [Hor87] Eric J. Horvitz. Reasoning about beliefs and actions under computational resource constraints. In *Proceedings of the 1987 Workshop on Uncertainty in Artificial Intelligence*, 1987.
- [Hor88] Eric J. Horvitz. Reasoning under varying and uncertain resource constraints. In *Proceedings of the Seventh National Conference on Artificial Intelligence*, August 1988.
- [How66a] N. Howard. The theory of metagames. *General Systems*, 11:167–200, 1966.
- [How66b] R. A. Howard. Information value theory. *IEEE Transactions on Systems, Man, and Cybernetics*, pages 22–27, August 1966.
- [HPC93] Steve Hanks, Martha E. Pollack, and Paul R. Cohen. Benchmarks, test beds, controlled experimentation, and the design of agent architectures. *AI Magazine*, 14(4):17–42, Winter 1993.
- [HPS93] Michael Huhns, M. P. Papazoglou, and G. Schlageter. *Proceedings of International Conference on Intelligent and Cooperative Information Systems*. IEEE Computer Society Press, Los Alamitos, CA, 1993.
- [HR78] Allen R. Hanson and Edward M. Riseman. VISIONS: A computer system for interpreting scenes. In Allen R. Hanson and Edward M. Riseman, editors, *Computer Vision Systems*, pages 303–333. Academic Press, 1978.
- [HR80] Frederick Hayes-Roth. Towards a framework for distributed AI. *SIGART Newsletter*, (73):51–52, October 1980.
- [HR85] Barbara Hayes-Roth. A blackboard architecture for control. *Artificial Intelligence*, 26(3):251–321, July 1985. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 503–540, Morgan Kaufmann, 1988.).
- [HR89a] Barbara Hayes-Roth. Control in blackboard systems. In V. Jagannathan, Rajendra Dodhiawala, and Lawrence S. Baum, editors, *Blackboard Architectures and Applications*, pages 3–7. Academic Press, Inc., 1989.

- [HR89b] Barbara Hayes-Roth. A multi-processor interrupt-driven architecture for adaptive intelligent systems. In *Proceedings of the Third Annual AAAI Workshop on Blackboard Systems*, Detroit, August 1989. Also KSL-87-31.
- [HR91] R. P. Harrison and P. A. Ratcliffe. Towards standards for the validation of expert systems. *Expert Systems with Applications*, 2(4):251–258, 1991.
- [HRBvG95] Barbara Hayes-Roth, Lee Brownston, and Robert van Gent. Multiagent collaboration in directed improvisation. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 148–154, San Francisco, CA, 1995. MIT Press.
- [HREF<sup>+</sup>88a] F. Hayes-Roth, L. Erman, S. Fouse, J. Lark, and J. Davidson. ABE: A cooperative operating system and development environment. In Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 457–490. Morgan Kaufmann, 1988.
- [HREF<sup>+</sup>88b] Frederick Hayes-Roth, Lee D. Erman, Scott Fouse, Jay S. Lark, and James Davidson. ABE: A cooperative operation system and development environment. In Mark Richer, editor, *AI Tools and Techniques*. Ablex Publishing Corporation, 1988. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 457–489, Morgan Kaufmann, 1988.).
- [HRGJH86] Barbara Hayes-Roth, Alan Garvey, M. Vaughan Johnson, Jr., and Micheal Hewett. A modular and layered environment for reasoning about action. Technical Report KSL-86-38, Knowledge Systems Laboratory, Stanford University, 1986.
- [HRHR79] Barbara Hayes-Roth and Frederick Hayes-Roth. A cognitive model of planning. *Cognitive Science*, 3(4):275–310, October–December 1979.
- [HRHW<sup>+</sup>89] Barbara Hayes-Roth, Micheal Hewett, Richard Washington, Rattikorn Hewett, and Adam Seiver. Distributing intelligence within an individual. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*, pages 385–412. Pitman, 1989.
- [HRJGH86] Barbara Hayes-Roth, M. Vaughan Johnson, Jr., Alan Garvey, and Micheal Hewett. Application of the BB1 blackboard architecture to arrangement-assembly tasks. *International Journal for Artificial Intelligence in Engineering*, 1(2):85–94, 1986.
- [HRJGH88] Barbara Hayes-Roth, M. Vaughan Johnson, Alan Garvey, and Micheal Hewett. Building systems in the BB\* environment. In Robert Englemore and Tony Morgan, editors, *Blackboard Systems*, pages 543–561. Addison Wesley, Wokingham, England, 1988.
- [HRL77a] Frederick Hayes-Roth and Victor R. Lesser. Focus of attention in the Hearsay-II speech understanding system. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, pages 27–35, August 1977.
- [HRL77b] Frederick Hayes-Roth and Victor R. Lesser. Focus of attention in the Hearsay-II system. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, pages 27–35, Tbilisi, Georgia, USSR, August 1977.
- [HRWHH89] Barbara Hayes-Roth, Richard Washington, Rattikorn Hewett, and Micheal Hewett. Intelligent monitoring and control. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, pages 250–255, Detroit, August 1989.
- [HRWL83] F. Hayes-Roth, D. A. Waterman, and D. B. Lenat. *Building Expert Systems*. Addison Wesley, 1983.
- [HS95] Thomas Haynes and Sandip Sen. Evolving behavioral strategies in predators and prey. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 32–37, 1995.

- [HS96] Thomas Haynes and Sandip Sen. Evolving behavioral strategies in predators and prey. In Gerhard Weiß and Sandip Sen, editors, *Adaptation and Learning in Multiagent Systems*, Lecture Notes in Artificial Intelligence. Springer Verlag, Berlin, Spring 1996.
- [HSSW95a] Thomas Haynes, Sandip Sen, Dale Schoenefeld, and Roger Wainwright. Evolving a team. In E. V. Siegel and J. R. Koza, editors, *Working Notes for the AAAI Symposium on Genetic Programming*, Cambridge, MA, November 1995. AAAI.
- [HSSW95b] Thomas Haynes, Sandip Sen, Dale Schoenefeld, and Roger Wainwright. Evolving multiagent coordination strategies with genetic programming. *Artificial Intelligence*, 1995. (submitted for review).
- [HSW93] M. Henz, G. Smolka, and J. Würtz. Oz—A programming language for multi-agent systems. In *Proceedings of the Thirteenth International Joint Conference on Artificial Intelligence*, pages 404–409, Chambéry, France, August 1993.
- [HT95] Katsutoshi Hirayama and Jun’ichi Toyoda. Forming coalitions for breaking deadlocks. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 155–162, San Francisco, CA, 1995. MIT Press.
- [HTD90] James Hendler, Austin Tate, and Mark Drummond. AI planning : Systems and techniques. *AI Magazine*, 11(2):61–77, Summer 1990.
- [Hub88] Bernardo Huberman. *The Ecology of Computation*. Elsevier Science Publishers/North Holland, Amsterdam, 1988.
- [Hud86] Eva Hudlická. *Diagnosing Problem-Solving System Behavior*. PhD thesis, University of Massachusetts, February 1986. (Also published as Technical Report 86-03, Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts 01003, February 1986.).
- [Huh87] Michael Huhns, editor. *Distributed Artificial Intelligence*. Morgan Kaufmann, 1987.
- [Hul94] Ingemar Hulthage, editor. *Computational Organization Design*. AAAI Spring Symposium, 1994. Working Notes.
- [Hum90] Norman P. Hummon. Organizational structures and network processes. *Journal of Mathematical Sociology*, 15(2):149–161, 1990.
- [Hut95] William R. Hutchinson. Teaching an agent to speak and listen with understanding: Why and how? In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM ’95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [HW95] Thomas D. Haynes and Roger L. Wainwright. A simulation of adaptive agents in a hostile environment. In *Proceedings of the 1995 ACM Symposium on Applied Computing*, pages 318–323. ACM Press, 1995.
- [HWS94] Thomas Haynes, Roger Wainwright, and Sandip Sen. Evolving cooperation strategies. Technical Report UTULSA-MCS-94-10, The University of Tulsa, December 16, 1994.
- [HWS95] Thomas Haynes, Roger Wainwright, and Sandip Sen. Evolving cooperation strategies. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 450, San Francisco, CA, 1995. MIT Press. (poster).
- [HWSS95] Thomas Haynes, Roger Wainwright, Sandip Sen, and Dale Schoenefeld. Strongly typed genetic programming in evolving cooperation strategies. In Larry Eshelman, editor, *Proceedings of the Sixth International Conference on Genetic Algorithms*, pages 271–278, San Francisco, CA, 1995. Morgan Kaufmann Publishers, Inc.
- [IG90] Francois Felix Ingrand and Michael P. Georgeff. Managing deliberation and reasoning in real-time AI systems. In *Proceedings of the 1990 DARPA Workshop on Innovative Approaches to Planning, Scheduling, and Control*, pages 284–291, November 1990.



- [Int87] Intellicorp. *KEE V3.1 Reference Manual*, 1987.
- [IPLJ86] B. W. Irish, H. V. Parunak, P. W. Lozo, and R. Judd. A distributed heuristic strategy for material transportation. In *Proceedings 1986 Conference on Intelligent Systems and Machines*, Rochester, Michigan, 1986.
- [Ish92a] T. Ishida. The tower of babel: Towards organization-centered problem solving. In *Proc. 11th International Workshop on Distributed Artificial Intelligence*, pages 141–154, February 1992.
- [Ish92b] Toru Ishida. Tower of Babel: Towards organization-centered problem solving. In *Proceedings of the 11th Workshop on Distributed Artificial Intelligence*, pages 141–153, The Homestead, Michigan, 1992.
- [Ish95a] Toru Ishida. Parallel, distributed and multi-agent production systems – A research foundation for distributed artificial intelligence. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 416–422, San Francisco, CA, 1995. MIT Press. (invited speaker talk).
- [Ish95b] Toru Ishida. Two is not always better than one: Experiences in real-time bidirectional search. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 185–192, San Francisco, CA, 1995. MIT Press.
- [Iso95] Hideki Isozaki. Reasoning about belief based on common knowledge of observability of actions. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 193–200, San Francisco, CA, 1995. MIT Press.
- [IWY95] Tatsuaki Itoh, Takashi Watanabe, and Takahira Yamaguchi. Self organizational approach for integration of distributed expert systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 209–216, San Francisco, CA, 1995. MIT Press.
- [IY88] Keki B. Irani and Suk I. Yoo. A methodology for solving problems: Problem modeling and heuristic generation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 10(5):676–686, September 1988.
- [IY95] Akira Ito and Hiroyuki Yano. The emergence of cooperation in a society of autonomous agents. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 201–208, San Francisco, CA, 1995. MIT Press.
- [IYG90] Toru Ishida, Makoto Yokoo, and Les Gasser. An organizational approach to adaptive production systems. In *Proceedings of the National Conference on Artificial Intelligence*, pages 52–58, July 1990.
- [IZMG90] Adam Irgon, Jean Zolnowski, Karen J Murray, and Marvin Gersha. Expert system development - A retrospective of five systems. *IEEE Expert*, 5(3):25–40, June 1990.
- [Jag88] V. Jagannathan. Realizing the concurrent blackboard model. Technical Report BCS-G2010-61, Boeing Advanced Technology Center, P.O. Box 24346, MS 7L-64, Seattle, Washington 98124, March 1988.
- [Jai91] Raj Jain. *The art of computer systems performance analysis : techniques for experimental design, measurement, simulation, and modeling*. John Wiley and Sons, Inc., New York, 1991.
- [JBM75] F. Jelinek, L. R. Bahl, and R. L. Mercer. Design of a linguistic statistical decoder for the recognition of continuous speech. *IEEE Transactions on Information Theory*, IT-21(3):250–256, May 1975.

- [JCG87] Philip M. Johnson, Daniel D. Corkill, and Kevin Q. Gallagher. Integrating BB1-style control into the Generic Blackboard System. Technical Report 87-59, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, June 1987. (Presented at the AAAI-87 Blackboard Workshop, Seattle, Washington, July 13, 1987.).
- [JCL<sup>+</sup>95a] N. R. Jennings, J. M. Corera, L. Laresgoiti, E. H. Mamdani, F. Perriollat, P. Skarek, and L. Z. Varga. Using ARCHON to develop real-world DAI applications for electricity transportation management and particle accelerator control. *IEEE Expert*, 1995. Special issue on real world applications of DAI systems.
- [JCL95b] N. R. Jennings, J. M. Corera, and I. Laresgotti. Developing industrial multi-agent systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 423–431, San Francisco, CA, 1995. MIT Press. (invited speaker talk).
- [JD86] V. Jagannathan and Rajendra Dodhiawala. Distributed artificial intelligence: An annotated bibliography. *SIGART Newsletter*, (95):44–56, January 1986. (Also published as Chapter 12 of *Distributed Artificial Intelligence*, Michael N. Huhns, editor, Pitman, 1987.).
- [Jen84] Tom Jenkins. Meetings without meeting. *Working Woman*, pages 29–30, January 1984.
- [Jen93] Nick R. Jennings. Commitments and conventions: The foundation of coordination in multi-agent systems. *The Knowledge Engineering Review*, 8(3):223–250, 1993.
- [Jen95] Nick Jennings. Controlling cooperative problem solving in industrial multi-agent systems using joint intentions. *Artificial Intelligence*, 1995. to appear.
- [JGC87] Philip M. Johnson, Kevin Q. Gallagher, and Daniel D. Corkill. *GBB Reference Manual*. Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, GBB Version 1.1 edition, December 1987. (Published as Technical Report 87-120, Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts 01003, December 1987.).
- [JHR89] M. Vaughan Johnson, Jr. and Barbara Hayes-Roth. Simultaneous, dynamic integration of diverse reasoning methods. In V. Jagannathan, Rajendra Dodhiawala, and Lawrence S. Baum, editors, *Blackboard Architectures and Applications*, pages 57–73. Academic Press, Inc., 1989.
- [JJS88] M. Jarke, M. T. Jelassi, and M. F. Shakun. MEDIATOR: Towards a negotiation support system. In M. F. Shakun, editor, *Evolutionary Systems Design: Policy Making Under Complexity and Group Decision Support*. Holden-Day, Inc., Oakland, CA, 1988.
- [JK73] W. Jackobs and M. Kiefer. Robot decisions based on maximizing utility. In *Proceedings of the Third International Joint Conference on Artificial Intelligence*, pages 402–411, August 1973.
- [Joh88] R. Johansen. *Groupware: Computer Support for Business Teams*. Free Press, New York, NY, 1988.
- [Jos95] Anupam Joshi. To learn or not to learn. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 38–42, 1995.
- [JS79] Anita K. Jones and Karsten Schwans. TASK forces: Distributed software for solving problems of substantial size. In *4th International Conference on Software Engineering*, pages 315–330, Munich, Germany, September 1979.
- [JS95] Sanjay Jain and Arun Sharma. Team learning of formal languages. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.

- [JvdT96] J-F Raskin, Y-H Tan and L. van der Torre. Modeling deontic states in Petri nets. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [Kab95] Froduald Kabanza. Synchronizing multiagent plans using temporal logic specifications. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 217–224, San Francisco, CA, 1995. MIT Press.
- [Kad92] R. Kadia. Issues encountered in building a flexible software development environment: Lessons from the Arcadia project. In *Proceedings of the Fifth ACM SIGSOFT Symposium on Software Development Environments (SDE5)*, pages 169–180, Tyson’s Corner, VA, December 1992.
- [Kau90] H. Kautz. A circumscriptive theory of plan recognition. In Cohen *et al.*, editor, *Intentions in Communication*, pages 105–133. MIT Press, 1990.
- [KC82] J. F. Kelley and A. Chapanis. How professional persons keep their calendars: Implications for computerization. *Journal of Occupational Psychology*, 55:141–156, 1982.
- [KC87] André I. Khuri and John A. Cornell. *Response Surfaces: Designs and Analyses*. Marcel Dekker, New York, 1987.
- [KCL93] J. C. Kunz, G. P. Cohen, and R. E. Levitt. Modeling effects of organizational structure and communication tools on design team productivity. In Michael Prietula, editor, *AI and theories of groups & organizations: Conceptual and Empirical Research*. AAAI Workshop, 1993. Working Notes.
- [KCTL91] Subbarao Kambhampati, Mark Cutkosky, Marty Tenenbaum, and Soo Hung Lee. Combining specialized reasoners and general purpose planners: A case study. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 199–205, Anaheim, CA, July 1991.
- [KD89] Keiji Kanazawa and Thomas Dean. A model for projection and action. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, pages 985–990, August 1989.
- [KD93] J. R. Krebs and N. B. Davies. *An Introduction to Behavioural Ecology*. Blackwell Scientific Publications, 1993.
- [KD95] R. Khosla and T. Dillon. Distributed symbolic-subsymbolic agent architecture for configuring power network faults. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 451, San Francisco, CA, 1995. MIT Press. (poster).
- [KDK85] C. Kincaid, P. Dupont, and A. Kaye. Electronic calendars in the office: An assessment of user needs and current technology. *ACM Transactions on Office Information Systems*, 3(1):89–102, January 1985.
- [KEL91] Sarit Kraus, Eithan Ephrati, and Daniel Lehmann. Negotiation in a non-cooperative environment. *Journal of Experimental and Theoretical Artificial Intelligence*, 3(4):255–282, 1991.
- [Ker85] G. E. Kersten. NEGOT - group decision support system. *Information and Management*, 8:237–246, 1985.
- [Ket94] Steven Ketchpel. Forming coalitions in the face of uncertain rewards. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 414–419, Seattle, WA, 1994.
- [Key87] Charles Key. Cooperative planning in the pilot’s associate. In *Proceedings of the DARPA Knowledge-Based Planning Workshop*, pages 13.1–13.10, December 1987.

- [KF80] Philip Klahr and William S. Faight. Knowledge-based simulation. In *Proceedings of the First Annual National Conference on Artificial Intelligence*, pages 181–183, Stanford, CA, August 1980.
- [KG91] D. N. Kinny and M. P. Georgeff. Commitment and effectiveness of situated agents. In *Proceedings of the Twelfth International Joint Conference on Artificial Intelligence*, pages 82–88, Sydney, Australia, August 1991.
- [KG94] Taha Khedro and Michael R. Genesereth. Progressive negotiation for resolving conflicts among distributed heterogeneous cooperating agents. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 381–386, Seattle, WA, 1994.
- [KGBK78] R. E. Kahn, S. A. Gronemeyer, J. Burchfiel, and R. C. Kunzelman. Advances in packet radio technology. In *Proceedings of the IEEE*, pages 1468–1496, November 1978.
- [KGR96] David Kinny, Michael Georgeff, and Anand Rao. A methodology and modelling technique for systems of BDI agents. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [KH81] William A. Kornfeld and Carl E. Hewitt. The scientific community metaphor. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):24–33, January 1981. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 311–320, Morgan Kaufmann, 1988.).
- [KH95a] Daniel Kuokka and Larry Harada. Matchmaking for information agents. In Chris S. Mellish, editor, *Proceedings of the Fourteenth International Joint Conference on Artificial Intelligence*, pages 672–678, Montréal, Canada, August 1995.
- [KH95b] Daniel Kuokka and Larry Harada. On using KQML for matchmaking. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 239–245, San Francisco, CA, 1995. MIT Press.
- [KHH89] J. O. Kephart, T. Hogg, and B. A. Huberman. Dynamics of computational ecosystems: Implications for DAI. In Michael N. Huhns and Les Gasser, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [Kic95] Alexander Kick. An agent algebra for the formal description and verification of multi-agent systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 452, San Francisco, CA, 1995. MIT Press. (poster).
- [Kim87] Jin H. Kim. Distributed inference for plausible classification. *Pattern Recognition Letters*, 5:195–201, March 1987.
- [KIO95] Kazuhiro Kuwabara, Toru Ishida, and Nobuyasu Osato. Agentalk: Coordination protocol description for multiagent systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 455, San Francisco, CA, 1995. MIT Press. (poster).
- [Kis78] H. Kise. A solvable case of the one-machine scheduling problem with ready and due times. *Operations Research*, 26(1), 1978.
- [Kit94] James E. Kittock. The impact of locality and authority on emergent conventions: Initial observations. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 420–425, Seattle, WA, 1994.
- [KL82] J. B. Kadane and P. D. Larkey. Subjective probability and the theory of games. *Management Science*, 28(2):113–120, February 1982.
- [KL89] K. Kuwabara and V. R. Lesser. Extended protocol for multi-stage negotiation. In *Proceedings of the Ninth Workshop on Distributed AI*, September 1989.

- [KL94] Daniel Kuaokka and Brian Livezey. A collaborative parametric design agent. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 387–393, Seattle, WA, 1994.
- [Kle87] Jack P. C. Kleijnen. *Statistical Tools for Simulation Practitioners*. Marcel Dekker, New York, 1987.
- [Kle89] Keith S. Klemba. OpenView’s architectural models. In Branislav Meandzija and Jil Westcott, editors, *Integrated Network Management I*. Elsevier Science Publishers/North Holland, New York, 1989.
- [Kle90a] Mark Klein. *Conflict Resolution in Cooperative Design*. PhD thesis, University of Illinois at Urbana-Champaign, January 1990.
- [Kle90b] Mark Klein. Supporting conflict resolution in cooperative design systems. In *Proceedings of the 1990 Distributed AI Workshop*, October 1990.
- [Kle91] Mark Klein. Supporting conflict resolution in cooperative design systems. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6), December 1991. (Special Issue on Distributed AI).
- [Klu95] Matthias Klusch. Coalition-based cooperation between intelligent agents. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 453, San Francisco, CA, 1995. MIT Press. (poster).
- [KM95] Alois Knoll and J. Meinkoehn. Hierarchical and lateral coordination in MAS : An analysis of message traffic flow. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 225–232, San Francisco, CA, 1995. MIT Press.
- [KMB69] W. L. Kilmer, W. S. McCulloch, and J. Blum. A model of the vertebrate central command system. *International Journal of Man-Machine Studies*, 1(3):279–309, July 1969.
- [KMJ95] E. A. Kendall, M. T. Malkoun, and C. H. Jiang. A methodology for developing agent based systems. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [KN80] K. Konolige and N. J. Nilsson. Multiple-agent planning systems. In *Proceedings of AAAI-80*, pages 138–142, Stanford University, Palo Alto, CA., 1980. AAAI.
- [Kno91] Craig A. Knoblock. Search reduction in hierarchical problem solving. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, Anaheim, July 1991.
- [KO95] Satoshi Kurihara and Michio Okada. Self-organization based on coordinated actions of autonomous agents. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 454, San Francisco, CA, 1995. MIT Press. (poster).
- [Koh81] Walter H. Kohler. A survey of techniques for synchronization and recovery in decentralized computer systems. *Computing Surveys*, 13(2):149–183, June 1981.
- [Koh82] Heinz Kohler. *Intermediate Microeconomics Theory and Applications*. Scott, Foresman and Company, 1982.
- [Kol88] Janet L. Kolodner, editor. *Proceedings of a Workshop on Case-Based Reasoning (DARPA)*. Morgan Kaufmann, 1988.
- [Kol93] Janet L. Kolodner. *Case-Based Reasoning*. Morgan Kaufmann Publishers, 1993.
- [Kol95] Michael Kolb. A cooperation language. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 233–238, San Francisco, CA, 1995. MIT Press.

- [Kon82] Kurt Konolige. Circumscriptive ignorance. In *Proceedings of the National Conference on Artificial Intelligence*, pages 202–204, Pittsburgh, PA, August 1982.
- [Kon83] Kurt Konolige. A deductive model of belief. In *Proceedings of the Eighth International Joint Conference on Artificial Intelligence*, pages 377–381, Karlsruhe, Federal Republic of Germany, August 1983.
- [Kon84] K. Konolige. A deductive model of belief. In *Proceedings of the Eighth International Joint Conference on Artificial Intelligence*, pages 377–381, August 1984.
- [Kon86] K. Konolige. *A Deduction Model of Belief*. Pitman Publishers/Morgan Kaufman Publishers, San Mateo, CA., 1986.
- [Kor79] William A. Kornfeld. ETHER: A parallel problem solving system. In *Proceedings of the Sixth International Joint Conference on Artificial Intelligence*, pages 490–492, Tbilisi, Georgia, USSR, August 1979.
- [Kor89] Richard E. Korf. Multi-agent decision trees. In *Proceedings of the 1989 Distributed AI Workshop*, pages 293–306, September 1989.
- [Kor92] Richard E. Korf. A simple solution to pursuit games. In *Working Papers of the 11th International Workshop on Distributed Artificial Intelligence*, pages 183–194, February 1992.
- [Kov92] C. De Koven. An experiment in automating case-based knowledge acquisition. Master’s thesis, Concordia University, 1992.
- [Koz94] John R. Koza. Evolution of emergent cooperative behavior using genetic programming. In Ray Paton, editor, *Computing with Biological Metaphors*, pages 280–297. Chapman & Hall, London, UK, 1994.
- [KP89] Kurt Konolige and Martha E. Pollack. Ascribing plans to agents. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, pages 924–930, August 1989.
- [KR74] F. E. Kast and J. E. Rosenzweig. *Organization and Management*. McGraw-Hill, second edition, 1974.
- [KR78] Brian W. Kernighan and Dennis M. Ritchie. *The C Programming Language*. Prentice Hall, 1978.
- [KR89] Matthew J. Katz and Jeffrey S. Rosenschein. Plans for multiple agents. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [KR90] C. De Koven and T. Radhakrishnan. An experiment in group problem solving. In *Multi-User Interfaces and Applications, Greece*, September 1990.
- [Kra93] Sarit Kraus. Agents contracting tasks in non-collaborative environments. In *Proceedings of the National Conference on Artificial Intelligence*, pages 243–248, 1993.
- [Kro96] C. Krogh. The right way to analyse agents. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [KS75] S. R. Kimbleton and G. M. Schneider. Computer communication networks: Approaches, objectives, and performance considerations. *Computing Surveys*, 7(3):129–173, September 1975.
- [KS86] G. E. Kersten and T. Szapiro. Generalized approach to modelling negotiations. *European Journal of Operational Research*, 26:142–149, 1986.

- [KS91] Henry F. Korth and Abraham Silberschatz. *Database System Concepts*. McGraw Hill, New York, NY, 1991.
- [KS96] Matthias Klusch and Onn Shehory. Coalition formation among rational information agents. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [KSC<sup>+</sup>94] Henry Kautz, Bart Selman, Michael Coen, Steven Ketchpel, and Chris Ramming. An experiment in the design of software agents. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 438–443, Seattle, WA, 1994.
- [KSS95] Tetsuo Kinoshita, Kenji Sugawara, and Norio Shiratori. Agent-based framework for developing distributed systems. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [KSY85] J. F. Kurose, M. Schwartz, and Y. Yemini. A microeconomic approach to optimization of channel access policies in multiaccess networks. In *Proceedings of the Fifth International Symposium on Distributed Computing Systems*, pages 70–80, May 1985.
- [Kum95] Ram Kumar. Internet information resource discovery tools: Current and future trends. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [KvM90] Thomas Kreifelts and Frank v. Martial. A negotiation framework for autonomous agents. In *Proceedings of the Second European Workshop on Modeling Autonomous Agents and Multi Agent Worlds*, Paris, France, August 1990.
- [KW91a] Sarit Kraus and Jonathan Wilkenfeld. The function of time in cooperative negotiations. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 179–184, Anaheim, CA, July 1991.
- [KW91b] Sarit Kraus and Jonathan Wilkenfeld. Negotiations over time in a multi agent environment. In *Proceedings of the Twelfth International Joint Conference on Artificial Intelligence*, pages 56–61, 1991.
- [KW93] Sarit Kraus and John Wilkenfeld. Multiagent negotiation under time constraints. Technical Report CS 2649, Institute for Advanced Computer Studies, University of Maryland, College Park, 1993.
- [Lan85] Amy L. Lansky. Behavioral specification and planning for multiagent domains. Technical Report 360, SRI International, Menlo Park CA, 1985.
- [Lan90] Amy L. Lansky. Localized search for controlling automated reasoning. In *Proceedings of the Workshop on Innovative Approaches to Planning, Scheduling and Control*, pages 115–125, San Diego, November 1990. Morgan Kaufmann.
- [Lar78] Robert Edwa Larson. *Principles of Dynamic Programming*. M. Dekker, New York, 1978. pts 1 and 2.
- [Lat87] Bruno Latour. *Science in Action*. Harvard University Press, Cambridge, MA, 1987.
- [Lau88] P. R. Laughlin. Collective induction: Group performance, social combination processes and mutual majority and minority influences. *Journal of Personality and Social Psychology*, 54:254–267, 1988.
- [LC78a] Victor R. Lesser and Daniel D. Corkill. Cooperative distributed problem solving: A new approach for structuring distributed systems. Technical Report 78-7, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, May 1978.

- [LC78b] Victor R. Lesser and Daniel D. Corkill. Functionally-accurate, cooperative distributed systems. In *Proceedings of the Distributed Sensor Nets Workshop*, pages 21–26, Carnegie-Mellon University, Pittsburgh, Pennsylvania, December 1978.
- [LC79a] Victor R. Lesser and Daniel D. Corkill. The application of Artificial Intelligence techniques to cooperative distributed processing. In *Proceedings of the Sixth International Joint Conference on Artificial Intelligence*, pages 537–540, Tbilisi, Georgia, USSR, August 1979.
- [LC79b] John D. Lowrance and Daniel D. Corkill. The design of GRASPER 1.0: A programming language extension for graph processing. Technical Report 79-6, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, February 1979.
- [LC81] Victor R. Lesser and Daniel D. Corkill. Functionally accurate, cooperative distributed systems. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):81–96, January 1981.
- [LC83a] Victor R. Lesser and Daniel D. Corkill. The distributed vehicle monitoring testbed. *AI Magazine*, 4(3):63–109, Fall 1983.
- [LC83b] Victor R. Lesser and Daniel D. Corkill. The Distributed Vehicle Monitoring Testbed: A tool for investigating distributed problem solving networks. *AI Magazine*, 4(3):15–33, Fall 1983. (Also published in *Blackboard Systems*, Robert S. Englemore and Anthony Morgan, editors, pages 353–386, Addison-Wesley, 1988 and in *Readings from AI Magazine: Volumes 1–5*, Robert Englemore, editor, pages 69–85, AAAI, Menlo Park, California, 1988).
- [LC87a] V. R. Lesser and D. Corkill. DVMT A tool for investigation of distributed problem solving networks. In M. N. Huhns, editor, *Distributed Artificial Intelligence*. Morgan Kaufmann, 1987.
- [LC87b] Victor R. Lesser and Daniel D. Corkill. Distributed problem solving. In *Encyclopedia of Artificial Intelligence*, pages 245–251. John Wiley & Sons, 1987.
- [LC93a] Z. Lin and K. Carley. Does agent style matter in organizational decision making? In Michael Prietula, editor, *AI and theories of groups & organizations: Conceptual and Empirical Research*. AAAI Workshop, 1993. Working Notes.
- [LC93b] Z. Lin and K. Carley. Proactive or reactive: An analysis of the effect of agent style on organizational decision-making performance. *International Journal of Intelligent Systems in Accounting, Finance, and Management*, 2(4):271–289, December 1993. Special issue on “Mathematical and Computational Models of Organizations: Models and Characteristics of Agent Behavior”.
- [LCD87] Victor R. Lesser, Daniel D. Corkill, and Edmund H. Durfee. An update on the Distributed Vehicle Monitoring Testbed. Technical Report 87-111, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, December 1987.
- [LCHW88] Victor R. Lesser, Daniel D. Corkill, Joseph A. Hernandez, and Robert C. Whitehair. Details of goal processing in blackboard architectures. Technical Report 88-19, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, March 1988. (Presented at the AAAI-88 Blackboard Workshop, Seattle, Washington, August 24, 1988.).
- [LCK<sup>+</sup>94] R. E. Levitt, P. G. Cohen, J. C. Kunz, C. Nass, T. Christiansen, and Y. Jin. The virtual design team: Simulating how organizational structure and communication tools affect team performance. In K. M. Carley and M. J. Prietula, editors, *Computational Organization Theory*. Lawrence Erlbaum Associates, 1994.



- [LCN90] Hector J. Levesque, Philip R. Cohen, and José H. T. Nunes. On acting together. In *Proceedings of the Eighth National Conference on Artificial Intelligence*, pages 94–99, July 1990.
- [LCP<sup>+</sup>81] Victor Lesser, Daniel Corkill, Jasmina Pavlin, Larry Lefkowitz, Eva Hudlická, Richard Brooks, and Scott Reed. A high-level simulation testbed for cooperative distributed problem solving. Technical Report 81-16, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, June 1981. (Revised and shortened versions of this report appeared in *Proceedings of the Distributed Sensor Networks Workshop*, MIT Lincoln Laboratory, Lexington, Massachusetts, pages 247–270, January 1982, and in *Proceedings of the Third International Conference on Distributed Computer Systems*, pages 341–349, October 1982.).
- [LCP<sup>+</sup>82] Victor R. Lesser, Daniel D. Corkill, Jasmina Pavlin, Larry Lefkowitz, Eva Hudlická, Richard Brooks, and Scott Reed. A high-level simulation testbed for cooperative distributed problem solving. In *Proceedings of the 3rd International Conference on Distributed Computing Systems*, pages 341–349, October 1982.
- [LCS<sup>+</sup>88] Thomas J. Laffey, Preston A. Cox, James L. Schmidt, Simon M. Kao, and Jackson Y. Read. Real-time knowledge-based systems. *AI Magazine*, 9(1):27–45, Spring 1988.
- [LCWH89a] Victor R. Lesser, Daniel D. Corkill, Robert C. Whitehair, and John A. Hernandez. Focus of control through goal relationships. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, pages 497–503, Detroit, August 1989.
- [LCWH89b] Victor R. Lesser, Daniel D. Corkill, Robert C. Whitehair, and Joseph A. Hernandez. Goal relationships and their use in a blackboard architecture. In V. Jagannathan, Rajendra Dodhiawala, and Lawrence S. Baum, editors, *Blackboard architectures and applications*, volume 3 of *Perspectives in Artificial Intelligence*, chapter 1, pages 9–26. Academic Press, 1989.
- [LD95a] Jaeho Lee and Edmund H. Durfee. A microeconomic approach to intelligent resource sharing in multiagent systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 457, San Francisco, CA, 1995. MIT Press. (poster).
- [Ld95b] Michael Luck and Mark d’Inverno. Engagement and cooperation in motivated agent modelling. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [Ld95c] Michael Luck and Mark d’Inverno. A formal framework for agency and autonomy. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 254–268, San Francisco, CA, 1995. MIT Press.
- [Ld95d] Michael Luck and Mark d’Inverno. A formal view of social dependence networks. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [LE77] Victor R. Lesser and Lee D. Erman. A retrospective view of the Hearsay-II architecture. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, pages 790–800, Tbilisi, Georgia, USSR, August 1977.
- [LE80a] V. R. Lesser and L. D. Erman. Distributed interpretation: A model and an experiment. *IEEE Transactions on Computers*, C-29(12):1144–1163, December 1980.
- [LE80b] Victor R. Lesser and Lee D. Erman. Distributed interpretation: A model and experiment. *IEEE Transactions on Computers*, C-29(12):1144–1163, December 1980. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 120–139, Morgan Kaufmann, 1988.).

- [LeB86a] Thomas J. LeBlanc. Shared memory versus message-passing in a tightly-coupled multiprocessor: A case study. Technical Report Butterfly Project Report 3, University of Rochester Computer Science Department, January 1986.
- [LeB86b] Thomas J. LeBlanc. Shared memory versus message-passing in a tightly-coupled multiprocessor: A case study. Technical Report 3, University of Rochester Computer Science Department, January 1986.
- [Lei87] David Leinweber. Expert systems in space. *IEEE Expert*, 2(1):26–36, Spring 1987.
- [Len75] Douglas B. Lenat. Beings: Knowledge as interacting experts. In *Proceedings of the Fourth International Joint Conference on Artificial Intelligence*, pages 126–133, Stanford, CA, August 1975. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 161–168, Morgan Kaufmann, 1988.).
- [Les80a] Victor R. Lesser. Cooperative distributed problem solving and organizational self-design. *SIGART Newsletter*, page 46, October 1980. (Part of the “Report on the Workshop on Distributed AI”).
- [Les80b] Victor R. Lesser. Models of problem-solving. *SIGART Newsletter*, (73):51, October 1980. (Part of the “Report on the Workshop on Distributed AI”).
- [Les84] Victor R. Lesser. Panel on “Information Processors as Organizations”, Presented at the National Conference on Artificial Intelligence, Austin, Texas. (Description of the panel appeared in *Science*, 225:1136–1137, September 10, 1984), August 1984.
- [Les89] Victor R. Lesser. Coordination in distributed problem solving networks. DARPA proposal, COINS, University of Massachusetts, 1989.
- [Les91] Victor R. Lesser. A retrospective view of FA/C distributed problem solving. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6), December 1991. (Special Issue on Distributed AI).
- [Les95] Victor R. Lesser. Multiagent systems: An emerging subdiscipline of AI. *ACM Computing Surveys*, 27(3):340–342, September 1995.
- [LF77] John D. Lowrance and Daniel P. Friedman. Hendrix’s model for simultaneous actions and continuous processes: An introduction and implementation. *International Journal of Man-Machine Studies*, pages 537–581, 1977.
- [LF87] Amy L. Lansky and Dan Fogelson. Localized representation and planning methods for parallel domains. In *Proceedings of the National Conference on Artificial Intelligence*, pages 240–245, Seattle, WA, August 1987.
- [LFE74] V. S. Lesser, R. D. Fennell, L. D. Ertman, and D. R. Reddy. Organization of the hearsay II speech understanding systems. In *Proc. IEEE Symp. Speech Recognition, Carnegie-Mellon University, Pittsburgh, PA*, April 1974.
- [LFS96] Giordano Lanzola, Sabina Falasconi, and Mario Stefanelli. Cooperating agents implementing distributed patient management. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [LG90a] J. Lyons and C. Grossner. A blackbox expert system: Detailed design. DAI Technical Report DAI-1190-0007, Concordia University, Montreal Quebec, November 1990.
- [LG90b] J. Lyons and C. Grossner. A blackbox expert system: Preliminary design. DAI Technical Report DAI-0790-0005, Concordia University, Montreal Quebec, July 1990.
- [LG90c] J. Lyons and C. Grossner. A blackbox expert system: Software requirements specification. DAI Technical Report DAI-0490-0003, Concordia University, Montreal Quebec, April 1990.

- [LG90d] J. Lyons and C. Grossner. A blackbox expert system: User requirements. DAI Technical Report DAI-0190-0001, Concordia University, Montreal Quebec, January 1990.
- [LI86] C. A. Luckhardt and K. B. Irani. An algorithmic solution of N-person games. In *Proceedings of the National Conference on Artificial Intelligence*, pages 158–162, Philadelphia, PA, August 1986.
- [Lib96] Henry Libermann. Software agents: The MIT approach. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [Lim82] Willie Y-P. Lim. HISDL—A structure description language. *Communications of the ACM*, 25(11):823–830, November 1982.
- [Lin90] Theodore A. Linden. Transformational synthesis: An approach to large-scale planning applications. In *Proceedings of the Workshop on Innovative Approaches to Planning, Scheduling and Control*, pages 126–130, San Diego, November 1990. Morgan Kaufmann.
- [Lin94] Z. Lin. A theoretical evaluation of measures of organizational design: Interrelationship and performance predictability. In K. M. Carley and M. J. Prietula, editors, *Computational Organization Theory*, chapter 6, pages 113–159. Lawrence Erlbaum Associates, Hillsdale, NJ, 1994.
- [Lip89] Stanley B. Lippman. *C++ Primer*. Addison-Wesley, Reading, Mass., 1989.
- [Lit92] Glenn E. Littlepage. Impact of meeting procedures on meeting effectiveness. *Journal of Business and Psychology*, pages 361–369, Spring 1992.
- [LJN95] John E. Laird, Randolph M. Jones, and Paul E. Nielsen. Multiagent coordination in distributed interactive battlefield simulations. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 456, San Francisco, CA, 1995. MIT Press. (poster).
- [LK88] J. Lancaster and J. L. Kolodner. Varieties of learning from problem solving experience. In *Proceedings of the Tenth Annual Conference of the Cognitive Science Society*, pages 447–453, Montreal, Canada, 1988. The Cognitive Science Society.
- [LL67a] P. R. Lawrence and J. W. Lorsch. Organization and environment: Managing differentiation and integration. Technical report, 1967.
- [LL67b] Paul Lawrence and Jay Lorsch. *Organization and Environment*. Harvard University Press, Cambridge, MA, 1967.
- [LL88] Susan Lander and Victor Lesser. Negotiation among cooperating experts. In *Proceedings of the 1988 Distributed AI Workshop*, May 1988.
- [LL89] Susan Lander and Victor R. Lesser. A framework for the integration of cooperative knowledge-based systems. In *Proceedings of the 4th IEEE International Symposium on Intelligent Control*, pages 472–477, September 1989.
- [LL90] L. L. Lefkowitz and V. R. Lesser. Knowledge acquisition as knowledge assimilation. In B. Gaines and J. Boose, editors, *Machine Learning and Uncertain Reasoning Vol. 3*, pages 37–48. 1990.
- [LL92] S. Lander and V. R. Lesser. Negotiated search: Organizing cooperative search among heterogeneous expert agents. In *Proceedings of the Fifth International Symposium on Artificial Intelligence Applications in Manufacturing and Robotics*, December 1992.
- [LLC91] Susan E. Lander, Victor R. Lesser, and Margaret E. Connell. Knowledge-based conflict resolution for cooperation among expert agents. In D. Sriram, R. Logher, and S. Fukuda, editors, *Computer-Aided Cooperative Product Development*. Springer Verlag, 1991.

- [LLL90] Brigitte Lâassri, Hassan Lâassri, and Victor R. Lesser. Negotiation and its role in cooperative distributed problem solving. In *Proceedings of the Tenth International Workshop on Distributed AI*, Texas, October 1990. Also COINS TR-90-39.
- [LLN87] Jane W. S. Liu, Kwei-Jay Lin, and Swaminathan Natarajan. Scheduling real-time, periodic jobs using imprecise results. Technical report, Department of Computer Science, University of Illinois, 1987.
- [LLP94] Susan E. Lander, Victor R. Lesser, and M. V. Nagendra Prasad. Information sharing among heterogeneous reusable agents in cooperative distributed search. In *13th International Workshop on Distributed Artificial Intelligence (DAI-94)*, Seattle, WA, July 1994.
- [LM77] James A. Levin and James A. Moore. Dialogue-games: Metacommunication structures for natural language interaction. *Cognitive Science*, 1(4):395-420, October 1977. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 385-397, Morgan Kaufmann, 1988.).
- [LM82] Richard J. LeBlanc and Arthur B. Maccabe. The design of a programming language based on connectivity networks. In *Proceedings of the 3rd International Conference on Distributed Computing Systems*, pages 532-541, Miami/Ft. Lauderdale, Florida, October 1982.
- [LM86] A. Lewin and J. Minton. Determining organizational effectiveness: Another look and an agenda for research. *Management Science*, pages 514-538, 1986.
- [LM88] J. Lee and T. W. Malone. How can groups communicate when they use different languages? In *Proceedings of the Conference of Office Information Systems*, Palo Alto, CA., 1988.
- [LMM94] Yezdi Lashkari, Max Metral, and Pattie Maes. Collaborative interface agents. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 444-449, Seattle, WA, 1994.
- [LNR87] John E. Laird, Allen Newell, and Paul S. Rosenbloom. SOAR: An architecture for general intelligence. *Artificial Intelligence*, pages 1-64, 1987.
- [Lok96] G-J C. Lokhorst. Modelling business transactions: Deontic Petri nets. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [Lot94] Christopher G. Lott. Terrain flattening by autonomous robot: a genetic programming application. In John R. Koza, editor, *Genetic Algorithms at Stanford 1994*, pages 99-109, Stanford, California, 94305-3079 USA, Phone 415-329-1217 or 800-533-2670, December 1994. Stanford Bookstore.
- [Low79] John D. Lowrance. GRASPER 1.0 reference manual. Technical Report 78-20, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, February 1979.
- [Low83] John Douglas Lowrance. *Dependency-graph Models of Evidential Support*. PhD thesis, University of Massachusetts, February 1983. (Also published as Technical Report 82-26, Department of Computer and Information Science, University of Massachusetts, Amherst, Massachusetts 01003, Septemeber 1982.).
- [LP88] Victor R. Lesser and Jasmina Pavlin. Performing approximate processing to address real-time constraints. COINS Technical Report 87-126, University of Massachusetts, 1988.
- [LPD88] Victor R. Lesser, Jasmina Pavlin, and Edmund Durfee. Approximate processing in real-time problem solving. *AI Magazine*, 9(1):49-61, Spring 1988. (Also published in *Blackboard Architectures and Applications*, V. Jagannathan, R. Dodhiawala, and L. Baum, editors, Academic Press, 1989.).

- [LPS81] B. W. Lampson, M. Paul, and H. J. Siegart, editors. *Distributed Systems Architecture and Implementation*. Springer-Verlag, New York, 1981.
- [LR57] R. D. Luce and H. Raiffa. *Games and Decisions*. John Wiley and Sons, 1957.
- [LR58] R. Duncan Luce and Howard Raiffa. *Games and Decisions: Introduction and Critical Survey*. John Wiley & Sons, New York, 1958.
- [LR92] Ran Levy and Jeffrey S. Rosenschein. A game theoretic approach to the pursuit problem. In *Working Papers of the 11th International Workshop on Distributed Artificial Intelligence*, pages 195–213, February 1992.
- [LRN86] John Laird, Paul Rosenbloom, and Allen Newell. *Universal Subgoaling and Chunking: The automatic generation and learning of goal hierarchies*. Kluwer Academic Publishers, 1986.
- [LRP80] V. R. Lesser, S. Reed, and J. Pavlin. Quantifying and simulating the behavior of knowledge-based interpretation systems. In *Proceedings of the First Annual National Conference on Artificial Intelligence*, pages 111–115, Stanford, CA, August 1980.
- [LS95a] JyiShane Liu and Katia P. Sycara. Exploiting problem structure for distributed constraint optimization. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 246–254, San Francisco, CA, 1995. MIT Press.
- [LS95b] Andreas Lux and Donald Steiner. Understanding cooperation: an agent’s perspective. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 261–268, San Francisco, CA, 1995. MIT Press.
- [LSB79] Victor R. Lesser, Daniel Serrain, and Jeff Bonar. PCL: A process-oriented job control language. In *Proceedings of the First International Conference on Distributed Computing Systems*, pages 315–329, 1979.
- [LT88] Lu and Thompson. A distributed artificial intelligence approach to integrated engineering design. In *Proceedings of the First International Conference in Industrial and Engineering Applications of Artificial Intelligence and Expert Systems*, Tullahoma, Tennessee, 1988.
- [Lur87] M Luria. Goal conflict concerns. In *Proceedings of IJCAI-87*, pages 1025–1031, Milan, Italy, 1987. IJCAI.
- [LW78] R. Lacoss and R. Walton. Strawman design of a DSN to detect and track low flying aircraft. In *Proceedings of the Distributed Sensor Nets Workshop*, pages 41–52, 1978. Published by the Department of Computer Science, Carnegie-Mellon University, Pittsburgh, Pennsylvania.
- [LW90] K. S. Leung and M. H. Wong. An expert system shell using structured knowledge : An object oriented approach. *IEEE Computer*, 23(3):38–47, March 1990.
- [LW95] Britta Lenzmann and Ipke Wachsmuth. A user-adaptive interface agency for interaction with a virtual environment. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 43–46, 1995.
- [M. 96] P-Y. Schobbens M. Ryan, O. Rodriguez. Counterfactuals and updates as inverse modalities. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [MAE<sup>+</sup>65] John McCarthy, Paul W. Abrahams, Daniel J. Edwards, Timothy P. Hart, and Michael I. Levin. *Lisp 1.5 Programmer’s Manual*. MIT Press, second edition, 1965.
- [Mae89] Pattie Maes. The dynamics of action selection. In *IJCAI*, pages 991–997, 1989.

- [Mai85] Oded Z. Maimon. A multi-robot control experimental system with random parts arrival. In *Proceedings of the 1985 IEEE International Conference on Robotics and Automation*, pages 895–900, 1985.
- [Mal87a] T. W. Malone. Modelling coordination in organizations and markets. *Management Science*, 33(10), 1987.
- [Mal87b] Thomas W. Malone. Modeling coordination in organizations and markets. *Management Science*, 33(10):1317–1332, 1987. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 151–158, Morgan Kaufmann, 1988.).
- [Mal88] Thomas W. Malone. What is coordination theory? In *Proceedings of the National Science Foundation Coordination Theory Workshop*, February 1988.
- [Mal90] Thomas W. Malone. Organizing information processing systems: Parallels between human organizations and computer systems. In Scott P. Robertson, Wayne Zachary, and John B. Black, editors, *Cognition, Computing, and Cooperation*, pages 56–83. Ablex Publishing Corporation, 1990.
- [Mal94] Daniel Malmer. Hive: Development of a language among artificial life forms. In John R. Koza, editor, *Artificial Life at Stanford 1994*, pages 99–107, Stanford, California, June 1994. Stanford Bookstore.
- [Man79] William C. Mann. Design for dialogue comprehension. In *Proceedings of the Seventeenth Annual Meeting of the Association for Computational Linguistics*, August 1979.
- [Mar91] James G. March. How decisions happen in organizations. *Human-Computer Interaction*, 6:95–117, 1991.
- [Mat95] Maja J. Mataric. Learning in multi-robot systems. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 32–37, 1995.
- [Mat96] Maja Mataric. Merging co-evolution and reinforcement learning. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [Maz87] Murray S. Mazer. Exploring the use of distributed problem-solving in office support systems. In *Proceedings of the IEEE Computer Society Symposium on Office Automation*, pages 217–225, April 1987.
- [Maz91] Murray S. Mazer. Reasoning about knowledge to understand distributed AI systems. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6), December 1991. (Special Issue on Distributed AI).
- [MC87] Douglas J. MacIntosh and Susan E. Conry. SIMULACT: A generic tool for simulating distributed systems. In *Proceedings of the 1987 Eastern Simulation Conference*, 1987.
- [MC88] Kelly E. Murray and Daniel D. Corkill. Common Lisp object representation strategies: The UMass Parallel Implementation of Common Lisp implementation. Technical Report 88-35, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, April 1988.
- [MC91] Thomas W. Malone and Kevin Crowston. Toward an interdisciplinary theory of coordination. Center for Coordination Science Technical Report CCS-TR-120, MIT Sloan School of Management, 1991.
- [MC93] Mauro Manela and J. A. Campbell. Designing good pursuit problems as testbeds for Distributed AI: a novel application of Genetic Algorithms. In *Fifth European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Neuchâtel, Switzerland, August 24-27 1993.

- [MC94a] Thomas W. Malone and Kevin Crowston. The interdisciplinary study of coordination. *ACM Computing Surveys*, 26(1):87, March 1994.
- [MC94b] Geoffrey F. Miller and Dave Cliff. Co-evolution of pursuit and evasion I: Biological and game-theoretic foundations. Technical Report CSRP311, August 1994.
- [McC65] Warren S. McCulloch. *Embodiments of Mind*, chapter What's in the Brain that Ink may Character?, pages 387–397. MIT Press, 1965.
- [McD78] Drew McDermott. Planning and acting. *Cognitive Science*, (2):71–109, 1978.
- [McD82] Drew McDermott. A temporal logic for reasoning about processes and plans. *Cognitive Science*, 6:101–155, 1982.
- [McD86] Drew V. McDermott. A critique of pure reason. Technical Report YaleU/CSD/RR-480, Yale University, New Haven, CT, June 1986.
- [McK82] Bill McKelvey. *Organizational Systematics*. University of California Press, Berkeley, CA, 1982.
- [MCLP93] Thomas W. Malone, Kevin Crowston, Jintae Lee, and Brian Pentland. Tools for inventing organizations: Toward a handbook of organizational processes. Center for Coordination Science Technical Report 141, MIT Sloan School of Management, 1993.
- [MCM91] Douglas J. MacIntosh, Susan E. Conry, and Robert A. Meyer. Distributed automated reasoning: Issues in coordination, cooperation, and performance. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6), December 1991. (Special Issue on Distributed AI).
- [MD84] Drew McDermott and Ernest Davis. Planning routes through uncertain territory. *Artificial Intelligence*, pages 107–156, 1984.
- [MD90a] S. Maxwell and H. Delany. *Designing Experiments and Analyzing Data*. Wadsworth Publishing Company, 1990.
- [MD90b] Thomas A. Montgomery and Edmund H. Durfee. MICE users guide. Technical Report CSE-TR-64-90, Computer Science and Engineering Division, EE and CS Department, University of Michigan, Ann Arbor, MI 48109, June 1990.
- [MD90c] Thomas A. Montgomery and Edmund H. Durfee. Using MICE to study intelligent dynamic coordination. In *Proceedings of the Second Computer Society International Conference on Tools for Artificial Intelligence*, pages 438–444, November 1990.
- [MD92] Thomas A. Montgomery and Edmund H. Durfee. Search reduction in hierarchical distributed problem solving. In *Proceedings of the Eleventh International Workshop on Distributed AI*, Glen Arbor, Michigan, February 1992.
- [MDCL95] Craig Munday, Jirapun Dangedej, Tim Cross, and Dickson Lukose. Motivation and perception mechanisms in mobil agent for electronic commerce. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [MDH85] Yoram Moses, Danny Dolev, and Joseph H. Halpern. Cheating husbands and other stories: A case study of knowledge, action, and communication (preliminary version). In *Proceedings of the Fourth Annual ACM Symposium on Principles of Distributed Computing*, pages 215–223, August 1985.
- [MDS91] David J. Musliner, Edmund H. Durfee, and Kang G. Shin. Execution monitoring and recovery planning with time. In *Proceedings of the Seventh IEEE International Conference on AI Applications*, pages 385–388, February 1991.

- [Mee80] James R. Meehan. Everything you always wanted to know about authority structures but were unable to represent. In *Proceedings of the First Annual National Conference on Artificial Intelligence*, pages 212–214, Stanford, CA, August 1980.
- [MFGH86] T. W. Malone, R. E. Fikes, K. R. Grant, and M. T. Howard. Market-like task scheduling in distributed computing environments. Technical Report CISR-WP-139, MIT Center for Information Systems Research, March 1986.
- [MFH83] T. W. Malone, R. E. Fikes, and M. T. Howard. Enterprise: A market-like task scheduler for distributed computing environments. Working Paper CISR WP 111 (Sloan WP 1537-84), Center for Information Systems Research, Massachusetts Institute of Technology, Cambridge, MA 02139, 1983.
- [MFN<sup>+</sup>95] James Mayfield, Tim Finin, Rajkumar Narayanaswamy, Chetan Shah, William MacCartney, and Keith Goolsbey. The cyclic friends network: Getting cyc agents to reason together. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [MG92] Michael Mesterton-Gibbons. *An Introduction to Game-Theoretic Modelling*. Addison-Wesley, 1992.
- [MGR95] Yishay Mor, Claudia V. Goldman, and Jeffrey S. Rosenschein. Learn your opponent's strategy (in polynomial time)! In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 53–58, 1995.
- [MGT<sup>+</sup>87] Thomas W. Malone, Kenneth R. Grant, Franklyn A. Turbak, Stephen A. Brobst, and Michael D. Cohen. Intelligent information-sharing systems. *Communications of the ACM*, 30(5):390–402, 1987.
- [MH94] M. Masuch and Z. Huang. A logical deconstruction of organizational action: Formalizing J.D. Thompson's *organizations in action* in a multi-agent action logic. CCSOM Working Paper 94-120, Department of Statistics and Methodology, PSCW, University of Amsterdam, Netherlands, 1994.
- [MH95] J. Masthoff and R. Van Hoe. A view on the architecture and design of highly autonomous and situated agents. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 458, San Francisco, CA, 1995. MIT Press. (poster).
- [MHM85] Jr. David M. McKeown, Jr. Wilson A. Harvey, and John McDermott. Rule-based interpretation of aerial imagery. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PAMI-7(5):570–585, September 1985.
- [Mil78] James Grier Miller. *Living Systems*. McGraw-Hill, 1978.
- [Mil87] James S. Miller. *Multischeme: A parallel processing system based on MIT Scheme*. PhD thesis, Massachusetts Institute of Technology, Cambridge, MA, September 1987. (Also published as Technical Report MIT/LCS/TR-402, Laboratory for Computer Science, Massachusetts Institute of Technology, Cambridge, Massachusetts, 02139, September 1987.).
- [Mil90] L. A. Miller. Dynamic testing of knowledge bases using the heuristic testing. *Expert Systems with Applications*, 1(3):249–270, 1990.
- [Min73] H. Mintzberg. *The Nature of Managerial Work*. Harper & Row, New York, NY, 1973.
- [Min79a] Marvin Minsky. The society theory of thinking. In Patrick Henry Winston and Richard Henry Brown, editors, *Artificial Intelligence: An MIT Perspective*, pages 423–450. MIT Press, 1979.
- [Min79b] H. Mintzberg. *The Structuring of Organizations: A Synthesis of the Research*. Prentice Hall, Englewood Cliffs, N.J., 1979.



- [Min80] Marvin Minsky. K-lines: A theory of memory. *Cognitive Science*, pages 117–133, 1980.
- [Min83] H. Mintzberg. *Structure in fives: designing effective organizations*. Prentice-Hall, Englewood Cliffs, N.J., 1983.
- [Min85] Marvin Minsky. *The Society of Mind*. Simon and Schuster, 1985.
- [Mis91] Frederick C. Mish, editor. *Webster’s Ninth New Collegiate Dictionary*, Springfield, MA, 1991. Merriam Webster.
- [MJ89] Cindy L. Mason and Rowland R. Johnson. DATMS: A framework for distributed assumption based reasoning. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*, pages 293–317. Pitman, 1989.
- [MJS<sup>+</sup>88] C. Mason, R. Johnson, R. Searfus, D. Lager, and T. Canales. A seismic event analyzer for nuclear test ban treaty verification. In *Proceedings of the Third International Conference on Applications of Artificial Intelligence in Engineering*, August 1988.
- [MK93] Pattie Maes and Robyn Kozierok. Learning interface agents. In *Proceedings of the Eleventh National Conference on Artificial Intelligence*, pages 459–464, July 1993.
- [MKS87] T. W. McGuire, S. Kiesler, and J. Siegel. Group and computer-mediated discussion effects in risk decision making. *Journal of Personality and Social Psychology*, 5:917–930, 1987.
- [ML90] Theresa Moehlman and Victor Lesser. Cooperative planning and decentralized negotiation in multi-fireboss phoenix. In *Proceedings of the 1990 DARPA Workshop on Innovative Approaches to Planning, Scheduling, and Control*, pages 144–159, November 1990.
- [MLB92a] Theresa Moehlman, Victor Lesser, and Brandon Buteau. Decentralized negotiation: An approach to the distributed planning problem. *Group Decision and Negotiation*, 1(2):161–192, 1992.
- [MLB92b] Theresa A. Moehlman, Victor R. Lesser, and Brandon L. Buteau. Decentralized negotiation: An approach to the distributed planning problem. *IEEE Transactions on Systems, Man, and Cybernetics*, Special Issue on Control, Planning and Scheduling 1992. Submitted.
- [MLKn91] Xiyi Miao, Peter B. Luh, David L. Kleimnman, and David A. Casta non. Distributed stochastic resource allocation in teams. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(1):61–70, January/February 1991.
- [MMW89] P. R. Monge, C. McSween, and Wyer. A profile of meetings in corporate america. In *Results of the 3M Meeting Effectiveness Study*. Anneberg School of Communications, University of Southern California, Los Angeles, 1989.
- [MNH96] Hitoshi Matsubara, Itsuki Noda, and Kazuo Hiraki. Learning of cooperative actions in multiagent systems: a case study of pass play in soccer. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [Moe84] Terry M. Moe. The new economics of organization. *American Journal of Political Science*, 28(4):739–777, November 1984.
- [Mor87] Leora Morgenstern. Knowledge preconditions for actions and plans. In *Proceedings of the National Conference on Artificial Intelligence*, pages 867–874, Seattle, WA, August 1987. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 151–158, Morgan Kaufmann, 1988.).
- [Mou90] Bernard Moulin. Strategic planning in expert systems. *IEEE Expert*, 5(2):69–75, April 1990.

- [Mou95] Vassilis S. Moustakis. CEG: A case based decision modeling architecture. *European Journal of Operational Research*, 84:170–191, 1995.
- [MR77] J. W. Meyer and B. Rowan. Institutionalized organizations: Formal structures as myth and ceremony. *American Journal of Sociology*, 83:340–363, 1977.
- [MR95a] Dario Maio and Stefano Rizzi. Unsupervised multi-agent exploration of structured environments. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 269–275, San Francisco, CA, 1995. MIT Press.
- [MR95b] Yishay Mor and Jeffrey S. Rosenschein. Time and the prisoner’s dilemma. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 276–282, San Francisco, CA, 1995. MIT Press.
- [MRS82] Gordon I. McCalla, Larry Reid, and Peter F. Schneider. Plan creation, plan execution, and knowledge acquisition in a dynamic microworld. *International Journal of Man-Machine Studies*, 16:89–112, 1982.
- [MRtPRG87] James L. McClelland, David E. Rumelhart, and the PDP Research Group. *Parallel Distributed Processing: Explorations in the Microstructure of Cognition (2 Volumes)*. MIT Press, 1987.
- [MS58] James G. March and Herbert A. Simon. *Organizations*. John Wiley and Sons, New York, 1958.
- [MS83] J. W. Meyer and W. R. Scott. *Organizational Environments: Ritual and Rationality*. Sage Publications, Beverly Hills, CA, 1983.
- [MS84] T. W. Malone and S. A. Smith. Tradeoffs in designing organizations: Implications for new forms of human organizations and computer systems. Working Paper CISR WP 112 (Sloan WP 1541-84), Center for Information Systems Research, Massachusetts Institute of Technology, Cambridge, MA 02139, March 1984.
- [MS86] Bertrand R. Meunier and Melvin F. Shakan. *Compromise, Negotiation, and Group Decision*. R. Deidel Publishing Co., Dorbrecht, 1986.
- [MS87] N. Muscettola and S. Smith. A probabilistic framework for resource-constrained multi-agent planning. In *Proceedings of IJCAI-87*, pages 1063–1066, Milan, Italy, 1987. IJCAI.
- [MS89] R. P. Mohanty and M. K. Siddiq. Multiple projects-multiple resources-constrained scheduling: Some studies. *IJPR*, 27(2):261–280, 1989.
- [MS90] P. Mi and W. Scacchi. A knowledge-based environment for modeling and simulating software engineering processes. *IEEE Transactions on Knowledge and Data Engineering*, 2(3):283–294, September 1990.
- [MSC82] David McArthur, Randy Steeb, and Stephanie Cammarata. A framework for distributed problem solving. In *Proceedings of the Second National Conference on Artificial Intelligence*, pages 181–184, Pittsburgh, PA, August 1982.
- [MSHB87] Uttam Mukhopadhyay, Larry Stephens, Michael Huhns, and Ronald Bonnell. An intelligent system for document retrieval in distributed office environments. *Journal of the American Society for Information Science*, 37:123–135, 1987. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 565–577, Morgan Kaufmann, 1988.).
- [MSK<sup>+</sup>87] S. Matwin, S. Szpakowicz, G. Kersten, W. Michalowski, and Z. Koperczak. Logic-based tools for negotiation support. Technical Report 87-10, University of Ottawa, 1987.
- [MST83] John Mylopoulos, Tetsutaro Shibahara, and John K. Tsotsos. Building knowledge-based systems: The PSN experience. *IEEE Computer*, pages 83–89, October 1983.

- [MTS89] R. Mirchandaney, D. Towsley, and J. Stankovic. Analysis of the effects of delays on load sharing. *IEEE Transactions on Computers*, 38(11):1513–1525, November 1989.
- [MW89] Branislav Meandzija and Jil Westcott. *Integrated Network Management I*. Elsevier Science Publishers/North Holland, New York, 1989.
- [MW95] Tracy Mullen and Michael P. Wellman. A simple computational market for network information services. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 283–289, San Francisco, CA, 1995. MIT Press.
- [MYB87] T. W. Malone, J. Yates, and R. I. Benjamin. Electronic markets and electronic hierarchies. *Communications of the ACM*, 30:484–497, 1987.
- [NA90] Keung-Chi Ng and Bruce Abramson. Uncertainty management in expert systems. *IEEE Expert*, 5(2):29–48, April 1990.
- [NBH<sup>+</sup>85] Heinrich Niemann, Horst Bunke, Ingrid Hofman, Gerhard Sagerer, Friedrich Wolf, and Herbert Feistel. A knowledge based system for analysis of gated blood pool studies. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PAMI-7(3), May 1985.
- [Nea90] Richard E. Neapolitan. *Probabilistic Reasoning in Expert Systems*. John Wiley and Sons, 1990.
- [Nel82] William R. Nelson. REACTOR: An expert system for diagnosis and treatment of nuclear reactor accidents. In *AAAI*, pages 296–301, 1982.
- [New80] Allen Newell. Reasoning, problem solving, and decision processes: The problem space as a fundamental category. In R. Nickerson, editor, *Attention and Performance VIII*. Erlbaum, 1980.
- [New90] Allen Newell. *Unified Theories of Cognition*. Harvard University Press, 1990.
- [NF78] Penny Nii and Edward A. Feigenbaum. Rule based understanding of signals. In D. A. Waterman and Frederick Hayes-Roth, editors, *Pattern-Directed Inference Systems*, pages 483–501. Academic Press, 1978.
- [NF96] I. Nunes and J. Fiadeiro. Generalising transition systems for durative actions. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [NFAR82] H. Penny Nii, Edward A. Feigenbaum, John J. Anton, and A. J. Rockmore. Signal-to-symbol transformation: HASP/SIAP case study. *AI Magazine*, 3(2):23–35, Spring 1982.
- [NG89] Walter E. Natemeyer and Jay S. Gilberg, editors. *Classics of Organizational Behavior*. The Interstate Printers & Publishers, Inc., 1989.
- [NHLS94] Daniel E. Neiman, David W. Hildum, Victor R. Lesser, and Tuomas Sandholm. Exploring meta-level information in a distributed scheduling system. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 394–400, Seattle, WA, 1994.
- [Nii86a] H. Penny Nii. Blackboard application systems and a knowledge engineering perspective (part 2). *AI Magazine*, August 1986.
- [Nii86b] H. Penny Nii. Blackboard application systems: Blackboard systems from a knowledge engineering perspective. *AI Magazine*, 7(3):82–106, Conference 1986.
- [Nii86c] H. Penny Nii. Blackboard systems: The blackboard model of problem solving and the evolution of blackboard architectures. *AI Magazine*, 7(2):38–53, Summer 1986.

- [Nil79] N. J. Nilsson. A production system for automatic deduction. In J. E. Hayes, Donald Michie, and L. I. Mikulich, editors, *Machine Intelligence*, pages 101–126. Halsted Press, 1979.
- [Nil80a] Nils J. Nilson. Two heads are better than one. *SIGART Newsletter*, (73):43, October 1980.
- [Nil80b] Nils J. Nilsson. *Principles of Artificial Intelligence*. Tioga, 1980.
- [NL77] David A. Nadler and III Edward E. Lawler. *Perspectives on Behavior in Organizations*, chapter Motivation: A diagnostic approach, pages 26–34. McGraw-Hill, 1977.
- [NL88] Sergei Nirenburg and Victor Lesser. Providing intelligent assistance in distributed office environments. In Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 590–598. Morgan Kaufmann, 1988.
- [NMK83] Sanjai Narain, David McArthur, and Philip Klahr. Large-scale system development in several Lisp environments. In *Proceedings of the Eighth International Joint Conference on Artificial Intelligence*, pages 859–861, Karlsruhe, Federal Republic of Germany, August 1983.
- [NNO95] Hideyuki Nakashima, Itsuki Noda, and Ichiro Ohsawa. Organic programming for multi-agents. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 459, San Francisco, CA, 1995. MIT Press. (poster).
- [Nor92] Peter Norvig. *Paradigms of Artificial Intelligence Programming: Case Studies in Common Lisp*. Morgan-Kaufman, 1992.
- [NR94] V. Nilakant and H. Rao. Agency theory and uncertainty in organizations: An evaluation. *Organizational Studies*, 15(5):649–672, 1994.
- [NS72] A. Newell and H. A. Simon. *Human Problem Solving*. Prentice Hall, 1972.
- [NS79] W. E. Newman and R. F. Sproull. *Principles of Interactive Computer Graphics*. McGraw-Hill, 1979.
- [NS91] S. J. Noronha and V. V. S. Sarma. Knowledge-based approaches for scheduling problems: A survey. *IEEE Transactions on Knowledge and Data Engineering*, 3(2):160–171, June 1991.
- [NT74] Kumpati S. Narendra and M. A. L. Thathachar. Learning automata—A survey. *IEEE Transactions on Systems, Man, and Cybernetics*, 4(4):323–334, July 1974.
- [NT90] Chisato Numaoka and Mario Tokoro. OpenMind: Conversational model for open systems. In *Proceedings of Workshop on Object-Oriented Computing 1990*, March 1990.
- [NTR87] Dennis J. Nicklaus, Siu S. Tong, and Carol J. Russo. Engineous: A knowledge directed COmputer aided design shell. In *Proceedings of the Third Conference on Artificial Intelligence Applications*. IEEE Computer Society Press, February 1987.
- [Num95] Chisato Numaoka. Introducing blind hunger dilemma. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 290–296, San Francisco, CA, 1995. MIT Press.
- [Num96] Chisato Numaoka. Bacterial evolution algorithm for rapid adaptation. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [ODCS88] Mark F. Orelup, John R. Dixon, Paul R. Cohen, and Melvin K. Simmons. Dominic II: Meta-level control in iterative redesign. In *Proceedings of the Seventh National Conference on Artificial Intelligence*, pages 25–30, August 1988.

- [ODvdK88] J. M. Orbell, R. M. Dawes, and A. J. C. van der Kragt. Explaining discussion-induced cooperation. *Journal of Personality and Social Psychology*, pages 811–819, 1988.
- [OGMR90] Timothy J. O’Leary, Michael Gaul, Kathleen E. Moffitt, and A. Essam Radwan. Validating expert systems. *IEEE Expert*, 5(3):51–58, June 1990.
- [OGS95] Sascha Ossowski and Ana Garcia-Serrano. A knowledge-level model of coordination. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [O’H95] G. M. P. O’Hare. The agent factory: An environment for the fabrication of distributed artificial intelligence systems. In G. O’Hare and N. Jennings, editors, *Foundations of Distributed Artificial Intelligence*, chapter 17. Wiley Inter-Science, 1995. Forthcoming.
- [OKM83] Horst Oberquelle, Ingbert Kupka, and Susanne Maass. A view of human-machine communication and co-operation. *International Journal of Man-Machine Studies*, 19:309–333, 1983.
- [OL90] R. M. O’Keefe and S. Lee. An integrative model of expert system verification. *Expert Systems with Applications*, 1(3):231–236, 1990.
- [OPH89] P. S. Ow, M. J. Prietula, and W. Hsu. Configuring knowledge-based systems to organizational structures: Issues and examples in multiple agent support. In L. F. Pau, J. Motiwalla, Y. H. Pao, and H. H. Teh, editors, *Expert Systems in Economics, Banking, and Management*, pages 309–318. North-Holland, Amsterdam, 1989.
- [OPL94a] Tim Oates, M. V. Nagendra Prasad, and Victor Lesser. Networked information retrieval as distributed problem solving. In *Proceedings of the CIKM ’94 Workshop on Intelligent Information Agents*, Gaithersburg, Maryland, December 1994.
- [OPL94b] Tim Oates, M. V. Nagendra Prasad, and Victor R. Lesser. Cooperative information gathering: A distributed problem solving approach. Technical Report 94-66, Department of Computer Science, University of Massachusetts, September 1994.
- [OPLD95] Tim Oates, M. V. Nagendra Prasad, Victor R. Lesser, and Keith Decker. A distributed problem solving approach to cooperative information gathering. In *AAAI-95 Spring Symposium on Information Gathering from Heterogeneous, Distributed Environments*, Stanford University, Stanford, CA, March 1995.
- [OS88] P. S. Ow and S. F. Smith. Viewing scheduling as an opportunistic problem solving process. In R. G. Jeroslow, editor, *Annals of Operations Research: Approaches to Intelligent Decision Support*. Baltzer Scientific Publishing Co., 1988.
- [OS90] O’Reilly-Staff. X toolkit intrinsics reference manual. *O’Reilly and Associates*, 5, 1990.
- [Osa95] Ei-Ichi Osawa. A metalevel coordination strategy for reactive cooperative planning. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 297–303, San Francisco, CA, 1995. MIT Press.
- [OSH88] P. S. Ow, S. F. Smith, and R. E. Howie. CSS: A cooperative scheduling system. In M. D. Oliff, editor, *Expert Systems and Intelligent Manufacturing*. Elsevier Science Publishing Co., 1988.
- [OST96] Toru Ohira, Ryusuke Sawatari, and Mario Tokoro. Distributed interaction with computon. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [Ous94] John K. Ousterhout. *Tcl and the Tk Toolkit*. Addison-Wesley, Reading, MA, 1994.
- [PA90] Walton A. Perkins and Avis Austin. Adding temporal reasoning to expert system building environments. *IEEE Expert*, 5(1):23–30, February 1990.

- [Pad92] John F. Padgett. The alchemist of contingency theory. *American Journal of Sociology*, 97(5):1462–1470, March 1992.
- [Pan92] R. R. Panko. Patterns of managerial communication. *Journal of Organizational Computing*, 2(1):95–122, 1992.
- [Pan93] R. R. Panko. Why are we meeting? In *Proceedings of the 27th Hawaiian International Conference of System Sciences*, volume IV, pages 169–178, 1993.
- [Par85] H. Van Dyke Parunak. Manufacturing experience with the contract net. In *Proceedings of the 1985 Distributed Artificial Intelligence Workshop*, pages 67–91, December 1985.
- [par87] H. Van Dyke parunak. Manufacturing experience with the contract net. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, Research Notes in Artificial Intelligence, chapter 10, pages 285–310. Pitman, 1987.
- [Par90a] H. Van Dyke Parunak. Distributed AI and manufacturing control: Some issues and insights. In Y. Demazeau and J.-P. Muller, editors, *Decentralized AI*, pages 81–101. North Holland, 1990.
- [Par90b] H. Van Dyke Parunak. Toward a formal model of inter-agent control. In *Proceedings of the 1990 Distributed AI Workshop*, October 1990.
- [Pat95] Dhiraj K. Pathak. Service agents for customer self-service. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Pav83] Jasmina Pavlin. Predicting the performance of distributed knowledge-based systems: A modeling approach. In *Proceedings of the Third National Conference on Artificial Intelligence*, pages 314–319, Washington, D.C., August 1983.
- [Pav84] Jasmina Pavlin. A model for analysis and simulation of distributed knowledge-based systems. In *Proceedings of the Sixth European Conference on Artificial Intelligence*, September 1984.
- [Pav85] Jasmina Pavlin. *A Model for Prediction and Description of Knowledge Based System Behavior*. PhD thesis, University of Massachusetts, May 1985.
- [Pax78] W. H. Paxton. *The Executive System*. D. E. Walker, Elsevier, North-Holland, 1978.
- [PB96] Martha E. Pollack and Cristina Bicchieri. The potential for the evolution of cooperation among web agents. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [PC84] Jasmina Pavlin and Daniel D. Corkill. Selective abstraction of AI system activity. In *Proceedings of the National Conference on Artificial Intelligence*, pages 264–268, Austin, TX, August 1984.
- [PC88] K. Parsaye and M. Chignell. *Expert Systems for Experts*. John Wiley, 1988.
- [PCL87] H. Edward Pattison, Daniel D. Corkill, and Victor R. Lesser. Instantiating descriptions of organizational structures. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, Research Notes in Artificial Intelligence, chapter 3, pages 59–96. Pitman, 1987.
- [Pea85] Judea Pearl. How to do with probabilities what people say you can't. In *Proceedings of the Second Conference on Artificial Intelligence Applications*, pages 6–12. IEEE Computer Society Press, December 1985.
- [Pea86] Judea Pearl. A constraint - propagation approach to probabilistic reasoning. In L. N. Kanal and J. F. Lemmer, editors, *Uncertainty in Artificial Intelligence*, pages 357–369. Elsevier Science Publishers, P.O. Box 1991, 1000 BZ Amsterdam, The Netherlands, 1986.

- [Pea87] Judea Pearl. Distributed revision of composite beliefs. *Artificial Intelligence*, 33(2):173–215, 1987.
- [PEG95] Terry Payne, Peter Edwards, and Claire Green. Experience with rule induction and k-nearest neighbour methods for interface agents that learn. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [Per86] Charles Perrow. *Complex Organizations*. Random House, New York, 1986.
- [Per90] C. R. Perrault. An application of default logic to speech act theory. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*. MIT Press, 1990.
- [Pet84] J. H. Peterson. A comparison of exact approaches for solving the multiple constrained resource project scheduling problem. *Management Science*, 30(7):854–867, 1984.
- [Pet92] Charles J. Petrie, Jr., editor. *Enterprise Integration Modeling: Proceedings of the First International Conference*. MIT Press, 1992.
- [Pfe83] J. Pfeffer. Organizational demography. In L. L. Cummings and B. M. Staw, editors, *Research in Organizational Behavior*, volume 5. JAI Press, Greenwich, CT, 1983.
- [Pfe91] J. Pfeffer. Organizational theory and structural perspectives on management. *Journal of Management*, 17(4):789–803, December 1991.
- [Pfe93] J. Pfeffer. Barriers to the advance of organizational science: Paradigm development as an independent variable. *Academy of Management Review*, 18(4):599–620, 1993.
- [PG88] Don Peacock and Mark Giuffrida. Big brother: A network services expert. In *USENIX Conference*, Summer 1988.
- [PGGR93a] A. Preece, P. Gokulchander, C. Grossner, and T. Radhakrishnan. Modeling rule base structure for expert system quality assurance. In *Thirteenth International Joint Conference on Artificial Intelligence: Workshop on Validation of knowledge-Based Systems*, August 1993.
- [PGGR93b] A. Preece, C. Grossner, P. Gokulchander, and T. Radhakrishnan. Structural validation of expert systems: Experience using a formal model. In *Eleventh National Conference on Artificial Intelligence (AAAI 93): Workshop on Validation and Verification of knowledge-Based Systems*, July 1993.
- [PGGR94a] A. Preece, C. Grossner, P. Gokulchander, and T. Radhakrishnan. Structural validation of expert systems: Experience using a formal model. In *World Congress on Expert Systems*, February 1994.
- [PGGR94b] A. Preece, C. Grossner, P. Gokulchander, and T. Radhakrishnan. Structural validation of rule-based systems using a structural model. *Submitted for review to IEEE Knowledge and Data Engineering*, February 1994.
- [Phi92] John L. Phillips, Jr. *How to Think About Statistics*. W. H. Freeman and Company, 1992.
- [PHR87] William J. Pardee and Barbara Hayes-Roth. Intelligent real-time control of material processing. Technical Report 1, Rockwell International Science Center, Palo Alto, CA, February 1987.
- [PIKL85] H. Van Dyke Parunak, Bruce W. Irish, James Kindrick, and Peter W. Lozo. Fractal actors for distributed manufacturing control. In *Proceedings of the Second IEEE Conference on AI Applications*, pages 653–660, December 1985.
- [Pit90a] K. Pitula. Cooperative distributed problem solving and the game of distributed blackbox. Master’s thesis, Concordia University, 1990.

- [Pit90b] Kristina Pitula. Cooperative problem solving through the game of distributed blackbox. Master's thesis, Concordia University, Montreal, Quebec, June 1990.
- [PL83] Jasmina Pavlin and Victor R. Lesser. Task allocation in distributed problem solving systems. Technical Report ?, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, May 1983.
- [PL95] M. V. Nagendra Prasad and Victor R. Lesser. Distributed case-based learning. work in progress, 1995.
- [PL96] M. V. Nagendra Prasad and Victor Lesser. Learning situation-specific coordination in generalized partial global planning. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [PLL95a] M. V. Nagendra Prasad, Victor Lesser, and Susan Lander. Learning experiments in a heterogeneous multi-agent system. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 59–64, 1995.
- [PLL95b] M. V. Nagendra Prasad, Victor R. Lesser, and Susan Lander. On reasoning and retrieval in distributed case bases. In *1995 IEEE International Conference on Systems, Man and Cybernetics*, Vancouver, Canada, October 1995.
- [PLL95c] M. V. Nagendra Prasad, Victor R. Lesser, and Susan Lander. Reasoning and retrieval in distributed case bases. *Journal of Visual Communication and Image Representation, Special Issue on Digital Libraries*, 1995. Also as UMASS CS Technical Report 95-27, 1995.
- [PLL95d] M. V. Nagendra Prasad, Victor R. Lesser, and Susan E. Lander. Learning organizational roles in a heterogeneous multi-agent system. Technical Report Tech Report 95-35, UMass Computer Science, 1995.
- [PLL96] M. V. Nagendra Prasad, Victor Lesser, and Susan Lander. Learning organizational roles in a heterogeneous multiagent system. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [PLS92] M. P. Papazoglou, S. C. Laufmann, and T. K. Sellis. An organizational framework for cooperating intelligent information systems. *Intelligent and Cooperative Information Systems*, 1(1):169–202, March 1992.
- [Pol90] M. E. Pollack. Plans as complex mental attitudes. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communications*. MIT Press, 1990.
- [Por91] Roger Porkess. *The Harper Collins Dictionary of Statistics*. Harper Collins, 1991.
- [PR90] Martha E. Pollack and Marc Ringuette. Introducing Tileworld: Experimentally evaluating agent architectures. In *Proceedings of the Eighth National Conference on Artificial Intelligence*, pages 183–189, Boston, July 1990.
- [PRG90a] K. Pitula, T. Radhakrishnan, and C. Grossner. Distributed blackbox : A testbed for distributed problem solving. In *IEEE Int'l Phoenix Conference on Computers and Communications*, pages 143–147, March 1990.
- [PRG90b] K. Pitula, T. Radhakrishnan, and C. Grossner. Distributed blackbox: A test bed for distributed problem solving. In *IEEE Int. Phoenix Conference on Computers and Communications*, March 1990.
- [Pri93] Michael Prietula, editor. *AI and theories of groups & organizations: Conceptual and Empirical Research*. AAAI Workshop, 1993. Workshop Working Notes.



- [PRK90] David W. Payton, J. Kenneth Rosenblatt, and David M. Keirse. Plan guided reaction. *IEEE Transactions on Systems, Man, and Cybernetics*, 20(6):1370–1382, November 1990.
- [Pru81] Dean G. Pruitt. *Negotiation Behavior*. Academic Press, 1981.
- [PS85] James L. Peterson and Abraham Silberschatz. *Operating System Concepts*. Addison Wesley, Reading, MA, 1985.
- [PS95] Douglas Pearson and Arun Sharma. Combining learning from instruction with recovery from incorrect knowledge. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [PS96] Henry Prakken and Giovanni Sartor. Argument-based extended logic programming with defeasible priorities. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [PT91a] Jeff Y-C Pan and Jay M. Tenenbaum. An intelligent agent framework for enterprise integration. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6):1391–1408, December 1991. (Special Issue on Distributed AI).
- [PT91b] Jeff Y-C Pan and Jay M. Tenenbaum. Toward an intelligent agent framework for enterprise integration. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 206–212, Anaheim, CA, July 1991.
- [PV83] Y. Paker and J.-P. Verjus, editors. *Distributed Computing Systems: Synchronization, Control, and Communication*. Academic Press, London, 1983.
- [Rai68] Howard Raiffa. *Decision Analysis: Introductory Lectures on Choice Under Uncertainty*. Addison-Wesley, Reading, Mass., 1968.
- [Rai82] H. Raiffa. *The Art and Science of Negotiation*. Harvard University Press, Cambridge, Mass., 1982.
- [Rai86] Malcolm Richard Railey. *Dynamic Control Structures For Cooperating Processes*. PhD thesis, Univ. of Illinois at Urbana-Champaign, January 1986.
- [Rao95] Anand S. Rao. Integrated agent architecture: Execution and recognition of mental-states. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [Rao96a] A. Rao. Operational semantics of BDI agents. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [Rao96b] Anand S. Rao. AgentSpeak(L): BDI agents speak out in a logical computable language. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [Rap89] A. Rapoport. Prisoner’s dilemma. In J. Eatwell, M. Milgate, and P. Newman, editors, *The New Palgrave: Game Theory*, pages 199–204. Macmillan, London, 1989.
- [Ras89] Eric Rasmusen. *Games and Information: An Introduction to Game Theory*. Basil Blackwell, 1989.
- [RAS96] Shounak Roychowdhury, Neeraj Arora, and Sandip Sen. Effects of local information on group behavior. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [RB89a] J. S. Rosenschein and J. S. Breese. Communication-free interactions among rational agents: A probabilistic approach. In L. Gasser and M. N. Huhns, editors, *Distributed Artificial Intelligence, Vol. II*. Pitman Publishing Ltd., 1989.

- [RB89b] Jeffrey S. Rosenschein and John S. Breese. Communication-free interactions among rational agents: A probabilistic approach. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*, pages 99–118. Pitman, 1989.
- [RD94] Simon Raik and Bohdan Durnota. The evolution of sporting strategies. In Russel J. Stonier and Xing Huo Yu, editors, *Complex Systems: Mechanisms of Adaption*, pages 85–92, Amsterdam, Netherlands, 1994. IOS Press. Presents a model for cooperation between members of teams.
- [Rea87] Rammohan K. Reagade. Metagames and metasystems. In John P. van Gigch, editor, *Decisionmaking about Decisionmaking*. Abacus Press, 1987.
- [Rea94] Leslie A. Real, editor. *Behavioral Mechanisms in Evolutionary Ecology*. The University of Chicago Press, 1994.
- [Reb83] Rene Reboh. Extracting useful advice from conflicting expertise. In *IJCAI*, pages 145–150, 1983.
- [Rey94] Craig W. Reynolds. Competition, coevolution and the game of tag. In *Artificial Life IV*. MIT Press, 1994.
- [RF90] William N. Robinson and Stephen Fickas. Negotiation freedoms for requirements engineering. Technical Report CIS-TR-90-04, Department of Computer and Information Science, University of Oregon, April 1990.
- [RFS90a] Michel Roboam, Mark Fox, and Katia Sycara. Enterprise management network architecture—distributed knowledge base support. Technical Report CMU-RI-TR-90-21, Carnegie-Mellon University, Pittsburgh, PA, 1990.
- [RFS90b] Michel Roboam, Mark Fox, and Katia Sycara. Enterprise management network architecture—the organization layer. Technical Report CMU-RI-TR-90-22, Carnegie-Mellon University, Pittsburgh, PA, 1990.
- [RG66] A. Rapoport and M. Guyer. A taxonomy of 2 x 2 games. In *Yearbook of the Society for General Systems Research XI*, pages 203–214. 1966.
- [RG84] Jeffrey S. Rosenschein and Michael R. Genesereth. Communication and cooperation. Technical Report HPP-84-5, Stanford Heuristic Programming Project, 1984.
- [RG85] Jeffrey S. Rosenschein and Michael R. Genesereth. Deals among rational agents. In *Proceedings of the Ninth International Joint Conference on Artificial Intelligence*, pages 91–99, Los Angeles, CA, August 1985. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 227–234, Morgan Kaufmann, 1988.).
- [RG87] Jeffrey S. Rosenschein and Michael R. Genesereth. Communication and cooperation among logic-based agents. In *Proceedings of the Sixth Phoenix Conference on Computers and Communications*, pages 594–600, Scottsdale, AZ, February 1987.
- [RG90] Mark Rokey and Sven Grenander. Planning for pace telerobotics : The remote mission specialist. *IEEE Expert*, 5(3):8–15, June 1990.
- [RG95] Anand S. Rao and Michael P. Georgeff. BDI agents: from theory to practice. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 312–319, San Francisco, CA, 1995. MIT Press.
- [RGG86] Jeffrey S. Rosenschein, Matthew L. Ginsberg, and Michael R. Genesereth. Cooperation without communication. In *Proceedings of the National Conference on Artificial Intelligence*, pages 51–57, Philadelphia, PA, August 1986.
- [Ric82] Elaine Rich. *Artificial Intelligence*. McGraw-Hill, 1982.

- [Ric83] Elaine Rich. *Artificial Intelligence*. McGraw Hill, New York, New York, 1983.
- [Ric86] James P. Rice. Poligon: A system for parallel problem solving. Technical Report KSL-86-19, Knowledge Systems Laboratory, Computer Science Department, Stanford University, Palo Alto, CA 94304, April 1986.
- [RL80] Scott Reed and Victor R. Lesser. Division of labor in honey bees and distributed focus of attention. Technical Report 80-17, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, September 1980.
- [RN95] Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Prentice Hall, 1995.
- [Rob91] William N. Robinson. A decision theoretic perspective of multiagent requirements negotiation. In *Proceedings of the AAAI Workshop on Cooperation Among Heterogeneous Intelligent Systems*, Anaheim, California, July 1991.
- [Rob93] Stephen P. Robbins. *Organizational Behavior: Concepts, Controversies, and Applications*. Prentice Hall, 1993.
- [Ros82] Jeffery S. Rosenschein. Synchronization of multi-agent plans. In *Proceedings of the National Conference on Artificial Intelligence*, pages 115–119, Pittsburgh, PA, August 1982. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 187–191, Morgan Kaufmann, 1988.).
- [Ros83] Stan Rosenschein. Reasoning about distributed action. *SIGART Newsletter*, 84:7, 1983.
- [Ros85] Jeffery Solomon Rosenschein. *Rational Interaction: Cooperation Among Intelligent Agents*. PhD thesis, Stanford University, October 1985.
- [Ros86] Jeffrey S. Rosenschein. *Rational Interaction Cooperation Among Intelligent Agents*. PhD thesis, Stanford Univ., 1986.
- [Ros88a] Jeffrey S. Rosenschein. The role of knowledge in logic-based rational interactions. In *Proceedings of the Seventh Phoenix Conference on Computers and Communications*, pages 497–504, Scottsdale, AZ, February 1988.
- [Ros88b] Sheldon Ross. *A First Course in Probability*. Macmillan, New York, 1988.
- [Ros90] David A. Rosenblitt. *Supporting Collaborative Planning: The Plan Integration Problem*. PhD thesis, Massachusetts Institute of Technology, Cambridge, MA 02139, 1990.
- [Ros95] Jeffrey S. Rosenschein. Multiagent planning as a social process: Voting, privacy, and manipulation. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 431, San Francisco, CA, 1995. MIT Press. (invited speaker talk).
- [RS81] C. Rieger and S. Small. Toward a theory of distributed word expert natural language parsing. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):43–51, January 1981.
- [RS83] Jeffrey S. Rosenschein and Vineet Singh. The utility of meta-level effort. Technical Report HPP-83-20, Stanford Heuristic Programming Project, Stanford University, March 1983.
- [RS84] Krithivasan Ramamritham and John A. Stankovic. Dynamic task scheduling in hard real-time distributed systems. *IEEE Software*, 1(3):65–75, July 1984.
- [RSC87] Steve Rowley, Howard Shrobe, and Robert Cassels. Joshua: Uniform access to heterogeneous knowledge structures or why joshing is better than conniving or planning. In *Proceedings of the National Conference on Artificial Intelligence*, pages 48–52, Seattle, WA, July 1987.

- [RSS90] Krithi Ramamritham, John A. Stankovic, and Perng-Fei Shiah. Efficient scheduling algorithms for real-time multiprocessor systems. *IEEE Transactions on Parallel and Distributed Systems*, 1(2):184–195, April 1990.
- [RSZ89] Krithi Ramamritham, John A. Stankovic, and Wei Zhao. Distributed scheduling of tasks with deadlines and resource requirements. *IEEE Transactions on Computers*, 38(8):1110–1123, August 1989.
- [Rum76] D. E. Rumelhart. Toward an interactive model of reading. Technical Report 56, Center for Human Information Processing, University of California, San Diego, California, 1976.
- [Rus85] Stuart Russell. The compleat guide to MRS. Technical Report KSL No. 85-12, Knowledge Systems Laboratory, Computer Science Department, Stanford University, Stanford, CA 94305, June 1985.
- [Rya95] Conor Ryan. GPRobots and GPTeams - competition, co-evolution and co-operation in genetic programming. In E. S. Siegel and J. R. Koza, editors, *Working Notes for the AAAI Symposium on Genetic Programming*, pages 86–93, Menlo Park, CA, 10–12 November 1995. AAAI.
- [Rya96] Mark Ryan. Feature oriented specifications. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [RZ94] J. S. Rosenschein and G. Zlotkin. Designing conventions for automated negotiation. *AI Magazine*, pages 29–46, Fall 1994.
- [SAB<sup>+</sup>82] Mark Stefik, Jan Aikins, Robert Balzer, John Beniot, Lawrence Birnbaum, Frederick Hayes-Roth, and Earl Sacerdoti. The organization of expert systems, A tutorial. *Artificial Intelligence*, 18:135–172, 1982.
- [Sac74] Earl D. Sacerdoti. Planning in a hierarchy of abstraction spaces. *Artificial Intelligence*, 5:115–135, 1974.
- [Sac75] Earl D. Sacerdoti. The nonlinear nature of plans. In *Proceedings of the Fourth International Joint Conference on Artificial Intelligence*, pages 206–214, Tbilisi, Georgia, USSR, August 1975.
- [Sac77] Earl D. Sacerdoti. *A Structure for Plans and Behavior*. Elsevier, 1977.
- [Sac78] Earl D. Sacerdoti. What language understanding research suggests about Distributed Artificial Intelligence. In *Proceedings of the Distributed Sensor Networks Workshop*, pages 8–11, 1978. Published by the Department of Computer Science, Carnegie-Mellon University, Pittsburgh, Pennsylvania.
- [Sac79] Earl D. Sacerdoti. Problem solving tactics. In *Proceedings of the Sixth International Joint Conference on Artificial Intelligence*, pages 1077–1085, Tbilisi, Georgia, USSR, August 1979.
- [Sac80] Earl D. Sacerdoti. Problem solving tactics. *AI Magazine*, pages 7–15, 1980.
- [Sak88] Markku Sakkinen. On the darker side of C++. In *Proceedings of European Conference on Object-Oriented Programming 1988*, pages 162–176, 1988.
- [SAL95] J. Alfredo Sánchez, Flávio S. Azevedo, and John J. Leggett. PARAgente: Exploring the issues in agent-based user interfaces. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 320–327, San Francisco, CA, 1995. MIT Press.
- [San93a] Toumas Sandholm. An implementation of the contract net protocol based on marginal cost calculation. In *Proceedings of the National Conference on Artificial Intelligence*, pages 256–262, 1993.

- [San93b] Tuomas Sandholm. An implementation of the contract net protocol based on marginal cost calculations. In *Proceedings of the Eleventh National Conference on Artificial Intelligence*, pages 256–262, Washington, July 1993.
- [SB87a] David Smith and Martin Broadwell. Plan coordination in support of expert systems. In *Proceedings of the DARPA Knowledge-based Planning Workshop*, Austin, Texas, December 1987.
- [SB87b] David Smith and Martin Broadwell. Plan coordination in support of expert systems integration. In *Proceedings of the DARPA Knowledge-Based Planning Workshop*, pages 12.1–12.6, December 1987.
- [SB88] David Smith and Martin Broadwell. The pilot’s associate—an overview. In *Proceedings of the SAE Aerotech Conference*, Los Angeles, CA, May 1988.
- [SB95] Weiming Shen and Jean-Paul Barthes. DIDE: A multi-agent environment for engineering design. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 344–351, San Francisco, CA, 1995. MIT Press.
- [SBM95] Barry G. Silverman, Nabil Bedewi, and Alfredo Morales. Intelligent agents in software reuse repositories. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM ’95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [SC85a] Katia Sycara-Cyranski. Arguments of persuasion in labor mediation. In *Proceedings of IJCAI-85*, pages 294–296, Los Angeles, CA, 1985.
- [SC85b] Katia Sycara-Cyranski. Arguments of persuasion in labour mediation. In *Proceedings of the Ninth International Joint Conference on Artificial Intelligence*, pages 294–296, Los Angeles, CA, August 1985.
- [SC85c] Katia Sycara-Cyranski. Persuasive argumentation in resolution of collective bargaining impasses. In *Proceedings of the Seventh Annual Conference of the Cognitive Science Society*, pages 356–360, Irvine, CA, 1985.
- [SC95a] Tuomas W. Sandholm and Robert H. Crites. On multiagent Q-learning in a semi-competitive domain. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 71–77, 1995.
- [SC95b] Craig Van Slyke and Rosann Webb Collins. Trust between users and their intelligent agents. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM ’95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [SC95c] Ira A. Smith and Philip R. Cohen. Toward a semantics for a speech act based agent communications language. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM ’95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [SC95d] Lorenzo Sommaruga and Nadia Catenazzi. From practice to theory in designing autonomous agents. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [SC95e] Lorenzo Sommaruga and Nadia Catenazzi. Some principles in designing cooperative architectures of autonomous agents. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 461, San Francisco, CA, 1995. MIT Press. (poster).
- [Sch60] Thomas C. Schelling. *The Strategy of Conflict*. Harvard University Press, Cambridge, MA, 1960.
- [Sch86] Eric Schoen. The CAOS system. Technical Report STAN-CS-86-1125, Computer Science Department, Stanford University, Stanford, CA 94305, March 1986.
- [Sch87a] Roger C. Schank. What is AI, anyway? *AI Magazine*, 8(4):59–65, Winter 1987.

- [Sch87b] M. J. Schoppers. Universal plans for reactive robots in unpredictable environments. In *Proceedings of the Tenth International Joint Conference on Artificial Intelligence*, pages 1039–1046, Milan, Italy, August 1987.
- [Sch96a] Jurgen Schmidhuber. Theory of multi-agent learning in unrestricted environments. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [Sch96b] P-Y. Schobbens. A comparative logic for preferences. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [SCHR<sup>+</sup>88a] R. Steeb, S. Cammarata, F. Hayes-Roth, P. Thorndyke, and R. Wesson. Architectures for distributed air-traffic control. In Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 90–101. Morgan Kaufman, 1988.
- [SCHR<sup>+</sup>88b] R. Steeb, S. Cammarata, F. A. Hayes-Roth, P. W. Thorndyke, and R. B. Wesson. Distributed intelligence for air fleet control. In Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 90–101. Morgan Kaufmann, 1988.
- [SCN<sup>+</sup>86] R. Steeb, S. Cammarata, S. Narain, J. Rothenburg, and W. Giarla. Cooperative intelligence for remotely piloted vehicle fleet control. Technical Report R-3408-ARPA, Rand Corporation, October 1986.
- [Sco87] W. Richard Scott. *Organizations: Rational, Natural, and Open Systems*. Prentice-Hall, Inc., Englewood Cliffs, NJ, 1987.
- [SD81] Reid G. Smith and Randall Davis. Frameworks for cooperation in distributed problem solving. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):61–70, January 1981. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 61–70, Morgan Kaufmann, 1988.).
- [SD91] Sandip Sen and Edmund H. Durfee. A formal study of distributed meeting scheduling: Preliminary results. In *Proceedings of the ACM Conference on Organizational Computing Systems '91*, pages 55–68, Atlanta, GA, November 1991.
- [SD92a] Sandip Sen and Edmund H. Durfee. Automated meeting scheduling among heterogeneous agents. In *Working Papers of the AAAI-92 Workshop on Cooperation among Heterogeneous Intelligent Agents*, pages 116–120, July 1992.
- [SD92b] Sandip Sen and Edmund H. Durfee. A formal analysis of communication and commitment in distributed meeting scheduling. In *Working Papers of the 11th International Workshop on Distributed Artificial Intelligence*, pages 333–342, February 1992.
- [SD92c] Y.-P. So and E. H. Durfee. A distributed problem solving infrastructure for computer network management. *International Journal of Intelligent and Cooperative Information Systems*, 1(2):363–392, June 1992.
- [SD93a] Sandip Sen and Edmund H. Durfee. Avoiding and resolving conflicts among cooperative scheduling agents. In *Proc. of IJCAI-93 Workshop on Computational Models of Conflict Management in Cooperative Problem Solving*, August 1993.
- [SD93b] Sandip Sen and Edmund H. Durfee. Dependent subtask processing in a contract-net for manufacturing. In *Proc. of AAAI-93 Workshop on Intelligent Manufacturing Technology*, pages 7–13, July 1993.
- [SD93c] Sandip Sen and Edmund H. Durfee. The effects of search bias on flexibility in distributed scheduling. In *Working Papers of the 12th International Workshop on Distributed Artificial Intelligence*, pages 321–334, May 1993.

- [SD93d] Sandip Sen and Edmund H. Durfee. Using temporal abstractions and cancellations for efficiency in automated meeting scheduling. In *Proc. of the International Conference on Intelligent and Cooperative Information Systems*, pages 163–172, May 1993.
- [SD94a] Sandip Sen and Edmund H. Durfee. Adaptive surrogate agents. In *Working Papers of the 13th International Workshop on Distributed Artificial Intelligence*, pages 320–333, July 1994.
- [SD94b] Sandip Sen and Edmund H. Durfee. On the design of an adaptive meeting scheduler. In *Proc. of the Tenth IEEE Conference on AI Applications*, pages 40–46, March 1994.
- [SD94c] Sandip Sen and Edmund H. Durfee. The role of commitment in cooperative negotiation. *International Journal of Intelligent and Cooperative Information Systems*, 3(1):67–81, 1994.
- [SD95a] Sandip Sen and Edmund H. Durfee. Unsupervised surrogate agents and search bias change in flexible distributed scheduling. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 336–343, San Francisco, CA, 1995. MIT Press.
- [SD95b] Jaime Simão Sichman and Yves Demazeau. Exploiting social reasoning to deal with agency level inconsistency. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 352–359, San Francisco, CA, 1995. MIT Press.
- [SDKM86] J. Siegel, V. Dubrovsky, S. Kiesler, and T. McGuire. Group processes in computer-mediated communication. *Organizational Behavior and Human Decision Processes*, 37:157–187, 1986.
- [SDL<sup>+</sup>88] William Swartout, Thomas Dean, Theodore Linden, Richard Fikes, Peter Cheeseman, and Drew McDermott. DARPA Santa Cruz workshop on planning. *AI Magazine*, pages 115–131, 1988.
- [SDSC95] Carla Simone, Monica Divitini, Kjeld Schmidt, and Peter Carstensen. A multi-agent approach to the design of coordination mechanisms. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 460, San Francisco, CA, 1995. MIT Press. (poster).
- [SE87] Kang G. Shin and Mark E. Epstein. Intertask communications in an integrated multirobot system. *IEEE Journal of Robotics and Automation*, RA3(2):90–100, April 1987.
- [Sea69] J. R. Searle. *Speech Acts: An Essay in the Philosophy of Language*. Cambridge University Press, Cambridge, 1969.
- [Sea70] John R. Searle. *Speech Acts: An Essay in the Philosophy of Language*. Cambridge University Press, 1970.
- [Sei85] Charles L. Seitz. The cosmic cube. *Communications of the ACM*, pages 22–33, 1985.
- [Sek94] Mahendra Sekaran. Reinforcement learning based agent coordination. Master’s thesis, University of Tulsa, Tulsa, OK., April 1994.
- [Sen93] Sandip Sen. *Predicting Tradeoffs in Contract-Based Distributed Scheduling*. PhD thesis, University of Michigan, December 1993.
- [Sen95] Sandip Sen, editor. *Working Notes of the IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, Montréal Quebec, August 1995.
- [SF88] Baldev Singh and Ira R. Forman. Coordination systems: Manual interactions. Technical Report STP-229-88, Microelectronics and Computer Technology Corporation, Austin, TX, September 1988.

- [SF89a] N. Sadeh and M. S. Fox. Preference propagation in temporal/capacity constraint graphs. Technical report CMU-RI-TR-89-2, Robotics Institute, Carnegie Mellon University, January 1989.
- [SF89b] Arvind Sathi and Mark S. Fox. Constraint-directed negotiation of resource reallocations. In Michael N. Huhns and Les Gasser, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [SFB<sup>+</sup>87] M. Stefik, G. Foster, D. G. Bobrow, K. Kahn, S. Lanning, and L. Suchman. Beyond the chalkboard: Using computers to support collaboration and problem solving in meetings. *Communications of the ACM*, 30:32–47, 1987.
- [SG90] Michael J Shaw and James A. Gentry. Inductive learning for risk classification. *IEEE Expert*, 5(1):47–53, February 1990.
- [SGP78] E. Spangler, M. A. Gordon, and R. M. Pipkin. Token women: An empirical test of kanter’s hypothesis. *American Journal of Sociology*, 85:160–170, 1978.
- [SH87a] S. F. Smith and J. E. Hynynen. Integrated decentralization of production management: An approach for factory scheduling. In *Proceedings ASME Annual Winter Conference: Symposium on Integrated and Intelligent Manufacturing*, Boston, MA, December 1987.
- [SH87b] Stephen F. Smith and Juha E. Hynynen. Integrated decentralization of production management: An approach for factory scheduling. In C. R. Liu, A. Requicha, and S. Chandrasekar, editors, *Intelligent and Integrated Manufacturing Analysis and Synthesis*, pages 427–439. The American Society of Mechanical Engineers, New York, 1987.
- [Sha88] F. M. Shakun. *Evolutionary Systems Design: Policy Making Under Complexity and Group Decision Support Systems*. Holden-Day, Oakland, CA., 1988.
- [Sha96] Michael J. Shaw. Cooperative problem solving and learning in multi-agent information systems. *International Journal of Computational Intelligence and Organizations*, 1996.
- [Shi84a] S. B. Shimanoff. Coordinating group interaction via communication rules. In L. A. Samovar R. S. Cathcart, editor, *Small Group Communication*. Wm. C. Brown, Dubuque, Iowa, 1984.
- [Shi84b] S. B. Shimanoff. *Small Group Communication*, chapter Coordinating Group Interaction via Communication Rules. Wm. C. Brown, Dubuque, Iowa, 1984.
- [Sho76] E. H. Shortliffe. *Computer-Based Medical Consultations: MYCIN*. American Elsevier, 1976.
- [Sho91] Yoav Shoham. AGENT0: A simple agent language and its interpreter. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 704–709, Anaheim, July 1991.
- [Sho93] Yoav Shoham. Agent-oriented programming. *Artificial Intelligence*, 60:51–92, 1993.
- [Shu82] Martin Shubik. *Game Theory in the Social Sciences: Concepts and Solutions*. MIT Press, 1982.
- [Sia91] S. Sian. Adaptation based on cooperative learning in multi-agent systems. In Y. Demazeau and J.-P. Muller, editors, *Decentralize AI*, volume 2, pages 257–272. Elsevier Science Publications, 1991.
- [Sig95] Karl Sigmund, editor. *Games of Life: Explorations in Ecology, Evolution and Behaviour*. Penguin Books, 1995.
- [Sim57] Herbert A. Simon. *Models of Man*. John Wiley & Sons, New York, 1957.
- [Sim68] Herbert A. Simon. *The Sciences of the Artificial*. MIT Press, Cambridge, MA, 1968.



- [Sim69] Herbert A. Simon. *The Sciences of the Artificial*. MIT Press, 1969.
- [Sim71] W. E. Simkin. *Mediation and the Dynamics of Collective Bargaining*. The Bureau of National Affairs, Washington, D.C., 1971.
- [Sim73] Herbert A. Simon. The structure of ill-structured problems. *Artificial Intelligence*, 4:181–201, 1973.
- [Sim78] Herbert A. Simon. Information-processing theory of human problem solving. In W. K. Estes, editor, *Human Information Processing*, pages 271–295. Lawrence Erlbaum Associates, Hillsdale, NJ, 1978. Handbook of Learning and Cognition, Volume 5.
- [Sin90] Munindar P. Singh. The effect of agent control strategy on the performance of a DAI pursuit problem. In *Working Papers of the 10th International Workshop on Distributed Artificial Intelligence*, October 1990.
- [Sin91] Munindar Singh. Towards a formal theory of communication for multiagent systems. In *Proceedings of the Twelfth International Joint Conference on Artificial Intelligence*, August 1991.
- [SJ93] J. W. Smith and T. R. Johnson. A stratified approach to specifying, designing, and building knowledge systems. *IEEE Expert*, 8(3):15–25, June 1993.
- [SK95] Onn Shehory and Sarit Kraus. Task allocation via coalition formation among autonomous agents. In Chris S. Mellish, editor, *Proceedings of the Fourteenth International Joint Conference on Artificial Intelligence*, pages 655–661, Montréal, Canada, August 1995.
- [SKY89] K. Sugihara, T. Kikuno, and N. Yoshida. A meeting scheduler for office automation. *IEEE Transactions on Software Engineering*, 15(10):1141–1146, October 1989.
- [SL93a] T. Sugawara and V. Lesser. On-line learning of coordination plans. In *Working Papers of the 12th International Workshop on Distributed Artificial Intelligence*, May 1993.
- [SL93b] Toshiharu Sugawara and Victor R. Lesser. On-line learning of coordination plans. Computer Science Technical Report 93–27, University of Massachusetts, 1993.
- [SL95a] Tuomas Sandholm and Victor Lesser. Issues in automated negotiation and electronic commerce: Extending the contract net framework. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 328–335, San Francisco, CA, 1995. MIT Press.
- [SL95b] Tuomas W. Sandholm and Victor R. Lesser. Coalition formation among bounded rational agents. In Chris S. Mellish, editor, *Proceedings of the Fourteenth International Joint Conference on Artificial Intelligence*, pages 662–669, Montréal, Canada, August 1995.
- [SL95c] Tuomas W. Sandholm and Victor R. Lesser. Equilibrium analysis of the possibilities of unenforced exchange in multiagent systems. In Chris S. Mellish, editor, *Proceedings of the Fourteenth International Joint Conference on Artificial Intelligence*, pages 694–701, Montréal, Canada, August 1995.
- [SL95d] Toshiharu Sugawara and Victor Lesser. Learning coordination plans in distributed OS environments. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 462, San Francisco, CA, 1995. MIT Press. (poster).
- [Sla87] Stephen Slade. *The T Programming Language: A Dialect of Lisp*. Prentice Hall, 1987.
- [Sla91] Stephen Slade. Case-based reasoning: A research paradigm. *AI Magazine*, 12(1):42–55, Spring 1991.

- [SM89] Larry M. Stephens and Matthias B. Merx. Agent organization as an effector of DAI system performance. In *Working Papers of the 9th International Workshop on Distributed Artificial Intelligence*, September 1989.
- [SM90] Larry M. Stephens and Matthias B. Merx. The effect of agent control strategy on the performance of a DAI pursuit problem. In *Proceedings of the 1990 Distributed AI Workshop*, October 1990.
- [SM95a] Jude Shavlik and Richard Maclin. Learning from instruction and experience in competitive situations. In Diana Gordon, editor, *Working Notes of the ICML '95 Workshop on Agents that Learn from Other Agents*, Tahoe City, CA, 1995.
- [SM95b] H. Sorensen and M. McElligott. PSUN: A profiling system for usenet news. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Smi78] Reid Garfield Smith. *A Framework for Problem Solving in a Distributed Processing Environment*. PhD thesis, Stanford University, December 1978. A revised version was published by UMI Research Press.
- [Smi79] Reid G. Smith. A framework for distributed problem solving. In *Proceedings of the Sixth International Joint Conference on Artificial Intelligence*, pages 836–841, Tbilisi, Georgia, USSR, August 1979.
- [Smi80] Reid G. Smith. The contract net protocol: High-level communication and control in a distributed problem solver. *IEEE Transactions on Computers*, C-29(12):1104–1113, December 1980.
- [Smi84] J. M. Smith. The evolution of animal intelligence. In C. Hookway, editor, *Minds, Machines and Evolution*. Cambridge University Press, 1984.
- [Smi85] Reid G. Smith. Report on the 1984 distributed artificial intelligence workshop. *AI Magazine*, 6(3):234–243, Fall 1985.
- [Smi87] S. F. Smith. A constraint-based framework for reactive management of factory schedules. In M. Oliff, editor, *Expert Systems and the Leading Edge in Production Planning and Control*. Benjamin Cummings Publishers, 1987.
- [Smi89] D. E. Smith. Controlling backward inference. *Artificial Intelligence*, 39(2):145–208, June 1989.
- [Smi91] Mark Smith. Using massively-parallel supercomputers to model stochastic spatial predator-prey systems. Technical Report EPCC-TR91-06, April 17 1991.
- [SMR86] Arvind Sathi, Thomas E. Morton, and Steven F. Roth. Callisto: An intelligent project management system. *AI Magazine*, 7(5):34–52, 1986.
- [Sny91] M. A. Snyder. Reply: A commentary on the paper by Jain and Binford. *CVGIP: Graphic Models and Image Understanding*, 53(1):118–119, January 1991.
- [SO85] S. F. Smith and P. S. Ow. The use of multiple problem decompositions in time-constrained planning tasks. In *Proceedings 9th International Joint Conference on Artificial Intelligence*, Los Angeles, California, August 1985.
- [SOL<sup>+</sup>86] S. F. Smith, P. S. Ow, C. LePape, B. McLaren, and N. Muscettola. Integrating multiple scheduling perspectives to generate detailed production plans. In *Proceedings SME Conference on AI in Manufacturing*, Long Beach, CA, September 1986.
- [Spr77] Robert F. Sproull. *Strategy Construction using a Synthesis of Heuristic and Decision-theoretic Methods*. PhD thesis, Stanford University, Stanford, CA 94305, July 1977. (Also published as Technical Report CSL-77-2, Xerox Palo Alto Research Center, Palo Alto, California, July, 1977.).

- [SR77] Elliot M. Soloway and Edward M. Riseman. Levels of pattern description in learning. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, pages 801–811, Tbilisi, Georgia, USSR, August 1977.
- [SR87a] J. Stankovic and K. Ramamritham. *Tutorial on Hard Real-Time Systems*. IEEE Computer Society Press, 1987.
- [SR87b] J. A. Stankovic and K. Ramamritham. The design of the spring kernel. In *Proceedings of the 1987 Real-time Systems Symposium*, December 1987.
- [SR90] John A. Stankovic and Krithi Ramamritham. Editorial: What is predictability for real-time systems? *The Journal of Real-Time Systems*, 2:247–254, 1990.
- [SRC85] John A. Stankovic, Krithivasan Ramamritham, and Sheng-Chang Cheng. Evaluation of a flexible task scheduling algorithm for distributed hard real-time systems. *IEEE Transactions on Computers*, C-34(12):1130–1143, December 1985.
- [Sri87] N. S. Sridharan. Report on the 1986 workshop on distributed artificial intelligence. *AI Magazine*, 8(3):75–85, Fall 1987.
- [SRN90] T. K. Sellis, N. Roussopoulos, and R. T. Ng. Efficient compilation of large rule bases using logical access paths. *Information Systems*, 15(1):73–84, 1990.
- [SRS93] Chia Shen, Krithi Ramamritham, and John A. Stankovic. Resource reclaiming in multiprocessor real-time systems. *IEEE Transactions on Parallel and Distributed Systems*, 4(4):382–398, April 1993.
- [SRSF90a] Katia Sycara, Steve Roth, Norman Sadeh, and Mark Fox. An investigation into distributed constraint-directed factory scheduling. In *Proceedings of the Sixth IEEE Conference on Artificial Intelligence Applications*, 1990.
- [SRSF90b] Katia Sycara, Steve Roth, Norman Sadeh, and Mark Fox. Managing resource allocation in multi-agent time-constrained domains. In *Proceedings of the 1990 DARPA Planning Workshop*, pages 240–250, November 1990.
- [SRSF91] Katia Sycara, Steve Roth, Norman Sadeh, and Mark Fox. Distributed constrained heuristic search. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6):1446–1461, December 1991. (Special Issue on Distributed AI).
- [SS84] David Stemple and Tim Sheard. Specification and verification of abstract database types. In *Proceedings of the 3rd ACM SIGACT-SIGMOD Symposium on Principles of Database Systems*, pages 248–257, Waterloo, Ontario, April 1984.
- [SS86] Prakash P. Shenoy and Glenn Shafer. Propagating belief functions with local computations. *IEEE Expert*, 1986.
- [SS94a] Mahendra Sekaran and Sandip Sen. Learning with friends and foes. In *Sixteenth Annual Conference of the Cognitive Science Society*, pages 800–805, 1994.
- [SS94b] Mahendra Sekaran and Sandip Sen. To help or not to help. In *Proceedings of the Seventeenth Annual Conference of the Cognitive Science Society*, pages 736–741, Pittsburgh, PA, July 1994.
- [SS95a] Sandip Sen and Mahendra Sekaran. Multiagent coordination with learning classifier systems. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 84–89, 1995.
- [SS95b] Sandip Sen and Mahendra Sekaran. Reciprocity: A foundational principle for emergent cooperative behavior, 1995. submitted to the Fourteenth International Joint Conference on Artificial Intelligence.

- [SS95c] Sandip Sen and Mahendra Sekaran. Using reciprocity to adapt to others. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 90–94, 1995.
- [SS95d] John W. Sheppard and Steven L. Salzberg. Combining genetic algorithms with memory based reasoning. In Larry Eshelman, editor, *Proceedings of the Sixth International Conference on Genetic Algorithms*, pages 452–459, San Francisco, CA, 1995. Morgan Kaufmann Publishers, Inc.
- [SSC<sup>+</sup>95] Len Seligman, Kenneth P. Smith, Patricia L. Carbone, Cristopher Elsaesser, Paul E. Lehner, Michael S. V. Turner, and April Milner. Intelligent monitoring of distributed, heterogeneous data sources. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [SSH94] Sandip Sen, Mahendra Sekaran, and John Hale. Learning to coordinate without sharing information. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 426–431, Seattle, WA, 1994.
- [SST95] Andrea Schaerf, Yoav Shoham, and Moshe Tennenholtz. Adaptive load balancing: a study in co-learning. In Sandip Sen, editor, *IJCAI-95 Workshop on Adaptation and Learning in Multiagent Systems*, pages 78–83, 1995.
- [ST92] Yoav Shoham and Moshe Tennenholtz. On the synthesis of useful social laws for artificial agent societies (preliminary report). In *Proceedings of the National Conference on Artificial Intelligence*, pages 276–281, San Jose, CA, July 1992.
- [Sta82] Lawrence A. Stabile. Frame-based computer network monitoring. In *AAAI*, pages 327–330, 1982.
- [Sta84a] J. Stankovic. A perspective on distributed computer systems. *IEEE Transactions on Computers*, C-33(12), December 1984.
- [Sta84b] J. A. Stankovic. A perspective on distributed computing systems. *IEEE Transactions on Computers*, C-33(12):1102–1115, December 1984.
- [Sta84c] J. A. Stankovic. Simulations of three adaptive, decentralized controlled, job scheduling algorithms. *Computer Networks*, 8:199–217, 1984.
- [Sta85] J. A. Stankovic. An application of bayesian decision theory to decentralized control of job scheduling. *IEEE Transactions on Computers*, C-34(2):117–130, February 1985.
- [Sta86] Sharon A. Stansfield. ANGY: A rule-based expert system for automatic segmentation of coronary vessels from digital subtracted angiograms. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PAMI-8(2):188–199, March 1986.
- [Sta87] Susan Leigh Star. The structures of ill-structured solutions: Boundary objects and heterogeneous distributed problem solving. In Michael N. Huhns, editor, *Distributed Artificial Intelligence*, volume 1 of *Research Notes in Artificial Intelligence*. Pitman, 1987.
- [Sta89] Susan Leigh Star. The Structure of Ill-Structured Solutions: Boundary Objects and Heterogeneous Distributed Problem Solving. In L. Gasser and M. N. Huhns, editors, *Distributed Artificial Intelligence, Volume II*, pages 37–54. Pitman/Morgan Kaufmann, London, 1989.
- [Ste80] Mark Jeffrey Stefik. *Planning with Constraints*. PhD thesis, Stanford University, 1980.
- [Ste81a] Mark Stefik. Planning and meta-planning (MOLGEN: Part 2). *Artificial Intelligence*, 16:141–170, 1981.
- [Ste81b] Mark Stefik. Planning with constraints (MOLGEN: Part 1). *Artificial Intelligence*, 16:111–140, 1981.

- [Ste84] Guy L. Steele Jr. *Common Lisp: The language*. Digital Press, 1984.
- [Ste90] Luc Steels. Components of expertise. *AI Magazine*, 11(2):29–49, Summer 1990.
- [Ste92] W. Richard Stevens. *Advanced Programming in the UNIX Environment*. Addison-Wesley, 1992.
- [Ste95a] D. Steiner. IMAGINE. In G. O’Hare and N. Jennings, editors, *Foundations of Distributed Artificial Intelligence*, chapter 13. Wiley Inter-Science, 1995. Forthcoming.
- [Ste95b] Kenneth R. Stephens. Behavior analysis and intelligent agent theory. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM ’95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Sti87] Arthur L. Stinchcombe. *Constructing Social Theories*. University of Chicago Press, Chicago, 1987.
- [Sti90] Arthur L. Stinchcombe. *Information and Organizations*. University of California Press, Berkeley, CA, 1990.
- [Str78] A. Strauss. *Negotiations: Varieties, Processes, Contexts, and Social Order*. Jossey Bass, San Francisco, 1978.
- [Str87] B. Stroustrup. *The C++ Programming Language*. Addison Wesley, 1987.
- [Str92] David Straker. *C Style: Standards and Guidelines*. Prentice Hall, 1992.
- [Stu85] Christopher J. Stuart. An implemenation of a multi-agent plan synchronizer. In *Proceedings of the Ninth International Joint Conference on Artificial Intelligence*, pages 1031–1033, Los Angeles, CA, August 1985. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 216–219, Morgan Kaufmann, 1988.).
- [Sug95] Toshiharu Sugawara. Reusing past plans in distributed planning. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 360–367, San Francisco, CA, 1995. MIT Press.
- [Sus75] Gerald Jay Sussman. *A Computer Model of Skill Acquisition*. American Elsevier, 1975.
- [Suz87] Kiyoshi Suzaki. *The New Manufacturing Challenge*. The Free Press, 1987.
- [SV96] Peter Stone and Manuela Veloso. Learning: A case study in robotic soccer. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [SW83] Ron Sauers and Rick Walsh. On the requirements of future expert systems. In *IJCAI*, pages 110–115, 1983.
- [SW88] M. Shaw and A. B. Whinston. A distributed knowledge-based approach to flexible automation: The contract-net framework. *International Journal of Flexible Manufacturing Systems*, 1(1):85–104, 1988.
- [SW89] Michael J. Shaw and Andrew B. Whinston. Learning and adaptation in distributed artificial intelligence systems. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [Swa89] W. Swartout. DARPA workshop on planning. *AI Magazine*, 9(2):101–112, 1989.
- [Swa91] W. Swartout. Empirical study of expert system development. *Knowledge-Based Systems*, 4(4):101–112, December 1991.

- [Syc87a] Ekaterini P. Sycara. *Resolving Adversarial Conflicts: An Approach Integrating Case-Based and Analytic Methods*. PhD thesis, Georgia Institute of Technology, Atlanta, GA, June 1987.
- [Syc87b] Katia Sycara. Finding creative solutions in adversarial impasses. In *Proceedings of the Ninth Annual Conference of the Cognitive Science Society*, pages 997–1003, Seattle, WA., 1987.
- [Syc87c] Katia Sycara. Planning for negotiation: A case-based approach. In *DARPA Knowledge-Based Planning Workshop*, pages 11.1–11.10, December 1987.
- [Syc88a] Katia Sycara. Resolving goal conflicts via negotiation. In *Proceedings of the National Conference on Artificial Intelligence*, pages 245–250, August 1988.
- [Syc88b] Katia Sycara. Utility theory in conflict resolution. *Annals of Operations Research*, 12:65–84, 1988.
- [Syc88c] Katia P. Sycara. Multi-agent compromise via negotiation. In *Proceedings of the 1988 Distributed AI Workshop*, May 1988.
- [Syc89a] Katia Sycara. Argumentation: Planning other agents’ plans. In *Proceedings of IJCAI-89*, Detroit, MI, 1989.
- [Syc89b] Katia Sycara. Multi-agent compromise via negotiation. In M. Huhns and L. Gasser, editors, *Distributed Artificial Intelligence, Volume II*. Pittman Publishing Ltd and Morgan Kaufmann, 1989.
- [Syc89c] Katia Sycara. Negotiation planning: An AI approach. *European Journal of Operational Research*, 1989.
- [Syc89d] Katia P. Sycara. Multiagent compromise via negotiation. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [Syc95] Katia P. Sycara. Coordination of multiple intelligent agents in the infosphere. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 432, San Francisco, CA, 1995. MIT Press. (invited speaker talk).
- [Tam95] Milind Tambe. Agent and agent-group tracking in a real-time dynamic environment. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 368–375, San Francisco, CA, 1995. MIT Press.
- [Tan87] Steven L. Tanimoto. *The Elements of Artificial Intelligence*. Computer Science Press, Rockville, Maryland, 1987.
- [Tan93] Ming Tan. Multi-agent reinforcement learning: Independent vs. cooperative agents. In *Proceedings of the Tenth International Conference on Machine Learning*, pages 330–337, June 1993.
- [Tat75] Austin Tate. Interacting goals and their use. In *Proceedings of the Fourth International Joint Conference on Artificial Intelligence*, pages 215–218, August 1975.
- [Tat76] Austin Tate. NONLIN: A hierarchical non-linear planner. Technical report, Department of Artificial Intelligence, University of Edinburgh, 1976.
- [Tat77] Austin Tate. Generating project networks. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, pages 888–893, Tbilisi, Georgia, USSR, August 1977.
- [Tat84] Austin Tate. Goal structure — capturing the intent of plans. In T. O’Shea, editor, *ECAI-84: Proceedings of the Sixth European Conference on Artificial Intelligence*. Elsevier, 1984.

- [Tat87] Deborah G. Tatar. *A Programmer's Guide to Common Lisp*. Digital Press, 1987.
- [Tau93] Eric Taub. Sharing schedules. *MacUser*, pages 155–162, July 1993.
- [TBC<sup>+</sup>88] R. N. Taylor, F. C. Belz, L. A. Clarke, L. Osterweil, W. W. Selby, J. C. Wileden, A. L. Wolfe, and M. Young. Foundations for the Arcadia environment architecture. In *Proceedings of the ACM SIGSOFT/SIGPLAN Software Engineering Symposium on Practical Software Development Environments*, pages 1–12, Boston, MA, November 1988.
- [TCS96] Can Tuerker, Stefan Conrad, and Gunter Saake. Dynamically changing behaviour. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [Ten79] Robert R. Tenny. *Distributed Decision Making using a Distributed Model*. PhD thesis, Massachusetts Institute of Technology, Cambridge, MA, June 1979. (Also published as Technical Report LIDS-TD-938, Laboratory for Information and Decision Systems, Massachusetts Institute of Technology, Cambridge, Massachusetts, 02139, June 1979.).
- [Ter90] Kornel Terplan. Integrated network management. In Manu Malek and Mark Wall, editors, *Network Management and Control*, pages 31–57. Plenum Press, New York, 1990.
- [TFM95] Chelliah Thirunavukkarasu, Tim Finin, and James Mayfield. Secret agents - a security architecture for KQML. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [TGP89] David Taenzer, Murthy Ganti, and Sunil Podar. Object-oriented software reuse : The yoyo problem. *Journal of Object-Oriented Programming*, 2(3):30–35, 1989.
- [TH82] M. Turoff and S. R. Hiltz. Computer support for group versus individual decisions. *IEEE Transactions on Communications*, 30:82–90, 1982.
- [Tho67] James D. Thompson. *Organizations in Action*. McGraw-Hill, New York, 1967.
- [TJS96] Milind Tambe, Lewis Johnson, and Wei-Min Shen. Adaptive agent tracking in real-world multiagent domains: A preliminary report. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [TKOH96] Y. Tahara, F. Kumeno, A. Ohsuga, and S. Honiden. Formal semantics of agent evolution in the language Flage. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [TM89] Moshe Tennenholtz and Yoram Moses. On cooperation in a multi-entity model (preliminary report). In *IJCAI*, pages 918–923, 1989.
- [TMC81] Perry W. Thorndyke, David McArthur, and Stephanie Cammarata. Autopilot: A distributed planner for air fleet control. In *Proceedings of the Seventh International Joint Conference on Artificial Intelligence*, pages 171–177, Vancouver, Canada, August 1981.
- [TOR95] Walt Truszkowski, Jidé Odubiyi, and Ed Ruberton. An agent model in a multiagent system architecture for automating distributed systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 463, San Francisco, CA, 1995. MIT Press. (poster).
- [Tos84] Henry L. Tosi. *Theories of organization*. John Wiley and Sons, New York, second edition edition, 1984.
- [Tou84] David S. Touretzky. *Lisp: A gentle introduction to symbolic computation*. Harper & Row, 1984.

- [TS81a] Robert R. Tenney and Nils R. Sandell, Jr. Strategies for distributed decisionmaking. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(8):527–538, August 1981.
- [TS81b] Robert R. Tenney and Nils R. Sandell, Jr. Structures for distributed decisionmaking. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(8):517–527, August 1981. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 236–247, Morgan Kaufmann, 1988.).
- [Tub84] Stewart L. Tubbs. *A Systems Approach to Small Group Interaction*. Addison-Wesley, second edition, 1984.
- [TW84] A. Tate and A. M. Whiter. Planning with multiple resource constraints and an application to a naval planning problem. *IEEE Transactions on Computers*, pages 410–416, 1984.
- [TW86] Masaki Togai and Hiroyuki Watanabe. Expert system on a chip: An engine for real-time approximate reasoning. *IEEE Expert*, pages 55–62, 1986.
- [tWM85] Yung terng Wang and Robert J. T. Morris. Load sharing in distributed systems. *IEEE Transactions on Computers*, C-34(3):204–217, March 1985.
- [Tza95] Elpida S. Tzafestas. Beyond cooperation-ism and competition-ism (exploring social phenomena with a generalized tit-for-tat model). In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 464, San Francisco, CA, 1995. MIT Press. (poster).
- [Van94] W. Van de Velde. A constructivist view on knowledge engineering. In T. Cohn, editor, *Proceedings of the European Conference on Artificial Intelligence*, pages 729–734. John Wiley and Sons, 1994.
- [vB83] Gregor v. Bochmann. *Distributed Systems Design*. Springer-Verlag, Berlin, 1983.
- [VD95a] José M. Vidal and Edmund H. Durfee. Recursive agent modeling using limited rationality. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 376–383, San Francisco, CA, 1995. MIT Press.
- [VD95b] José M. Vidal and Edmund H. Durfee. Task planning agents in the UMDL. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Ver83] S. A. Vere. Planning in time: Windows and durations for activities and goals. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 5(3):246–267, May 1983.
- [VL90] R. R. Vane and P. E. Lehner. Hypergames and AI in automated adversarial planning. In *Proceedings of the 1990 DARPA Planning Workshop*, pages 198–206, November 1990.
- [vLFD96] B. van Linder F. Dignum. Modelling rational agents in a dynamic environment. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [VLV87] H. Velthuijsen, B. J. Lippolt, and J. C. Vonk. A parallel blackboard system for robot control. In *Proceedings of the Tenth International Joint Conference on Artificial Intelligence*, pages 1157–1159, August 1987.
- [vLvdHM96] B. van Linder, W. van der Hoek, and J-J Ch. Meyer. Formalising motivational attitudes. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [vM89] Frank von Martial. Multiagent plan relationships. In *Proceedings of the 1989 Distributed AI Workshop*, pages 59–72, September 1989.



- [vM90a] Frank v. Martial. A conversation model for resolving conflicts among distributed office activities. In *Proceedings of the Fifth Conference on Office Information Systems*, pages 99–108, Cambridge, MA, April 1990.
- [vM90b] Frank von Martial. Coordination of plans in multiagent worlds by taking advantage of the favor relation. In *Proceedings of the 1990 Distributed AI Workshop*, October 1990.
- [vM90c] Frank von Martial. Interactions among autonomous planning agents. In Y. Demazeau and J.-P. Muller, editors, *Decentralized AI*, pages 105–119. North Holland, 1990.
- [vM91] Frank v. Martial. *Coordinating Plans of Autonomous Agents via Relationship Resolution and Communication*. PhD thesis, Universit at des Saarlandes, 1991.
- [vM92] Frank v. Martial. *Coordinating Plans of Autonomous Agents*. Springer-Verlag, Berlin, 1992. Lecture Notes in Artificial Intelligence no. 610.
- [VMDS87] Ralph J. Verrilli, Kenneth L. Meunier, John R. Dixon, and Melvin K. Simmons. Iterative respecification management: A model for problem-solving networks in mechanical design. In *Proceedings of the ASME Computers in Engineering Conference*, August 1987.
- [VP88] J. F. Voss and T. A. Post. On the solving of ill-structured problems. In M. Chi R. Glaser and M. Farr, editors, *The Nature of Expertise*. Lawrence Erlbaum Associates, 1988.
- [Vro84] V. H. Vroom. A normative model of managerial decision making. In D. S. Pugh, editor, *Organization Theory*, chapter 16, pages 256–276. Penguin Books, 1984.
- [Š91] D. D. Šiljak. *Decentralized Control of Complex Systems*. Academic Press, Inc., Boston, 1991.
- [VS93] Angel Velez-Sosa. Division of labor in the central place foraging behavior problem. In John R. Koza, editor, *Artificial Life at Stanford 1993*, Stanford, California, December 1993. Stanford Bookstore.
- [Wag96a] Gerd Wagner. A logical and operational model of scalable knowledge- and perception-based agents. In Rudy van Hoe, editor, *Seventh European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, Eindhoven, The Netherlands, January 1996.
- [Wag96b] Gerd Wagner. A logical and operational model of scalable knowledge-based agents. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [WAHC94] D. L. Westbrook, S. D. Anderson, D. M. Hart, and P. R. Cohen. Common lisp instrumentation package: User manual. Technical Report 94-26, Department of Computer Science, University of Massachusetts, 1994.
- [Wal75] David Waltz. Understanding line drawings of scenes with shadows. In Patrick Winston, editor, *The Psychology of Computer Vision*, pages 19–91. McGraw-Hill, 1975.
- [Wal77] Richard Waldinger. Achieving several goals simultaneously. In *Machine Intelligence 8*, pages 94–136, 1977.
- [Wat83] Richard C. Waters. LetS: An expressional loop notation. A. I. Memo 680a, Artificial Intelligence Laboratory, Massachusetts Institute of Technology, Cambridge, MA 02139, February 1983.
- [Wat86] D. A. Waterman. *A Guide to Expert Systems*. Addison Wesley, 1986.
- [WB87] J. Whiteside and J. Bennett. Usability engineering: Our experience and evolution. In M. Helander, editor, *Handbook of Computer-Human Interaction*. North Holland, 1987.
- [WBL81] Jack C. Wileden, Peter C. Bates, and Victor R. Lesser. EDL: An event description language for distributed debugging. Technical Report 81-07, Department of Computer and Information Science, University of Massachusetts, Amherst, MA 01003, March 1981.

- [WD91] Michael P. Wellman and Jon Doyle. Preferential semantics for goals. In *Proceedings of the Ninth National Conference on Artificial Intelligence*, pages 698–703, Anaheim, CA, July 1991.
- [Web47] Max Weber. *The Theory of Social and Economic Organization*. Oxford University Press, New York, 1947. Translated and edited by A. M. Henderson and T. Parsons.
- [Wei93] Gerhard Weiß. Learning to coordinate actions in multi-agent systems. In *Proceedings of the International Joint Conference on Artificial Intelligence*, pages 311–316, August 1993.
- [Wel92] Michael P. Wellman. A general-equilibrium approach to distributed transportation planning. In *Proceedings of the Tenth National Conference on Artificial Intelligence*, pages 282–289, July 1992.
- [Wel93] Michael Wellman. A market-oriented programming environment and its application to distributed multicommodity flow problems. *Journal of Artificial Intelligence Research*, 1:1–23, 1993.
- [Wel94] Michael P. Wellman. A computational market model for distributed configuration design. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 401–407, Seattle, WA, 1994.
- [Wer89] Eric Werner. Cooperating agents: A unified theory of communication and social structure. In Les Gasser and Michael N. Huhns, editors, *Distributed Artificial Intelligence*, volume 2 of *Research Notes in Artificial Intelligence*. Pitman, 1989.
- [Wer90a] Keith Werkman. Knowledge-based model of negotiation using shareable perspectives. In *Proceedings of the Tenth International Workshop on Distributed AI*, Texas, October 1990.
- [Wer90b] Keith J. Werkman. *Multiagent Cooperative Problem-Solving through Negotiation and Sharing of Perspectives*. PhD thesis, Lehigh University, May 1990.
- [Wer90c] Keith James Werkman. *Multiagent Cooperative Problem Solving Through Negotiation and Perspective Sharing*. PhD thesis, Lehigh University, 1990.
- [Wer95] Keith J. Werkman. Negotiation amongst intelligent agents during information retrieval from the web. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Wes77] Robert B. Wesson. Planning in the world of the air traffic controller. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, pages 473–479, August 1977.
- [WH82] Leonard P. Wesley and Allen R. Hanson. The use of an evidential-based model for representing knowledge. In *Proceedings of the Workshop on Computer Vision: Representation and Control*, pages 14–25, August 1982.
- [WH84] Patrick Henry Winston and Berthold Klaus Paul Horn. *Lisp*. Addison-Wesley, second edition, 1984.
- [WHR89] Richard Washington and Barbara Hayes-Roth. Input data management in real-time AI systems. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, pages 250–255, Detroit, MI, August 1989.
- [WHRB<sup>+</sup>81] Robert Wesson, Frederick Hayes-Roth, John W. Burge, Cathleen Statz, and Carl A. Sunshine. Network structures for distributed situation assessment. *IEEE Transactions on Systems, Man, and Cybernetics*, 11(1):5–23, January 1981. (Also published in *Readings in Distributed Artificial Intelligence*, Alan H. Bond and Les Gasser, editors, pages 71–89, Morgan Kaufmann, 1988.).

- [Wil75] Oliver E. Williamson. *Markets and Hierarchies: Analysis and Antitrust Implications*. The Free Press, New York, 1975.
- [Wil81] Robert Wilensky. Meta-planning: Representing and using knowledge about planning in problem solving and natural language understanding. *Cognitive Science*, 5(3):197–233, July–September 1981.
- [Wil83] Robert Wilensky. *Planning and Understanding : A Computational Approach to Human Reasoning*. Addison-Wesley, Reading, Mass., 1983.
- [Wil84] D. Wilkins. Domain-independent planning: Representation and plan generation. *Artificial Intelligence*, 22:269–301, April 1984.
- [Wil86] John Davis Williams. *The Compleat Strategyst: Being a Primer on the Theory of Games of Strategy*. Dover, 1986.
- [Wil87] Robert Wilensky. *Common Lispcraft*. Norton, 1987.
- [Win77] Patrick Henry Winston. *Artificial Intelligence*. Addison-Wesley, 1977.
- [WL89a] Robert C. Whitehair and Victor R. Lesser. Approximate processing in the DVMT. Working paper, COINS, University of Massachusetts, 1989.
- [WL89b] Robert C. Whitehair and Victor R. Lesser. Real-time processing in the DVMT. In *Proceedings of the 1989 AAAI Workshop on Real-time AI*, 1989.
- [WLC95] Ipke Wachsmuth, Britta Lenzmann, and Yong Cao. VIENA: A multiagent interface to a virtual environment. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 465, San Francisco, CA, 1995. MIT Press. (poster).
- [WM94] Ian Watson and Farhi Marir. Case-based reasoning: A review. *The Knowledge Engineering Review*, 9(4), 1994.
- [Wob95] Wayne Wobcke. Plans and the revision of intentions. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [Woe95] Darrell Woelk. Developing infosleuth agents using Rosette: An actor based language. In Tim Finin and James Mayfield, editors, *Proceedings of the CIKM '95 Workshop on Intelligent Information Agents*, Baltimore, Maryland, 1995.
- [Woo77] W. A. Woods. Shortfall and density scoring strategies for speech understanding control. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, pages 13–26, August 1977.
- [Woo83] W. A. Woods. What’s important about knowledge representation. *IEEE Computer*, 16(10):22–27, October 1983.
- [Woo96] M. Wooldridge. A logic for BDI planning agents. In Pierre-Yves Schobbens, editor, *Working Notes of 3rd ModelAge Workshop: Formal Models of Agents*, Sesimbra, Portugal, January 1996.
- [WOS83] Thomas L. Williams, Paul J. Orgren, and Carl L. Smith. Diagnosis of multiple faults in a nationwide communication network. In *IJCAI*, pages 179–181, 1983.
- [WW95] Adam Walker and Michael Wooldridge. Understanding the emergence of conventions in multi-agent systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 384–389, San Francisco, CA, 1995. MIT Press.
- [Xia95] Y. Xiang. Distributed scheduling of multiagent communication. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 390–397, San Francisco, CA, 1995. MIT Press.

- [Yag95] Ronald R. Yager. Automated multiagent preference aggregation using fuzzy quantifiers. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 466, San Francisco, CA, 1995. MIT Press. (poster).
- [Yan86] Ronald M. Yannone. Expert systems in the fighter of the 1990s. *IEEE AES Magazine*, pages 12–16, February 1986.
- [YBS88] A. Yonezawa, J. Briot, and E. Shibayama. Object-oriented concurrent programming in ABCL/1. In Alan H. Bond and Les Gasser, editors, *Readings in Distributed Artificial Intelligence*, pages 434–444. Morgan Kaufmann, 1988.
- [YD91] Makoto Yokoo and Edmund H. Durfee. Distributed constraint optimization as a formal model of partially adversarial cooperation. Technical Report CSE-TR-101-91, Department of Electrical Engineering and Computer Science, University of Michigan, Ann Arbor, MI 48109, July 1991.
- [YDIK91] Makoto Yokoo, Edmund H. Durfee, Toru Ishida, and Kazuhiro Kuwabara. Distributed constraint satisfaction for formalizing distributed problem solving. Technical Report CSE-TR-102-91, Department of Electrical Engineering and Computer Science, University of Michigan, Ann Arbor, MI 48109, July 1991.
- [YDIK92] M. Yokoo, E. Durfee, T. Ishida, and K. Kuwabara. Distributed constraint satisfaction for formalizing distributed problem solving. In *Proceedings of the Twelfth International Conference on Distributed Computing Systems*, pages 614–621, 1992.
- [YH95] Michael Youssefmir and Bernardo A. Huberman. Resource contention in multiagent systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, pages 398–403, San Francisco, CA, 1995. MIT Press.
- [YHS85] Ju-Yuan D. Yang, Micheal N. Huhns, and Larry M. Stephens. An architecture for control and communications in distributed artificial intelligence systems. *IEEE Transactions on Systems, Man, and Cybernetics*, 15(3):13–23, May/June 1985.
- [YIK90] Makoto Yokoo, Toru Ishida, and Kazuhiro Kuwabara. Distributed constraint satisfaction for DAI problems. In *Proceedings of the 1990 Distributed AI Workshop*, Bandara, TX, October 1990.
- [Yok95] Makoto Yokoo. Asynchronous weak-commitment search for solving large-scale distributed constraint satisfaction problems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 467, San Francisco, CA, 1995. MIT Press. (poster).
- [YZ95] Hongxia Yang and Chengqi Zhang. Definition and application of a comprehensive framework for distributed problem solving. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.
- [Zad83] Lotfi A. Zadeh. Commonsense knowledge representation based on fuzzy logic. *IEEE Computer*, pages 61–64, October 1983.
- [Zha86] Wei Zhao. *A Heuristic Approach to Scheduling Hard Real-Time Tasks with Resource Requirements in Distributed Systems*. PhD thesis, University of Massachusetts, Amherst, Massachusetts 01003, February 1986.
- [Zis78] Michael D. Zisman. Use of production systems for modeling asynchronous, concurrent processes. In D. A. Waterman and Frederick Hayes-Roth, editors, *Pattern-Directed Inference Systems*, pages 53–68. Academic Press, 1978.
- [ZR89] Gilad Zlotkin and Jeffrey S. Rosenschein. Negotiation and task sharing among autonomous agents in cooperative domains. In *Proceedings of the Eleventh International Joint Conference on Artificial Intelligence*, pages 912–917, August 1989.

- [ZR90a] Gilad Zlotkin and Jeffrey S. Rosenschein. Blocks, lies, and postal freight: Nature of deception in negotiation. In *Proceedings of the 1990 Distributed AI Workshop*, October 1990.
- [ZR90b] Gilad Zlotkin and Jeffrey S. Rosenschein. Blocks, lies, and postal freight: The nature of deception in negotiation. In *Proceedings of the Tenth International Workshop on Distributed AI*, Texas, October 1990.
- [ZR90c] Gilad Zlotkin and Jeffrey S. Rosenschein. Negotiation and conflict resolution in non-cooperative domains. In *Proceedings of the National Conference on Artificial Intelligence*, pages 100–105, July 1990.
- [ZR91a] G. Zlotkin and J. S. Rosenschein. Incomplete information and deception in multi-agent negotiation. In *Proceedings of the Twelfth International Joint Conference on Artificial Intelligence*, pages 225–231, Sydney, Australia, August 1991.
- [ZR91b] Gilad Zlotkin and Jeffrey S. Rosenschein. Cooperation and conflict resolution via negotiation among autonomous agents in non-cooperative domains. *IEEE Transactions on Systems, Man, and Cybernetics*, 21(6), December 1991. (Special Issue on Distributed AI).
- [ZR94] Gilad Zlotkin and Jeffrey S. Rosenschein. Coalition, cryptography, and stability: Mechanisms for coalition formation in task oriented domains. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 432–437, Seattle, WA, 1994.
- [ZRS87a] W. Zhao, K. Ramamritham, and J. A. Stankovic. Scheduling tasks with resource requirements in hard real-time systems. *IEEE Transactions on Software Engineering*, May 1987.
- [ZRS87b] Wei Zhao, Krithi Ramamritham, and John A. Stankovic. Preemptive scheduling under time and resource constraints. *IEEE Transactions on Computers*, C-36(8):949–960, August 1987.
- [ZS96] Dajun Zeng and Katia Sycara. Bayesian learning in negotiation. In Sandip Sen, editor, *Working Notes for the AAAI Symposium on Adaptation, Co-evolution and Learning in Multiagent Systems*, Stanford University, CA, March 1996.
- [ZZ95a] Minjie Zhang and Chengqi Zhang. The consensus of uncertainties in distributed expert systems. In Victor Lesser, editor, *Proceedings of the First International Conference on Multi-Agent Systems*, page 468, San Francisco, CA, 1995. MIT Press. (poster).
- [ZZ95b] Minjie Zhang and Chengqi Zhang. Neural network strategies for solving synthesis problems in non-conflict cases in distributed expert systems. In Chengqi Zhang and Dickson Lukose, editors, *First Australian Workshop on Distributed Artificial Intelligence*, Canberra, Australia, 1995.